

# Language and Compiler Support for Stream Programs

Bill Thies

---

Computer Science and Artificial Intelligence Laboratory  
Massachusetts Institute of Technology

Thesis Defense  
September 11, 2008

**Date:** Wed, 17 Nov 1999

**From:** Saman Amarasinghe <saman@lcs.mit.edu>

**To:** Bill Thies <thies@mit.edu>

**Subject:** UROP Opportunities

Hi Bill,

I have a few UROP opportunities in the RAW project ...


Most of the projects can lead to an MENG thesis and beyond ...



# Acknowledgments


- **Project supervisors**
  - Prof. Saman Amarasinghe
  - Dr. Rodric Rabbah
- **Contributors to this talk**
  - Michael I. Gordon (Ph.D. student) – *led development of Raw backend*
  - Andrew A. Lamb (M.Eng) – *led development of linear optimizations*
  - Sitij Agrawal (M.Eng) – *led development of statespace optimizations*
- **Compiler developers**
  - Kunal Agrawal
  - Allyn Dimock
  - Steve Hall
  - Qiuyuan Jimmy Li
  - Jasper Lin
  - Michal Karczmarek
  - David Maze
  - Janis Sermulins
  - Phil Sung
  - Ceryen Tan
  - David Zhang
- **Application developers**
  - Basier Aziz
  - Matthew Brown
  - Jiawen Chen
  - Matthew Drake
  - Shirley Fung
  - Hank Hoffmann
  - Chris Leger
  - Ali Meli
  - Mani Narayanan
  - Satish Ramaswamy
  - Jeremy Wong
- **User interface developers**
  - Kimberly Kuo
  - Juan Reyes

# Multicores are Here





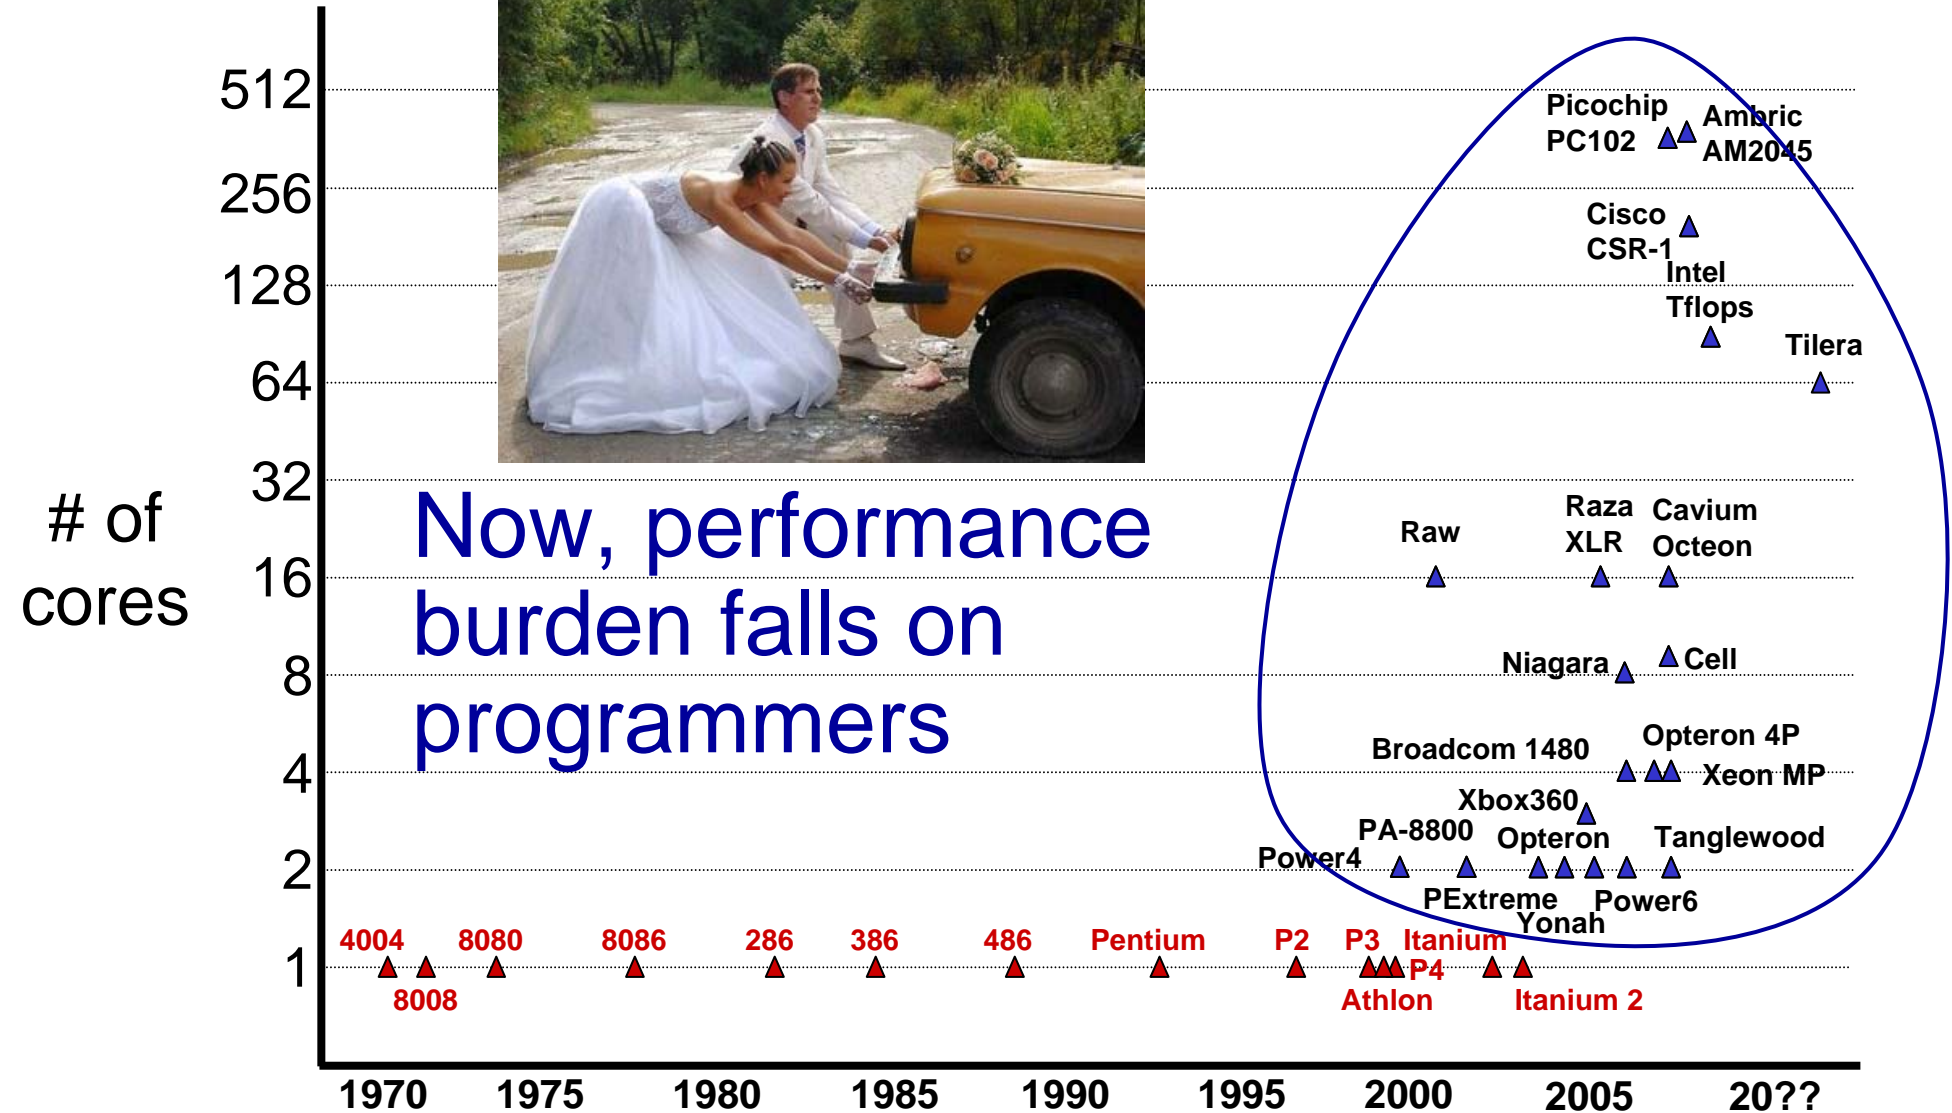
# Multicores are Here



# Multicores are Here



Now, performance burden falls on programmers




# Is Parallel Programming a New Problem?

- **No! Decades of research targeting multiprocessors**
  - Languages, compilers, architectures, tools...
- **What is different today?**
  - 1. Multicores vs. multiprocessors.** Multicores have:
    - New interconnects with non-uniform communication costs
    - Faster on-chip communication than off-chip I/O, memory ops
    - Limited per-core memory availability
  - 2. Non-expert programmers**
    - Supercomputers with >2048 processors today: 100 [\[top500.org\]](#)
    - Machines with >2048 cores in 2020: >100 million [\[ITU, Moore\]](#)
  - 3. Application trends**
    - Embedded: 2.7 billion cell phones vs 850 million PCs [\[ITU 2006\]](#)
    - Data-centric: YouTube streams 200 TB of video daily



# Streaming Application Domain

- **For programs based on streams of data**
  - Audio, video, DSP, networking, and cryptographic processing kernels
  - Examples: HDTV editing, radar tracking, microphone arrays, cell phone base stations, graphics




# Streaming Application Domain

- **For programs based on streams of data**


- Audio, video, DSP, networking, and cryptographic processing kernels
- Examples: HDTV editing, radar tracking, microphone arrays, cell phone base stations, graphics

- **Properties of stream programs**

- Regular and repeating computation
- Independent filters with explicit communication
- Data items have short lifetimes



# Brief History of Streaming



## Models of Computation

Petri Nets

Kahn Proc.

Networks

Synchronous Dataflow

Comp. Graphs

Communicating

Sequential Processes

## Modeling Environments

Ptolemy

Matlab/Simulink

Gabriel

Grape-II

etc.

## Languages / Compilers

Lucid

Id

Sisal

Erlang

Estérel

C

lazy


VAL

Occam

LUSTRE

pH

# Brief History of Streaming



## Models of Computation

Petri Nets  
Comp. Graphs

Kahn Proc.  
Communicating

Networks  
Sequential Processes

Synchronous Dataflow

## Modeling Environments

Ptolemy

Gabriel

Matlab/Simulink

Grape-II

etc.

## Languages / Compilers

Lucid

Id

Sisal

Erlang

Estérel

C

lazy

VAL

Occam

LUSTRE

pH


### Strengths

- Elegance
- Generality

### Weaknesses

- Unsuitable for static analysis
- Cannot leverage deep results from DSP / modeling community

# Brief History of Streaming



Models of Computation

Petri Nets  
Comp. Graphs

Kahn Proc.  
Communicating

Networks  
Sequential Processes

Synchronous Dataflow

Modeling Environments

Ptolemy  
Gabriel

Matlab/Simulink  
Grape-II

etc.

Languages / Compilers

Lucid  
C

Id  
lazy

VAL

Sisal  
Occam

Erlang  
LUSTRE

Estérel  
pH

StreamIt  
Cg

Brook  
StreamC

## Strengths

- Elegance
- Generality

## Weaknesses

- Unsuitable for static analysis
- Cannot leverage deep results from DSP / modeling community

*"Stream Programming"*



# StreamIt: A Language and Compiler for Stream Programs

- **Key idea:** design language that enables static analysis
  - **Goals:**
    1. Expose and exploit the parallelism in stream programs
    2. Improve programmer productivity in the streaming domain
  - **Project contributions:**
    - Language design for streaming [CC'02, CAN'02, PPOPP'05, IJPP'05]
    - Automatic parallelization [ASPLOS'02, G.Hardware'05, ASPLOS'06]
    - Domain-specific optimizations [PLDI'03, CASES'05, TechRep'07]
    - Cache-aware scheduling [LCTES'03, LCTES'05]
    - Extracting streams from legacy code [MICRO'07]
    - User + application studies [PLDI'05, P-PHEC'05, IPDPS'06]
- **7 years, 25 people, 300 KLOC**
  - **700 external downloads, 5 external publications**

# StreamIt: A Language and Compiler for Stream Programs

- **Key idea:** design language that enables static analysis
- **Goals:**
  1. Expose and exploit the parallelism in stream programs
  2. Improve programmer productivity in the streaming domain
- **I contributed to:**
  - Language design for streaming [CC'02, CAN'02, PPOPP'05, IJPP'05]
  - Automatic parallelization [ASPLOS'02, G.Hardware'05, ASPLOS'06]
  - Domain-specific optimizations [PLDI'03, CASES'05, TechRep'07]
  - Cache-aware scheduling [LCTES'03, LCTES'05]
  - Extracting streams from legacy code [MICRO'07]
  - User + application studies [PLDI'05, P-PHEC'05, IPDPS'06]
  - **7 years, 25 people, 300 KLOC**
  - **700 external downloads, 5 external publications**

# StreamIt: A Language and Compiler for Stream Programs

- **Key idea:** design language that enables static analysis
- **Goals:**
  1. Expose and exploit the parallelism in stream programs
  2. Improve programmer productivity in the streaming domain
- **This talk:**
  - Language design for streaming [CC'02, CAN'02, PPOPP'05, IJPP'05]
  - Automatic parallelization [ASPLOS'02, G.Hardware'05, ASPLOS'06]
  - Domain-specific optimizations [PLDI'03, CASES'05, TechRep'07]
  - Cache-aware scheduling [LCTES'03, LCTES'05]
  - Extracting streams from legacy code [MICRO'07]
  - User + application studies [PLDI'05, P-PHEC'05, IPDPS'06]
  - **7 years, 25 people, 300 KLOC**
  - **700 external downloads, 5 external publications**


# Part 1: Language Design

William Thies, Michal Karczmarek, Saman Amarasinghe (CC'02)

William Thies, Michal Karczmarek, Janis Sermulins, Rodric Rabbah,  
Saman Amarasinghe (PPoPP'05)

# StreamIt Language Basics


- **High-level, architecture-independent language**
  - Backend support for uniprocessors, multicores (Raw, SMP), cluster of workstations
- **Model of computation: synchronous dataflow** [Lee & Messerschmidt, 1987]
  - Program is a graph of independent *filters*
  - Filters have an atomic execution step with known input / output rates
  - Compiler is responsible for scheduling and buffer management
- **Extensions to synchronous dataflow**
  - Dynamic I/O rates
  - Support for sliding window operations
  - Teleport messaging [PPoPP'05]






# Representing Streams

- **Conventional wisdom: stream programs are graphs**
  - Graphs have no simple textual representation
  - Graphs are difficult to analyze and optimize
- **Insight: stream programs have structure**





*unstructured*




*structured*


# Structured Streams



**splitjoin**

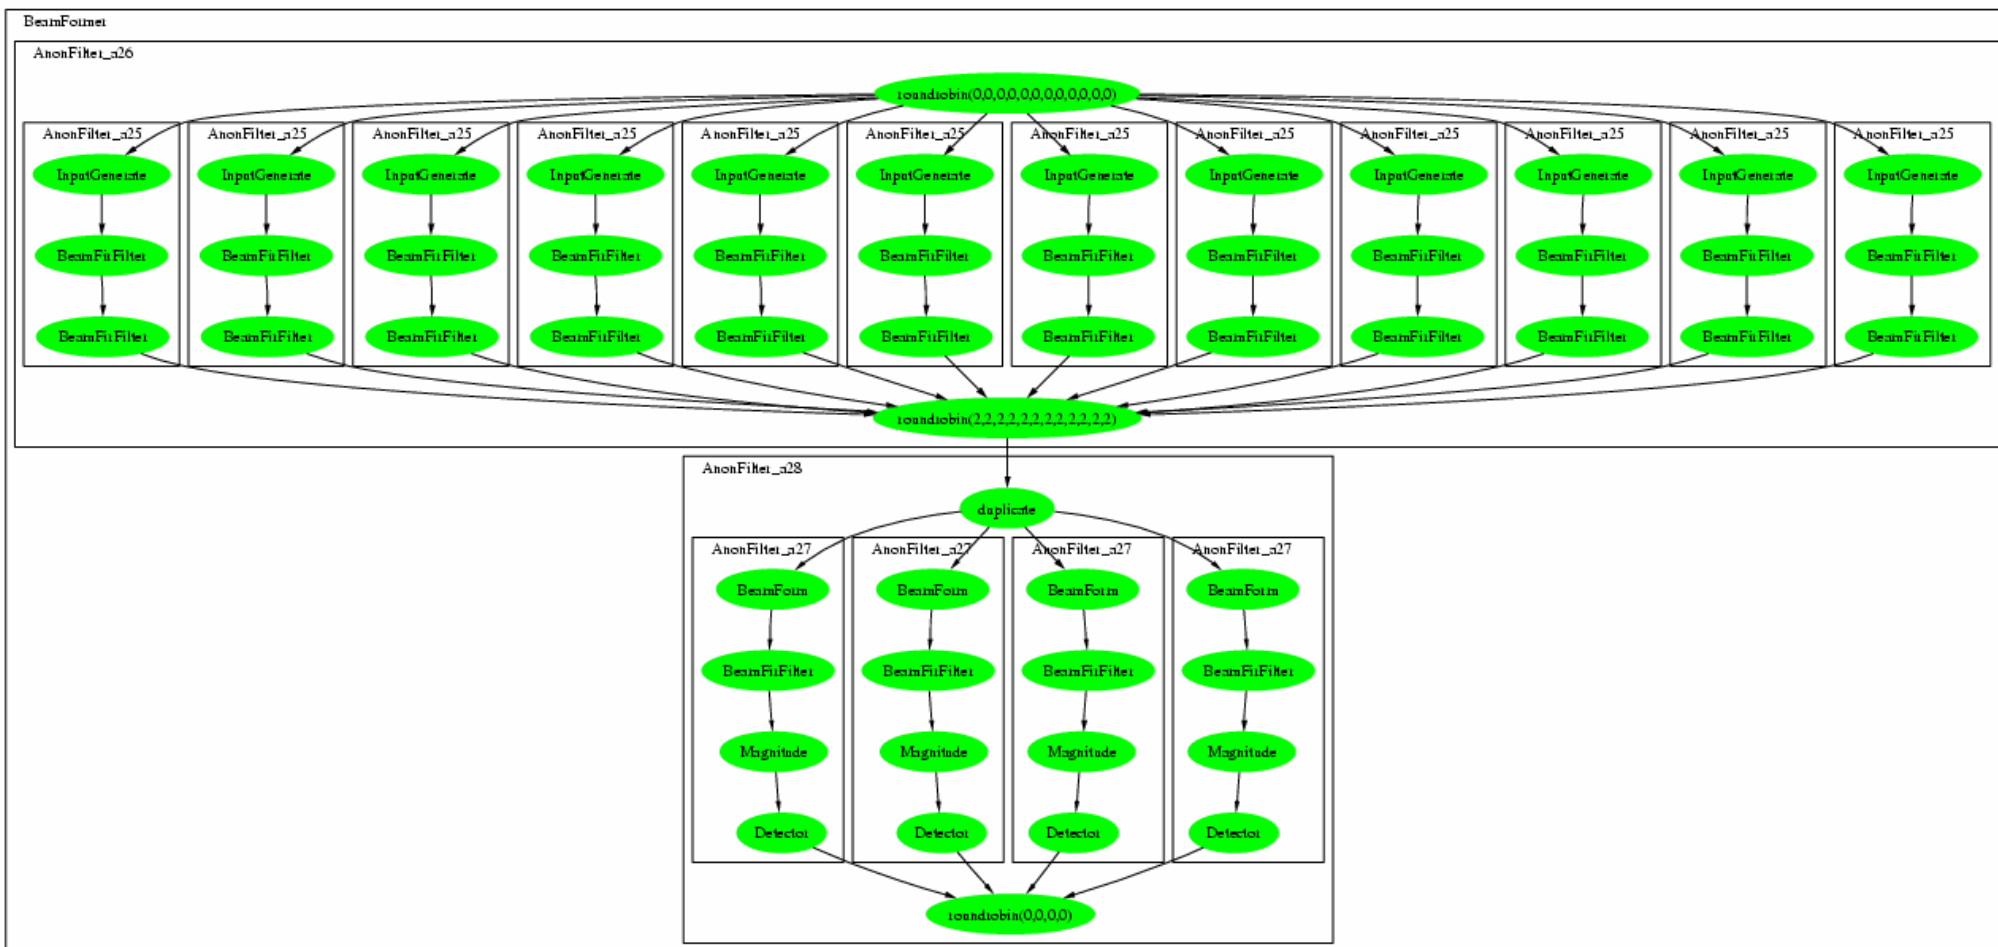


**feedback loop**




- Each structure is single-input, single-output
- Hierarchical and composable

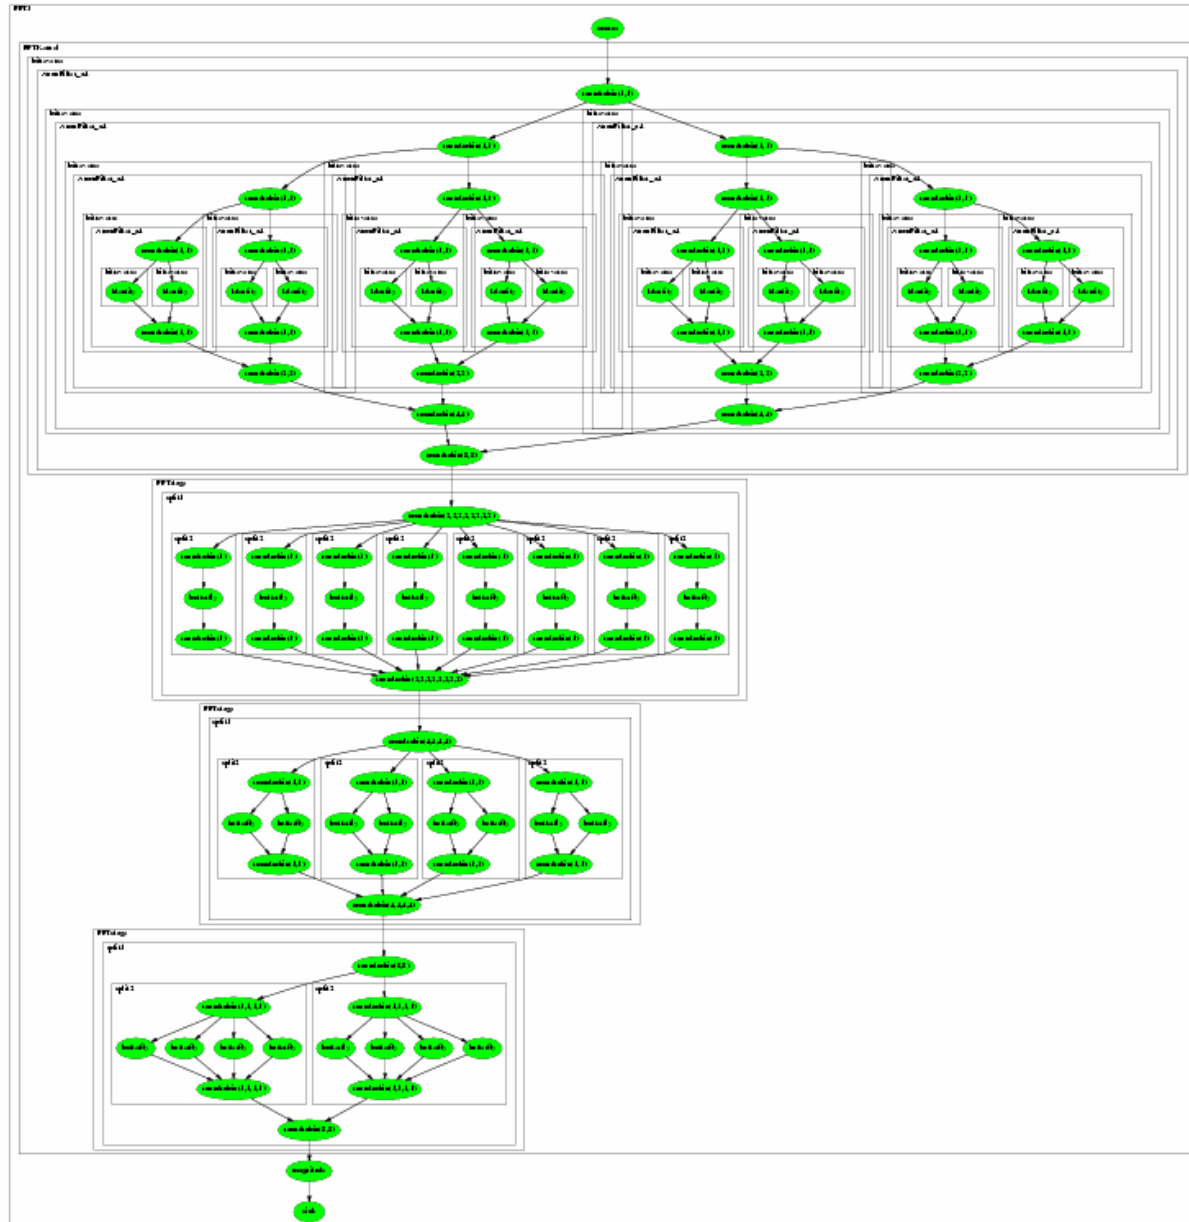
# Radar-Array Front End



# Filterbank

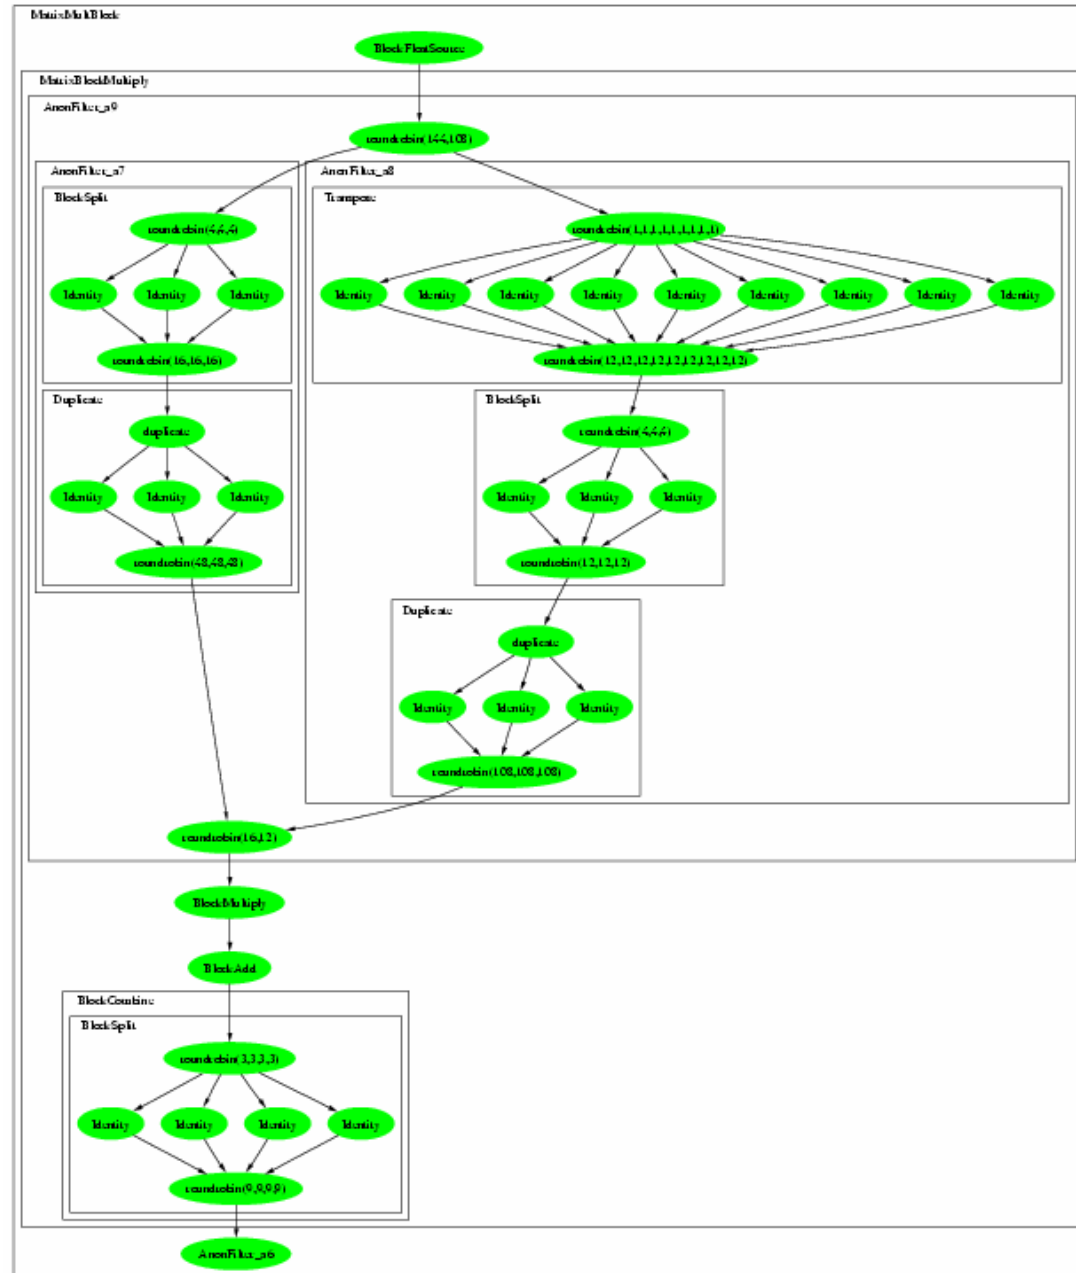


# FFT

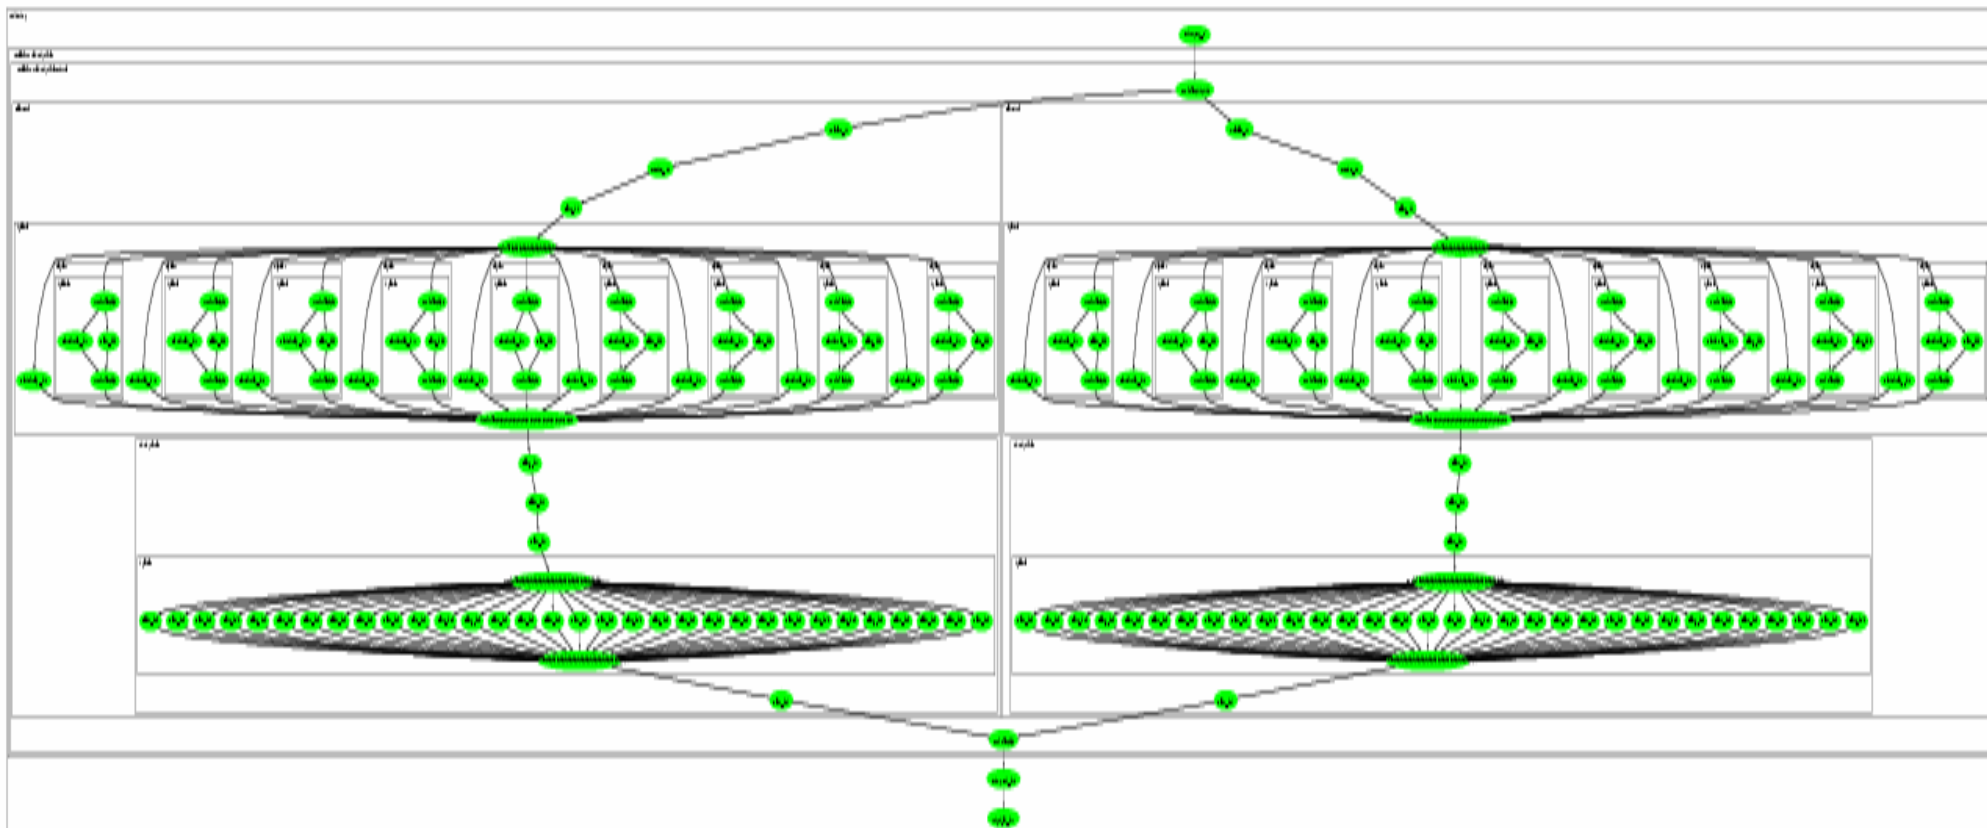





# Block Matrix Multiply



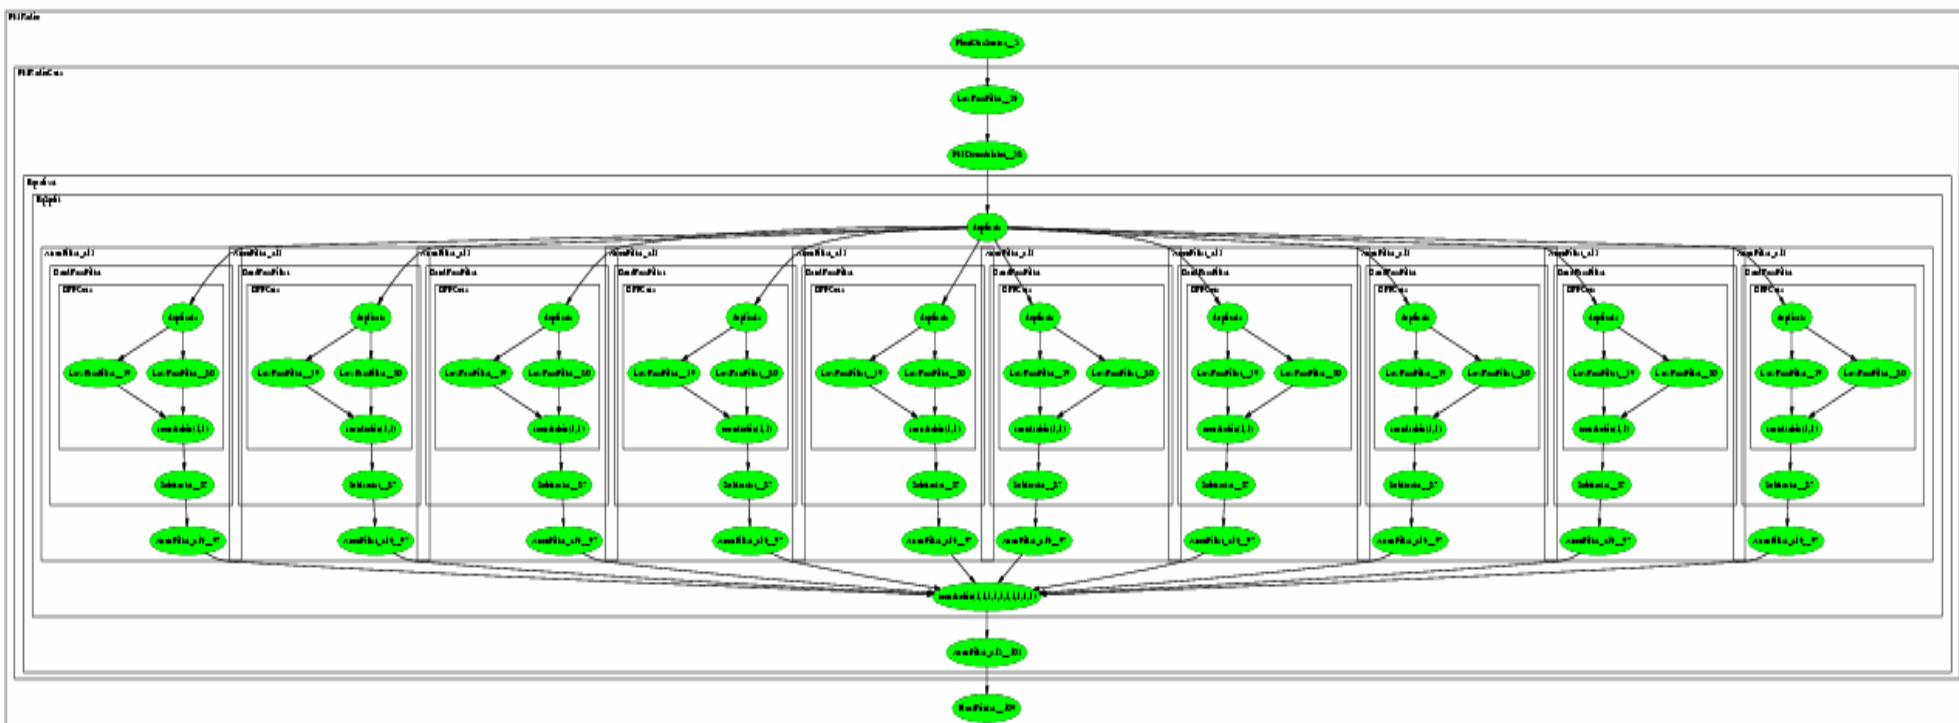
# MP3 Decoder




# Bitonic Sort



# FM Radio with Equalizer



# Ground Moving Target Indicator (GMTI)



**99 filters**

**3566 filter instances**



# Example Syntax: FMRadio

```
void->void pipeline FMRadio(int N, float lo, float hi) {
```

```
  add AtoD();
```

```
  add FMDemod();
```

```
  add splitjoin {  
    split duplicate;  
    for (int i=0; i<N; i++) {
```

```
      add pipeline {
```

```
        add LowPassFilter(lo + i*(hi - lo)/N);
```


```
        add HighPassFilter(lo + i*(hi - lo)/N);
```

```
      }
```

```
    }  
    join roundrobin();  
  }
```

```
  add Adder();
```

```
  add Speaker();  
}
```




# StreamIt Application Suite

- Software radio
- Frequency hopping radio
- Acoustic beam former
- Vocoder
- FFTs and DCTs
- JPEG Encoder/Decoder
- MPEG-2 Encoder/Decoder
- MPEG-4 (fragments)
- Sorting algorithms
- GMTI (Ground Moving Target Indicator)
- DES and Serpent crypto algorithms
- SSCA#3 (HPCS scalable benchmark for synthetic aperture radar)
- Mosaic imaging using RANSAC algorithm

Total size: 60,000 lines of code

# Control Messages

- Occasionally, low-bandwidth control messages are sent between actors
- Often demands precise timing
  - Communications: adjust protocol, amplification, compression
  - Network router: cancel invalid packet
  - Adaptive beamformer: track a target
  - Respond to user input, runtime errors
  - Frequency hopping radio
- Traditional techniques:
  - Direct method call (no timing guarantees)
  - Embed message in stream (opaque, slow)




# Idea 2: Teleport Messaging

- Looks like method call, but timed relative to data in the stream

```
TargetFilter x;  
if newProtocol(p) {  
  x.setProtocol(p) @ 2;  
}
```

```
void setProtocol(int p) {  
  reconfig(p);  
}
```

- Exposes dependences to compiler
- Simple and precise for user
  - Adjustable latency
  - Can send upstream or downstream



## Part 2: Automatic Parallelization


*Michael I. Gordon, William Thies, Saman Amarasinghe (ASPLoS'06)*

*Michael I. Gordon, William Thies, Michal Karczmarek, Jasper Lin, Ali S. Meli, Andrew A. Lamb, Chris Leger, Jeremy Wong, Henry Hoffmann, David Maze, Saman Amarasinghe (ASPLoS'02)*

# Streaming is an Implicitly Parallel Model

- **Programmer thinks about functionality, not parallelism**
- **More explicit models may...**
  - Require knowledge of target [\[MPI\]](#) [\[cG\]](#)
  - Require parallelism annotations [\[OpenMP\]](#) [\[HPF\]](#) [\[Cilk\]](#) [\[Intel TBB\]](#)
- **Novelty over other implicit models?**  
[\[Erlang\]](#) [\[MapReduce\]](#) [\[Sequoia\]](#) [\[pH\]](#) [\[Occam\]](#) [\[Sisal\]](#) [\[Id\]](#) [\[VAL\]](#) [\[LUSTRE\]](#)  
[\[HAL\]](#) [\[THAL\]](#) [\[SALSA\]](#) [\[Rosette\]](#) [\[ABCL\]](#) [\[APL\]](#) [\[ZPL\]](#) [\[NESL\]](#) [...]  
→ **Exploiting streaming structure for robust performance**


# Parallelism in Stream Programs



## Task parallelism

- Analogous to thread (fork/join) parallelism

# Parallelism in Stream Programs



## Task parallelism

- Analogous to thread (fork/join) parallelism

## Data parallelism


- Analogous to DOALL loops

## Pipeline parallelism


- Analogous to ILP that is exploited in hardware




# Baseline: Fine-Grained Data Parallelism



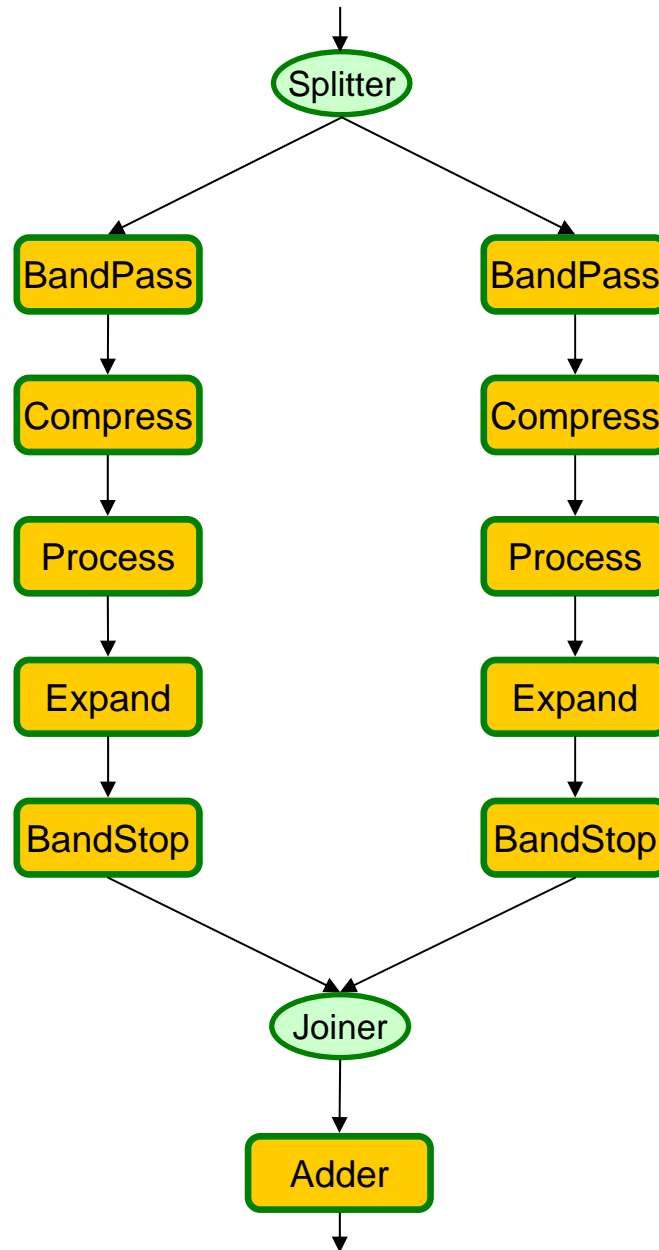
# Evaluation: Fine-Grained Data Parallelism




# Evaluation: Fine-Grained Data Parallelism




# Coarsening the Granularity



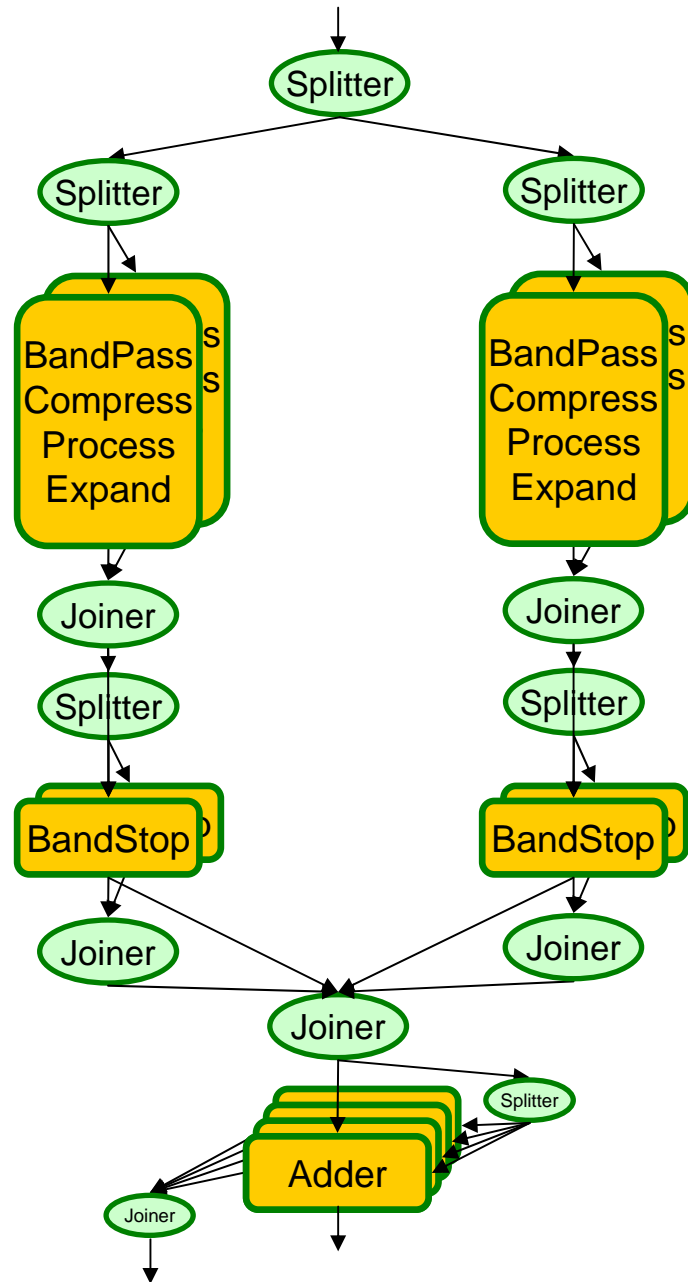
# Coarsening the Granularity




# Coarsening the Granularity



# Coarsening the Granularity




# Evaluation: Coarse-Grained Data Parallelism



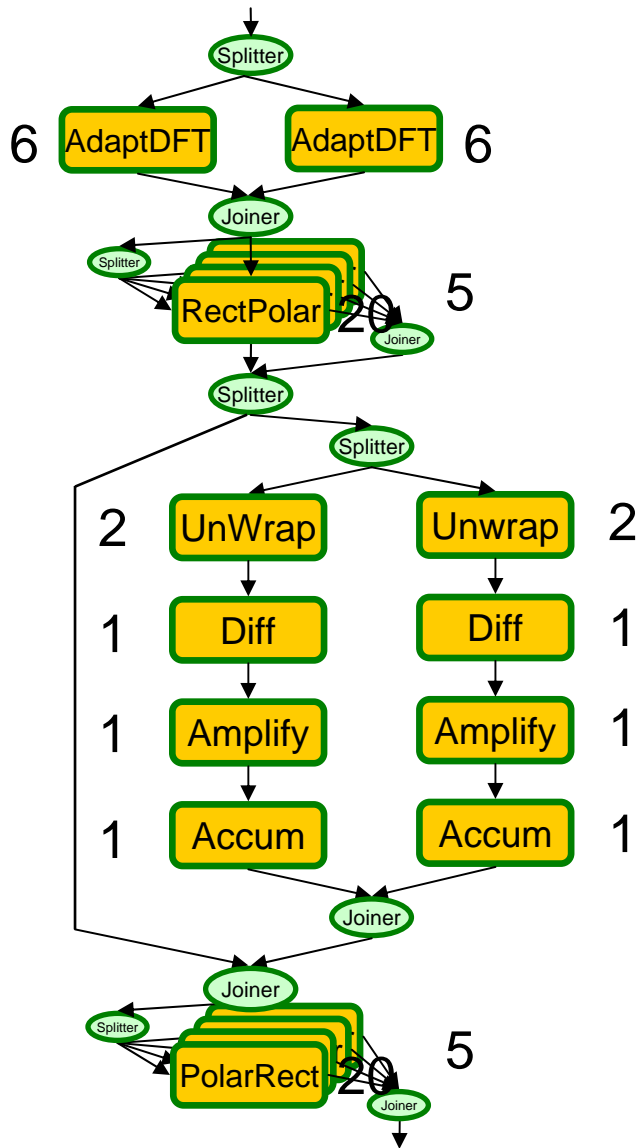


# Simplified Vocoder





Target a 4-core machine

# Data Parallelize





Target a 4-core machine

# Data + Task Parallel Execution




Target a 4-core machine

# We Can Do Better




Target a 4-core machine


# Coarse-Grained Software Pipelining



# Evaluation: Coarse-Grained Task + Data + Software Pipelining




# Evaluation: Coarse-Grained Task + Data + Software Pipelining



# Parallelism: Take Away

- **Stream programs have abundant parallelism**
  - However, parallelism is obfuscated in language like C
- **Stream languages enable new & effective mapping**



- In C, analogous transformations impossibly complex
- In StreamC or Brook, similar transformations possible  
[\[Khailany et al., IEEE Micro'01\]](#) [\[Buck et al., SIGGRAPH'04\]](#) [\[Das et al., PACT'06\]](#) [...]
- **Results should extend to other multicores**
  - Parameters: local memory, comm.-to-comp. cost
  - Preliminary results on Cell are promising [\[Zhang, dasCMP'07\]](#)




# Part 3: Domain-Specific Optimizations

*Andrew Lamb, William Thies, Saman Amarasinghe (PLDI'03)*

*Sitij Agrawal, William Thies, Saman Amarasinghe (CASES'05)*


# DSP Optimization Process

- Given specification of algorithm, minimize the computation cost




# DSP Optimization Process

- **Given specification of algorithm, minimize the computation cost**
  - Currently done by hand (MATLAB)




# DSP Optimization Process

- **Given specification of algorithm, minimize the computation cost**
  - Currently done by hand (MATLAB)
- **Can compiler replace DSP expert?**
  - Library generators limited [[Spiral](#)] [[FFTW](#)] [[ATLAS](#)]
  - Enable unified development environment




# Focus: Linear State Space Filters

- **Properties:**
  - Outputs are linear function of inputs and states
  - New states are linear function of inputs and states
- **Most common target of DSP optimizations**
  - FIR / IIR filters
  - Linear difference equations
  - Upsamplers / downsamplers
  - DCTs




# Focus: Linear State Space Filters



# Focus: Linear Filters

```
float->float filter Scale {  
  work push 2 pop 1 {  
    float u = pop();  
    push(u);  
    push(2*u);  
  }  
}
```


Linear  
dataflow  
analysis



# Focus: Linear Filters


```
float->float filter Scale {  
  work push 2 pop 1 {  
    float u = pop();  
    push(u);  
    push(2*u);  
  }  
}
```

Linear  
dataflow  
analysis





# Combining Adjacent Filters



# Combination Example

$\frac{6 \text{ mults}}{\text{output}}$



$u$

**Filter 1**

$$D = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

$y$

**Filter 2**

$$E = [4 \quad 5 \quad 6]$$

$z$

$\frac{1 \text{ mults}}{\text{output}}$



$u$

**Combined  
Filter**



$$G = [32]$$

$z$


# The General Case

- If matrix dimensions mis-match? *Matrix expansion*:

Original



Expanded



# The General Case

- If matrix dimensions mis-match? *Matrix expansion:*

$$A^e = A^n A_{\text{pre}}$$

$$B^e = \begin{bmatrix} A^n B_{\text{pre}} & A^{n-1} B & A^{n-2} B & \dots & B \end{bmatrix}$$

$$C^e = \begin{bmatrix} C A_{\text{pre}} \\ C A A_{\text{pre}} \\ \dots \\ C A^{n-1} A_{\text{pre}} \end{bmatrix}$$

$$D^e = \begin{bmatrix} C B_{\text{pre}} & D & 0 & 0 & \dots & 0 & 0 \\ C A B_{\text{pre}} & C B & D & 0 & \dots & 0 & 0 \\ C A^2 B_{\text{pre}} & C A B & C B & D & \dots & 0 & 0 \\ \dots & \dots & \dots & \dots & \dots & \dots & \\ C A^{n-1} B_{\text{pre}} & C A^{n-2} B & C A^{n-3} B & C A^{n-3} B & \dots & C B & D \end{bmatrix}$$

# The General Case

## Pipelines

$$\begin{aligned}
 A &= \begin{bmatrix} A_1 & 0 \\ B_2 C_1 & A_2 \end{bmatrix} & A_{\text{pre}} &= \begin{bmatrix} A_1^e & 0 \\ B_{\text{pre}2} C_1^e & A_{\text{pre}2} \end{bmatrix} \\
 B &= \begin{bmatrix} B_1 \\ B_2 D_1 \end{bmatrix} & B_{\text{pre}} &= \begin{bmatrix} B_1^e \\ B_{\text{pre}2} D_1^e \end{bmatrix} \\
 C &= \begin{bmatrix} D_2 C_1 & C_2 \end{bmatrix} & \overrightarrow{\text{initVec}} &= \begin{bmatrix} \overrightarrow{\text{initVec}_1} \\ \overrightarrow{\text{initVec}_2} \end{bmatrix} \\
 D &= D_2 D_1
 \end{aligned}$$

## Feedback Loops


$$\begin{aligned}
 \vec{x}_1 &= A_1 \vec{x}_1 + B_1 \vec{u}_1 = A_1 \vec{x}_1 + B_1 \vec{y} = A_1 \vec{x}_1 + B_1 (C_2 \vec{x}_2 + D_{2,1} \vec{u} + D_{2,2} C_3 \vec{x}_3) \\
 &= A_1 \vec{x}_1 + B_1 C_2 \vec{x}_2 + B_1 D_{2,1} \vec{u} + B_1 D_{2,2} C_3 \vec{x}_3 \\
 \vec{x}_2 &= A_2 \vec{x}_2 + B_2 \vec{u}_2 = A_2 \vec{x}_2 + B_{2,1} \vec{u} + B_{2,2} \vec{y}_3 = A_2 \vec{x}_2 + B_{2,1} \vec{u} + B_{2,2} C_3 \vec{x}_3 \\
 \vec{y}_2 &= C_2 \vec{x}_2 + D_2 \vec{u}_2 = C_2 \vec{x}_2 + D_{2,1} \vec{u} + D_{2,2} \vec{y}_3 = C_2 \vec{x}_2 + D_{2,1} \vec{u} + D_{2,2} C_3 \vec{x}_3 \\
 \vec{x}_3 &= A_3 \vec{x}_3 + B_3 \vec{u}_3 = A_3 \vec{x}_3 + B_3 \vec{y}_1 = A_3 \vec{x}_3 + B_3 (C_1 \vec{x}_1 + D_1 \vec{u}_1) \\
 &= A_3 \vec{x}_3 + B_3 (C_1 \vec{x}_1 + D_1 \vec{y}) = A_3 \vec{x}_3 + B_3 (C_1 \vec{x}_1 + D_1 (C_2 \vec{x}_2 + D_{2,1} \vec{u} + D_{2,2} C_3 \vec{x}_3)) \\
 &= A_3 \vec{x}_3 + B_3 C_1 \vec{x}_1 + B_3 D_1 C_2 \vec{x}_2 + B_3 D_1 D_{2,1} \vec{u} + B_3 D_1 D_{2,2} C_3 \vec{x}_3
 \end{aligned}$$

# The General Case


## Splitjoins

$$\begin{aligned}
 \mathbf{A} &= \begin{bmatrix} A_s & 0 & 0 & \dots & 0 \\ A_{1rs} & A_{1rr} & 0 & \dots & 0 \\ A_{2rs} & 0 & A_{2rr} & \dots & 0 \\ \dots & \dots & \dots & \dots & \dots \\ A_{krr} & 0 & 0 & \dots & A_{krs} \end{bmatrix} & \mathbf{B} &= \begin{bmatrix} B_s \\ B_{1r} \\ B_{2r} \\ \dots \\ B_{kr} \end{bmatrix} & \mathbf{C} &= \begin{bmatrix} C_{1s1} & C_{1r1} & 0 & \dots & 0 \\ C_{2s1} & C_{2r1} & 0 & \dots & 0 \\ \dots & \dots & \dots & \dots & \dots \\ C_{ks1} & 0 & 0 & \dots & C_{kr1} \\ \dots & \dots & \dots & \dots & \dots \\ C_{1sk} & C_{1rk} & 0 & \dots & 0 \\ C_{2sk} & C_{2rk} & 0 & \dots & 0 \\ \dots & \dots & \dots & \dots & \dots \\ C_{ksk} & 0 & 0 & \dots & C_{krk} \end{bmatrix} & \mathbf{D} &= \begin{bmatrix} D_{11} \\ D_{21} \\ \dots \\ D_{k1} \\ \dots \\ D_{1k} \\ D_{2k} \\ \dots \\ D_{kk} \end{bmatrix} \\
 \mathbf{C}_i &= \begin{bmatrix} C_{is1} & C_{ir1} \\ C_{is2} & C_{ir2} \\ \dots & \dots \\ C_{isexecutions} & C_{irexecutions} \end{bmatrix} & \mathbf{D}_i &= \begin{bmatrix} D_{i1} \\ D_{i2} \\ \dots \\ D_{iexecutions} \end{bmatrix} \\
 \mathbf{A}_{pre} &= \begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ 0 & A_{pre1rr} & 0 & \dots & 0 \\ 0 & 0 & A_{pre2rr} & \dots & 0 \\ \dots & \dots & \dots & \dots & \dots \\ 0 & 0 & 0 & \dots & A_{prekrr} \end{bmatrix} & \mathbf{B}_{pre} &= \begin{bmatrix} B_{pres} \\ B_{pre1r} \\ B_{pre2r} \\ \dots \\ B_{prekr} \end{bmatrix} & \overrightarrow{\text{initVec}} &= \begin{bmatrix} \vec{0} \\ \overrightarrow{\text{initVec}_{1r}} \\ \overrightarrow{\text{initVec}_{2r}} \\ \dots \\ \overrightarrow{\text{initVec}_{kr}} \end{bmatrix}
 \end{aligned}$$

# Floating-Point Operations Reduction

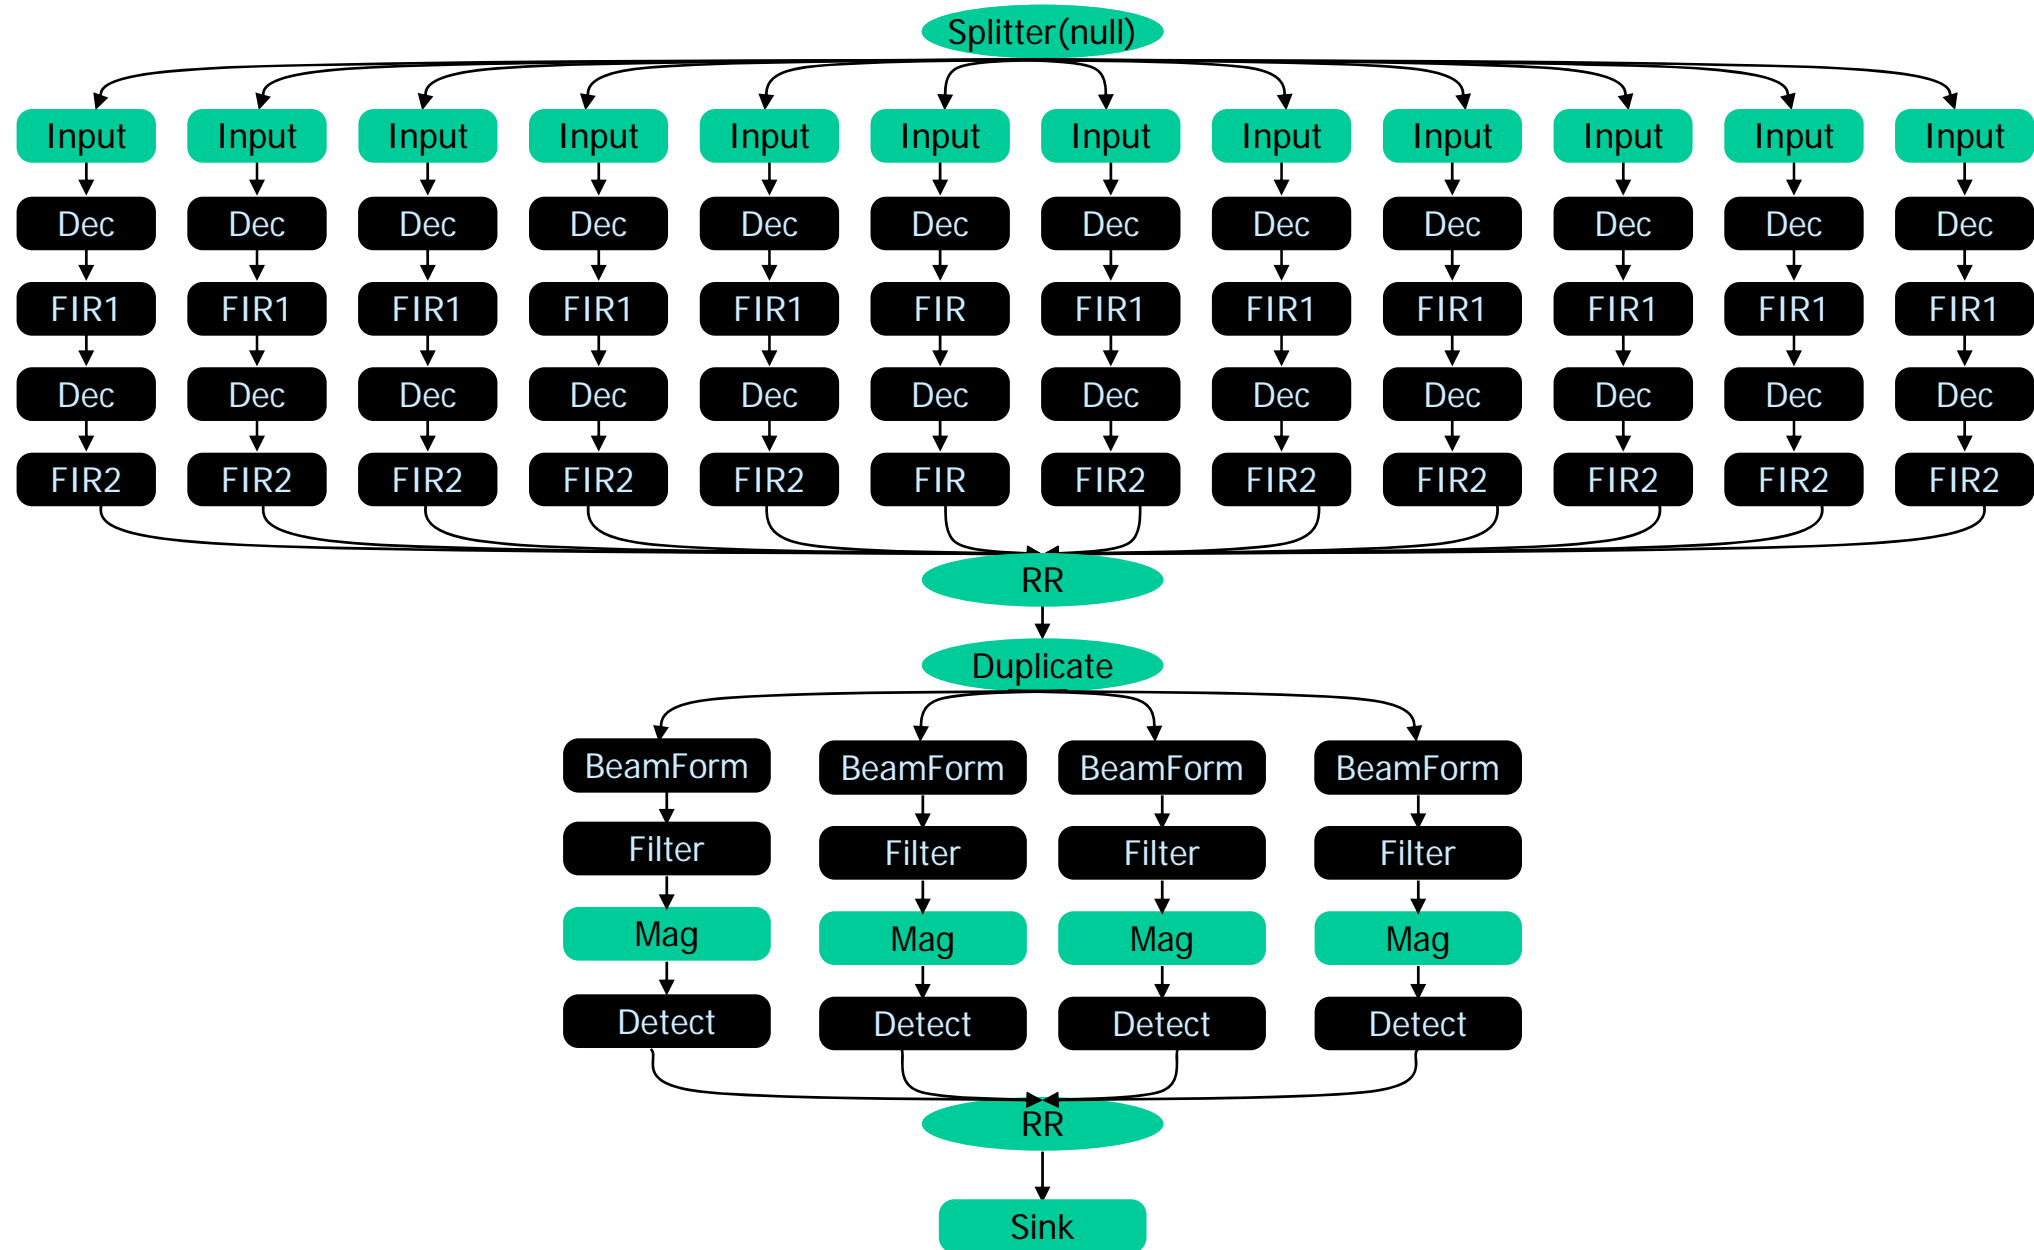


# Floating-Point Operations Reduction

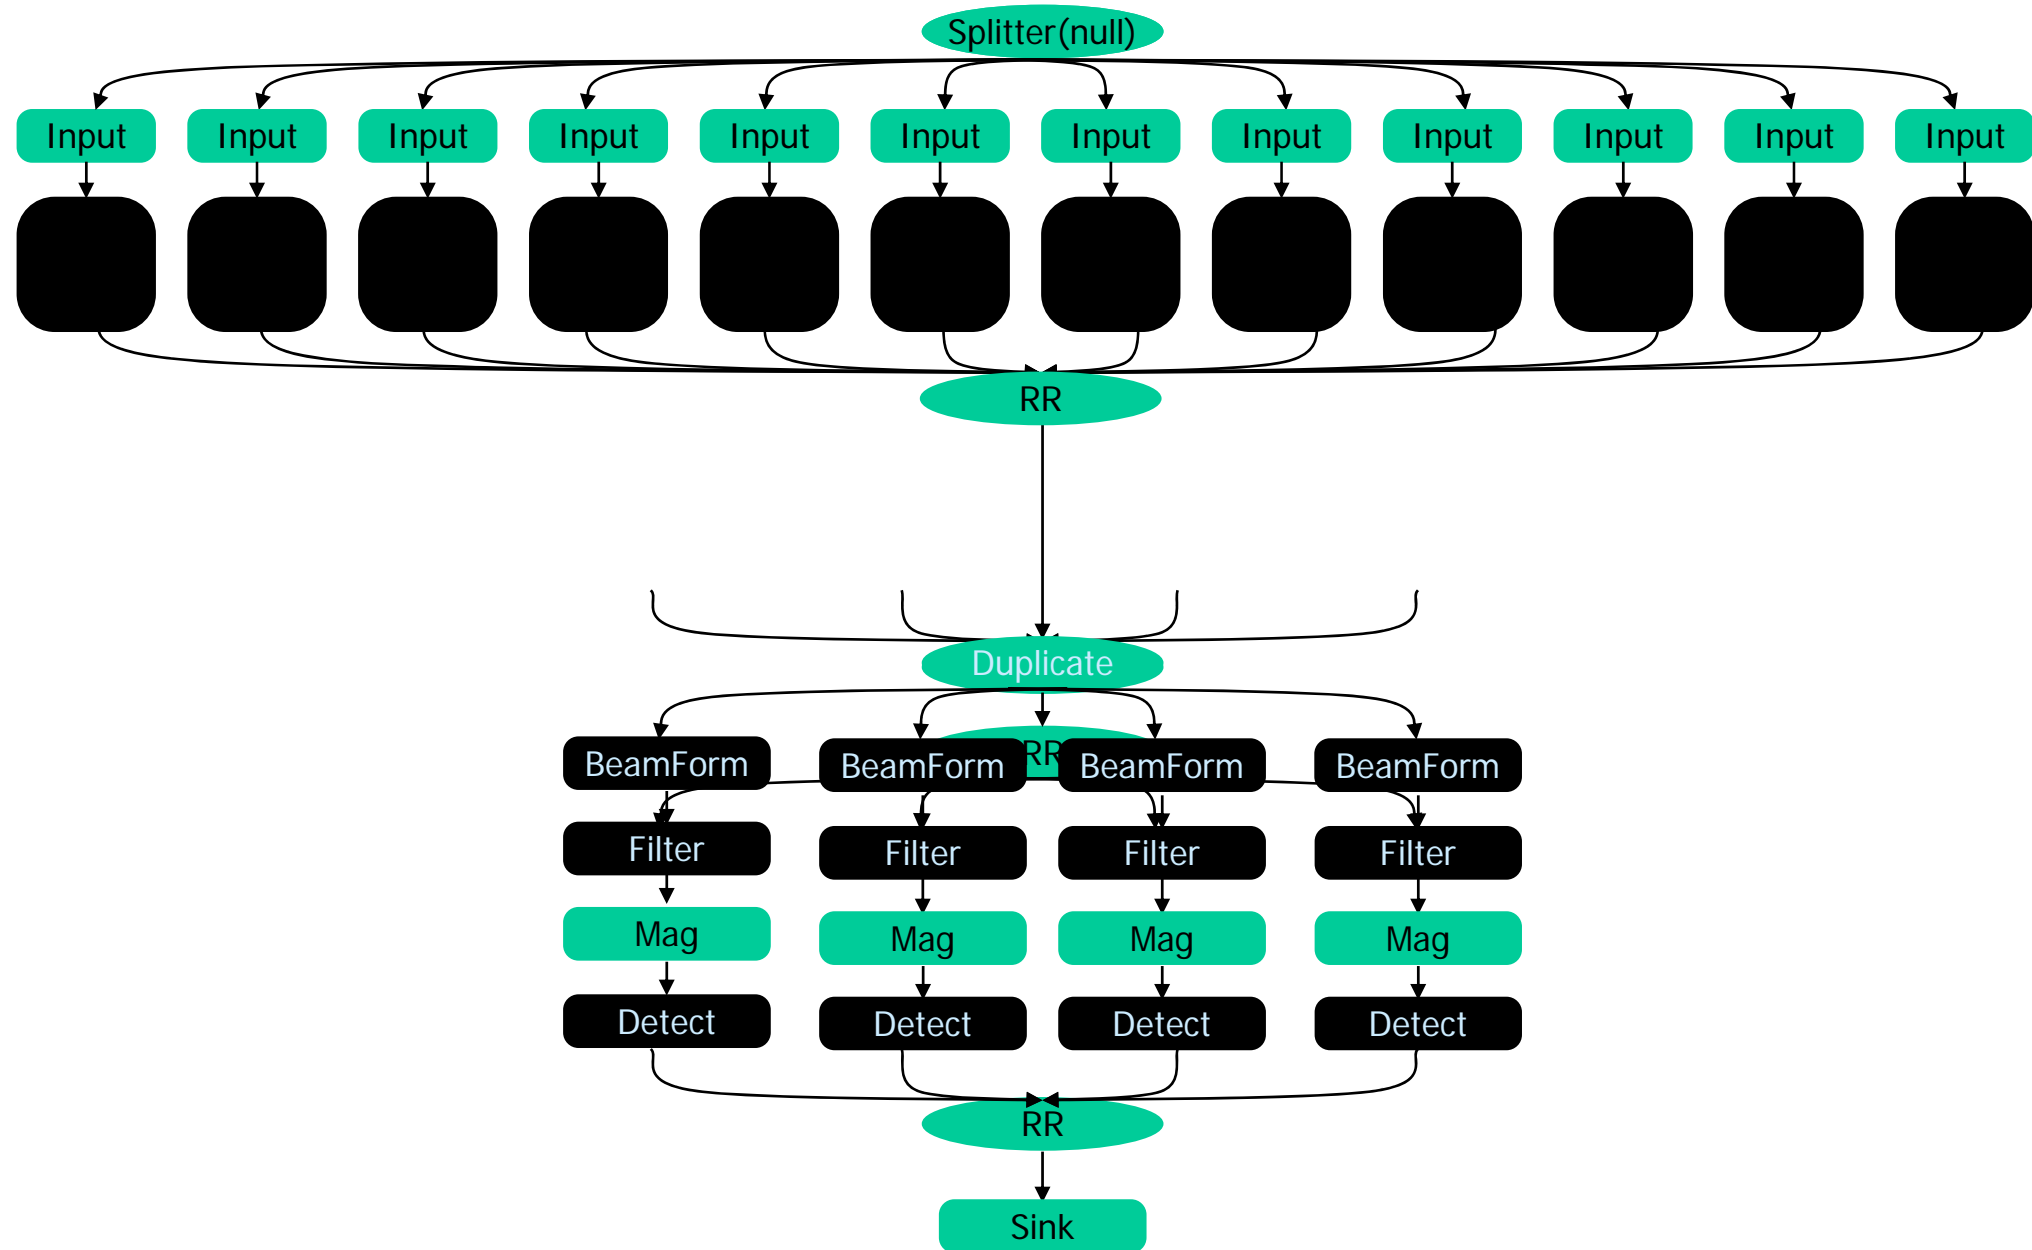




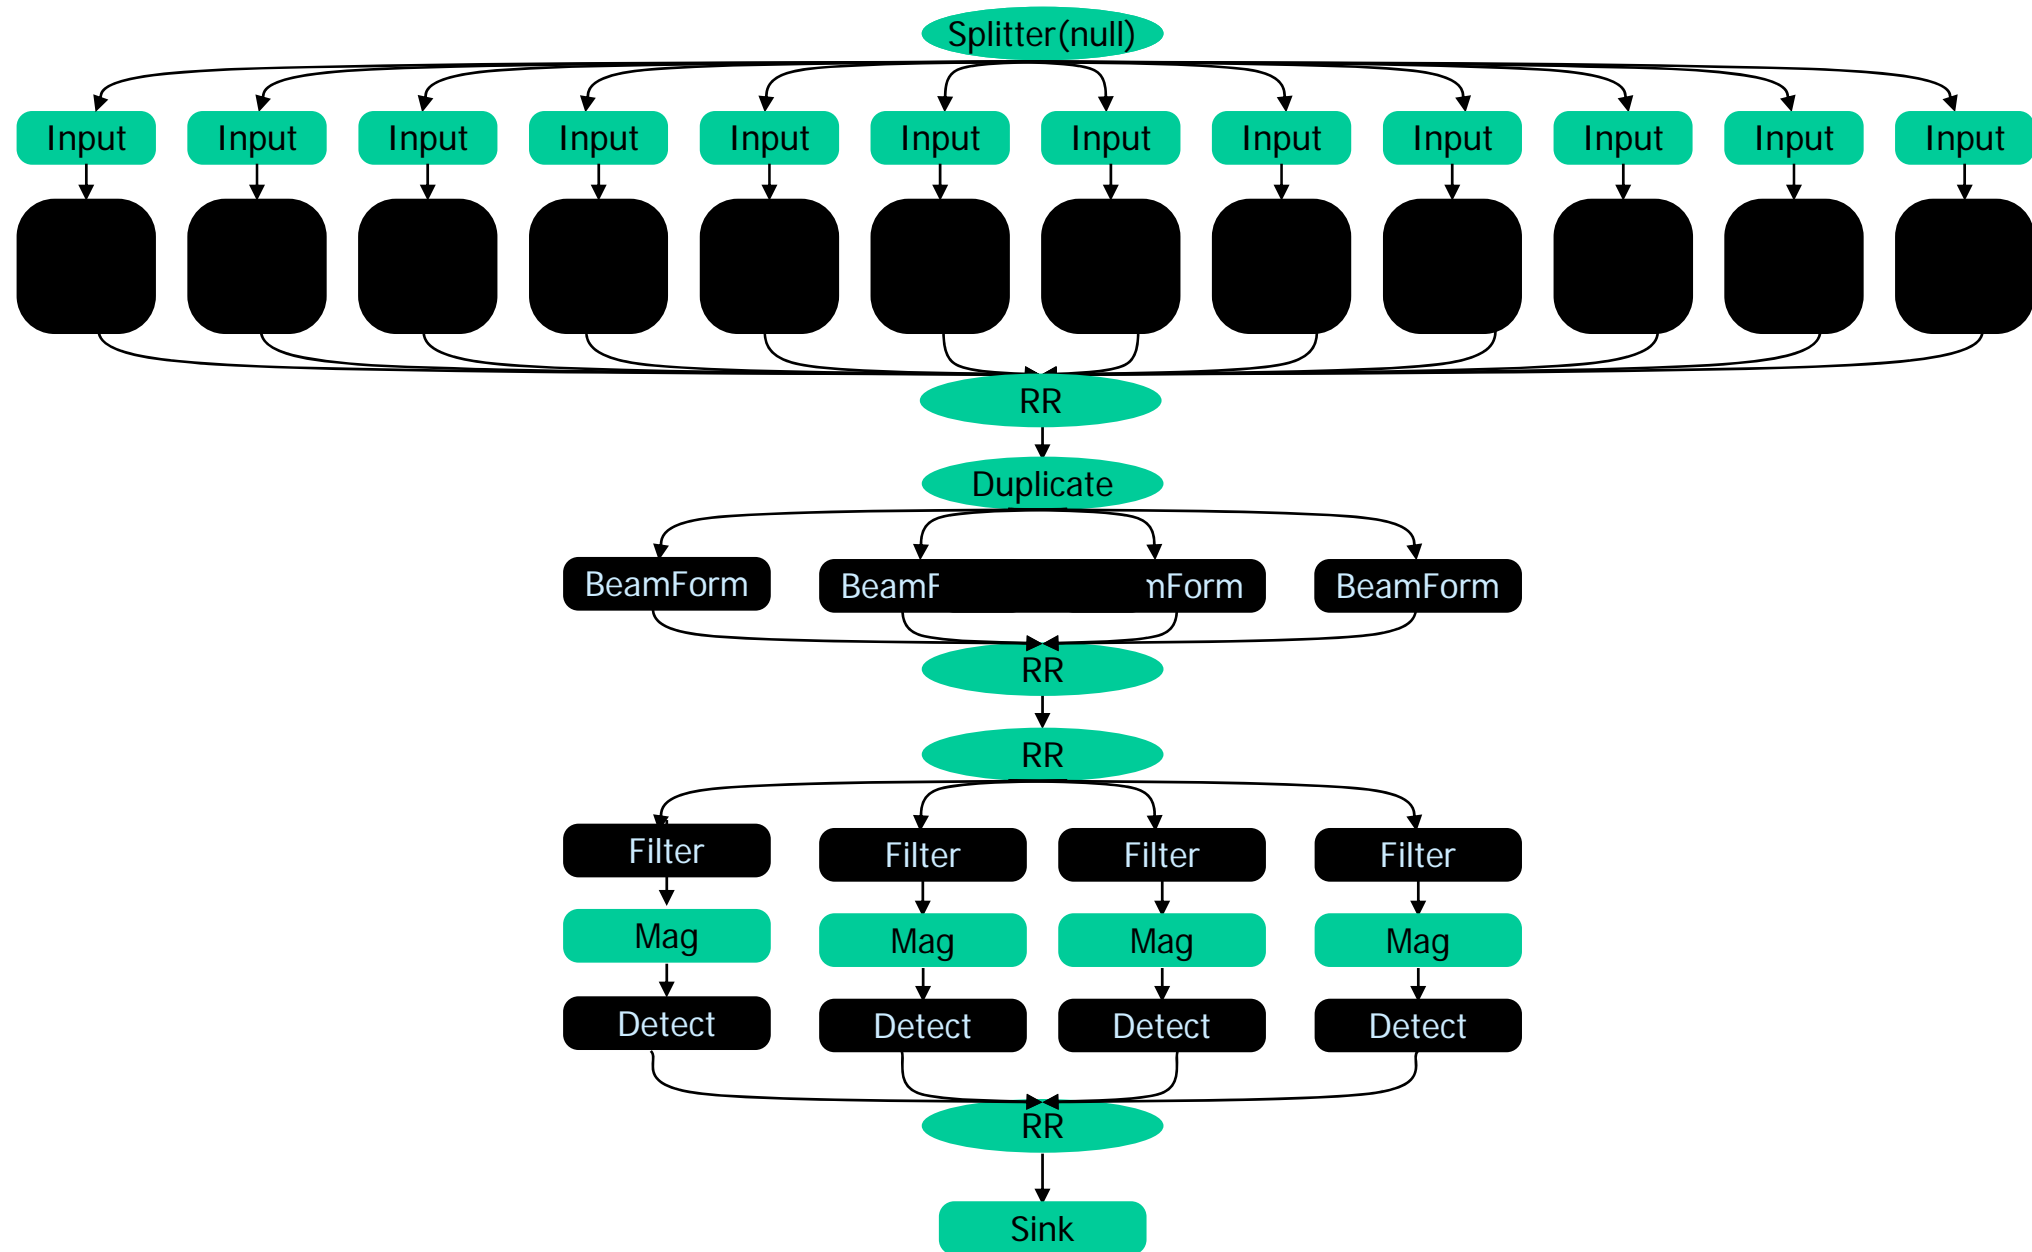
# Radar (Transformation Selection)




# Radar (Transformation Selection)



# Radar (Transformation Selection)




# Radar



Maximal Combination and  
Shifting to Frequency Domain



2.4 times as  
many FLOPS




Using Transformation  
Selection




half as many  
FLOPS

# Floating Point Operations Reduction




# Execution Speedup



On a Pentium IV


# Execution Speedup



On a Pentium IV

# StreamIt: Lessons Learned

- In practice, I/O rates of filters are often matched [\[LCTES'03\]](#)
  - Over 30 publications study an uncommon case (CD-DAT)



- Multi-phase filters complicate programs, compilers
  - Should maintain simplicity of only one atomic step per filter
- Programmers accidentally introduce mutable filter state

```
void>int filter SquareWave() {  
    work push 2 {  
        push(0);  
        push(1);  
    }  
}
```

*stateless*

```
void>int filter SquareWave() {  
    int x = 0;  
  
    work push 1 {  
        push(x);  
        x = 1 - x;  
    }  
}
```


*stateful*



# Future of StreamIt

- **Goal:** influence the next big language

## Origins of C++




# Research Trajectory

- **Vision: Make emerging computational substrates universally accessible and useful**


## 1. Languages, compilers, & tools for multicores

- I believe new language / compiler technology can enable scalable and robust performance
- Next inroads: expose & exploit flexibility in programs



## 2. Programmable microfluidics

- We have developed programming languages, tools, and flexible new devices for microfluidics
- Potential to revolutionize biology experimentation



## 3. Technologies for the developing world

- TEK: enable Internet experience over email account
- Audio Wiki: publish content from a low-cost phone
- uBox / uPhone: monitor & improve rural healthcare



# Conclusions

- **A parallel programming model will succeed only by luring programmers, making them do less, not more**
- **Stream programming lures programmers with:**
  - Elegant programming primitives
  - Domain-specific optimizations
- **Meanwhile, streaming is implicitly parallel**
  - Robust performance via task, data, & pipeline parallelism
- **We believe stream programming will play a key role in enabling a transition to multicore processors**

## *Contributions*

- Structured streams
- Teleport messaging
- Unified algorithm for task, data, pipeline parallelism
- Software pipelining of whole procedures
- Algebraic simplification of whole procedures
- Translation from time to frequency
- Selection of best DSP transforms