

StreamIt: A Language for Streaming Applications

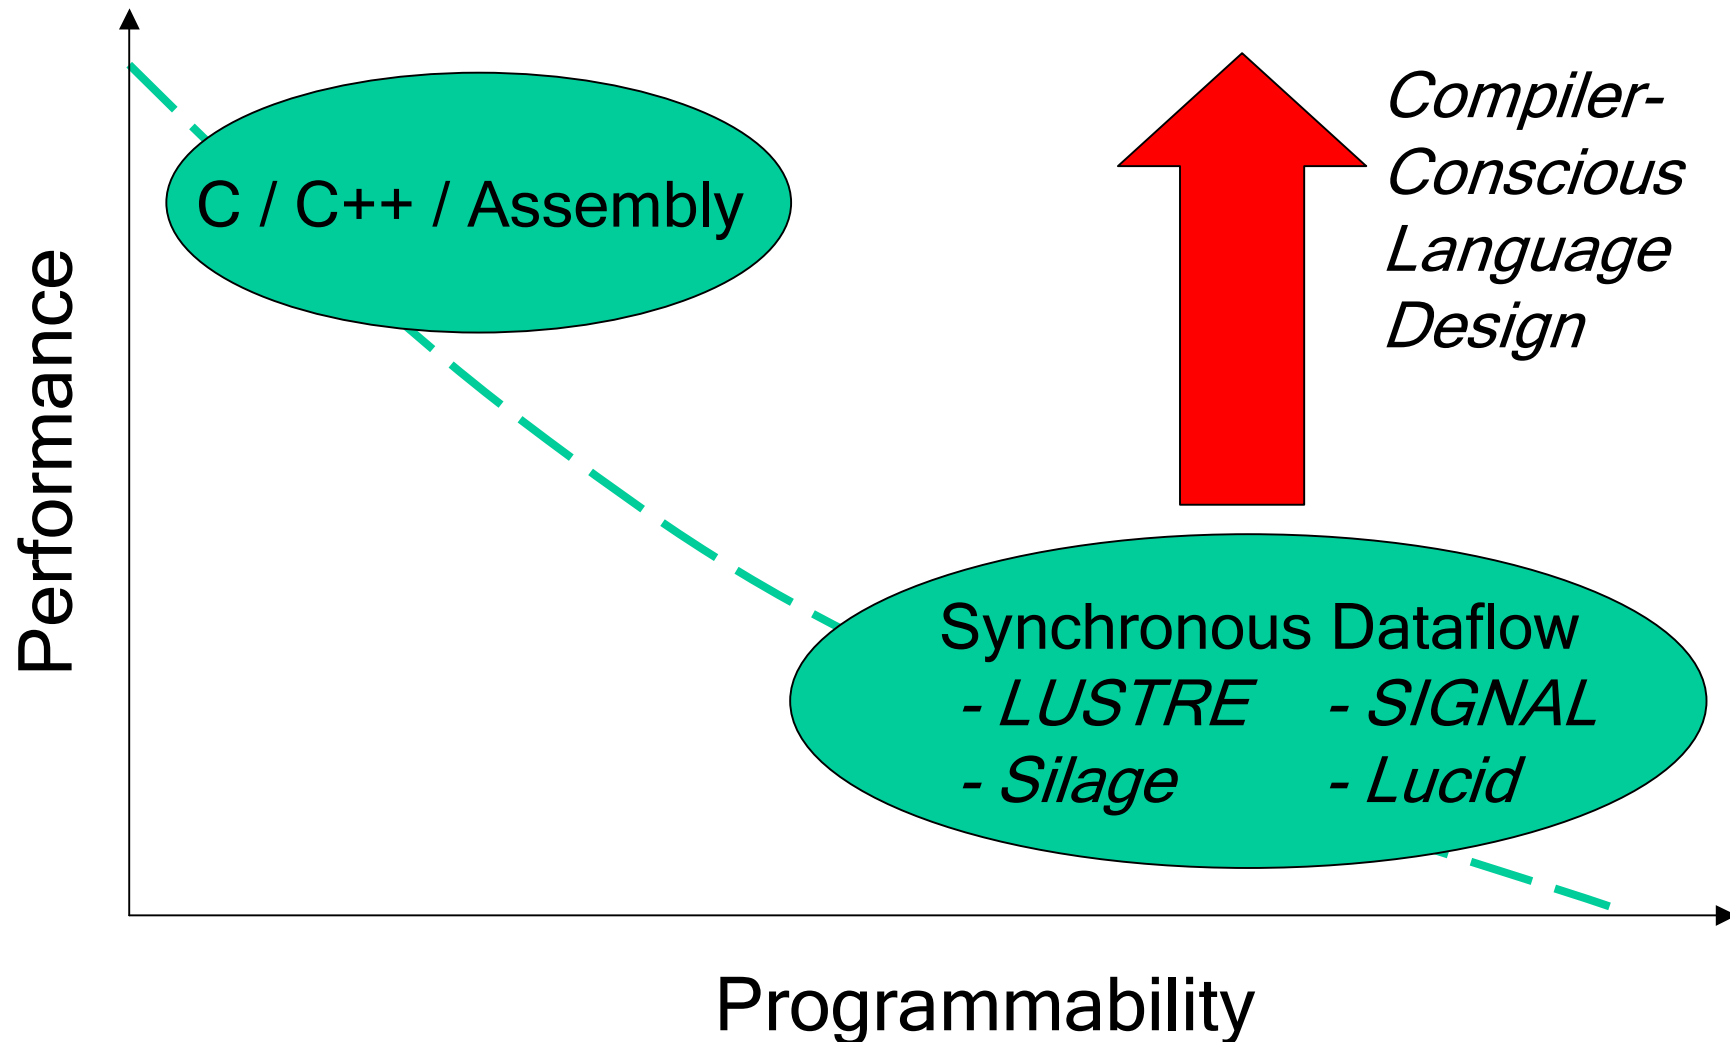
William Thies, Michal Karczmarek,
Michael Gordon, and Saman Amarasinghe

MIT Laboratory for Computer Science

Streaming Application Domain

- Based on streams of data
- Increasingly prevalent and important
 - Embedded systems
 - Cell phones, handheld computers, DSP's
 - Desktop applications
 - Streaming media
 - Software radio
 - Real-time encryption
 - Graphics packages
 - High-performance servers
 - Software routers
 - Cell phone base stations
 - HDTV editing consoles

Developing Stream Programs



The StreamIt Language

- Also a synchronous dataflow language
 - With a few extra features
 - Goals:
 - High performance
 - Improved programmer productivity
 - Language Contributions:
 - Structured model of streams
 - Messaging system for control
 - Automatic program morphing
- } **ENABLES
Compiler
Analysis &
Optimization**

Outline


- Design of StreamIt
 - Structured Streams
 - Messaging
 - Morphing
- Results
- Conclusions

Outline

- Design of StreamIt
 - Structured Streams
 - Messaging
 - Morphing
- Results
- Conclusions


Representing Streams

- Conventional wisdom: streams are graphs
 - Graphs have no simple textual representation
 - Graphs are difficult to analyze and optimize




Representing Streams

- Conventional wisdom: streams are graphs
 - Graphs have no simple textual representation
 - Graphs are difficult to analyze and optimize
- Insight: stream programs have structure



unstructured



structured


Structured Streams

- Hierarchical structures:


- Pipeline



- SplitJoin



- Feedback Loop



- Basic programmable unit: Filter




Structured Streams

- Hierarchical structures:


- Pipeline



- SplitJoin



- Feedback Loop




- Basic programmable unit: Filter



- Splits / Joins are compiler-defined




Representing Filters

- Autonomous unit of computation 
 - No access to global resources
 - Communicates through FIFO channels
 - `input.pop()` - `input.peek(index)` - `output.push(value)`
 - Peek / pop / push rates must be constant
- Looks like a Java class, with
 - An initialization function
 - A steady-state “work” function
 - Message handler functions
- Implementation has nothing to do with Java


Filter Example: LowPassFilter

```
class LowPassFilter extends Filter {  
    float[] weights;  
  
    void init(int N) {  
        weights = calcWeights(N);  
        setPush(1); setPop(1); setPeek(N);  
        setInput(Float.TYPE); setOutput(Float.TYPE);  
    }  
  
    void work() {  
        float result = 0;  
        for (int i=0; i<weights.length; i++) {  
            result += weights[i] * input.peek(i);  
        }  
        output.push(result);  
        input.pop();  
    }  
}
```




Filter Example: LowPassFilter

```
class LowPassFilter extends Filter {  
    float[] weights;  
  
    void init(int N) {  
        weights = calcWeights(N);  
        setPush(1); setPop(1); setPeek(N);  
        setInput(Float.TYPE); setOutput(Float.TYPE);  
    }  
  
    void work() {  
        float result = 0;  
        for (int i=0; i<weights.length; i++) {  
            result += weights[i] * input.peek(i);  
        }  
        output.push(result);  
        input.pop();  
    }  
}
```




Filter Example: LowPassFilter

```
class LowPassFilter extends Filter {  
    float[] weights;  
  
    void init(int N) {  
        weights = calcWeights(N);  
        setPush(1); setPop(1); setPeek(N);  
        setInput(Float.TYPE); setOutput(Float.TYPE);  
    }  
  
    void work() {  
        float result = 0;  
        for (int i=0; i<weights.length; i++) {  
            result += weights[i] * input.peek(i);  
        }  
        output.push(result);  
        input.pop();  
    }  
}
```




Filter Example: LowPassFilter

```
class LowPassFilter extends Filter {  
    float[] weights;  
  
    void init(int N) {  
        weights = calcWeights(N);  
        setPush(1); setPop(1); setPeek(N);  
        setInput(Float.TYPE); setOutput(Float.TYPE);  
    }  
  
    void work() {  
        float result = 0;  
        for (int i=0; i<weights.length; i++) {  
            result += weights[i] * input.peek(i);  
        }  
        output.push(result);  
        input.pop();  
    }  
}
```




Filter Example: LowPassFilter

```
class LowPassFilter extends Filter {  
    float[] weights;  
  
    void init(int N) {  
        weights = calcWeights(N);  
        setPush(1); setPop(1); setPeek(N);  
        setInput(Float.TYPE); setOutput(Float.TYPE);  
    }  
  
    void work() {  
        float result = 0;  
        for (int i=0; i<weights.length; i++) {  
            result += weights[i] * input.peek(i);  
        }  
        output.push(result);  
        input.pop();  
    }  
}
```



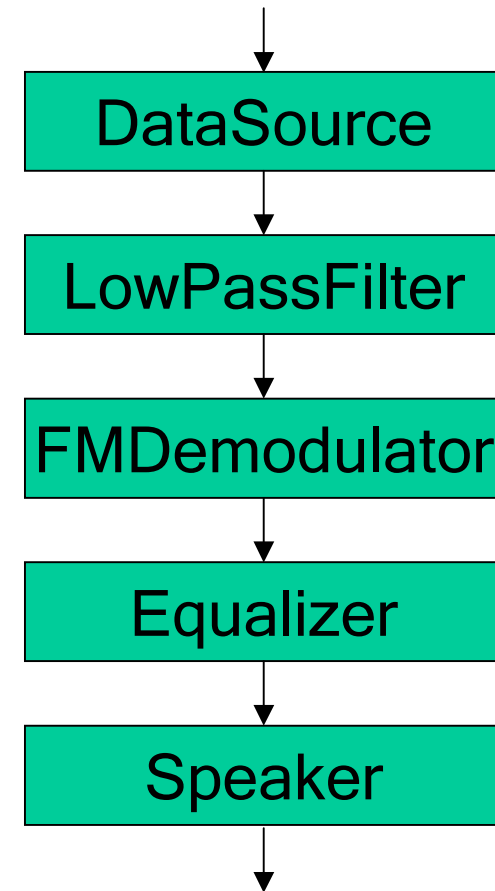
Pipeline Example: FM Radio

```
class FMRadio extends Pipeline {  
    void init() {  
        add(new DataSource());  
        add(new LowPassFilter());  
        add(new FMDemodulator());  
        add(new Equalizer(8));  
        add(new Speaker());  
    }  
}
```




Pipeline Example: FM Radio

```
class FMRadio extends Pipeline {  
    void init() {  
        add(new DataSource());  
        add(new LowPassFilter());  
        add(new FMDemodulator());  
        add(new Equalizer(8));  
        add(new Speaker());  
    }  
}
```




SplitJoin Example: Equalizer

```
class Equalizer extends Pipeline {  
    void init(int N) {  
        add(new SplitJoin() {  
            void init() {  
                setSplitter(Duplicate());  
                float freq = 10000;  
                for (int i = 0; i < N; i ++, freq*=2) {  
                    add(new BandPassFilter(freq, 2*freq);  
                }  
                setJoiner(RoundRobin(1));  
            });  
        add(new Adder(N));  
    }  
}
```




Why Structured Streams?

- Compare to structured control flow



GOTO statements




If / else / for statements

- Tradeoff:
 - PRO: - more robust - more analyzable
 - CON: - “restricted” style of programming

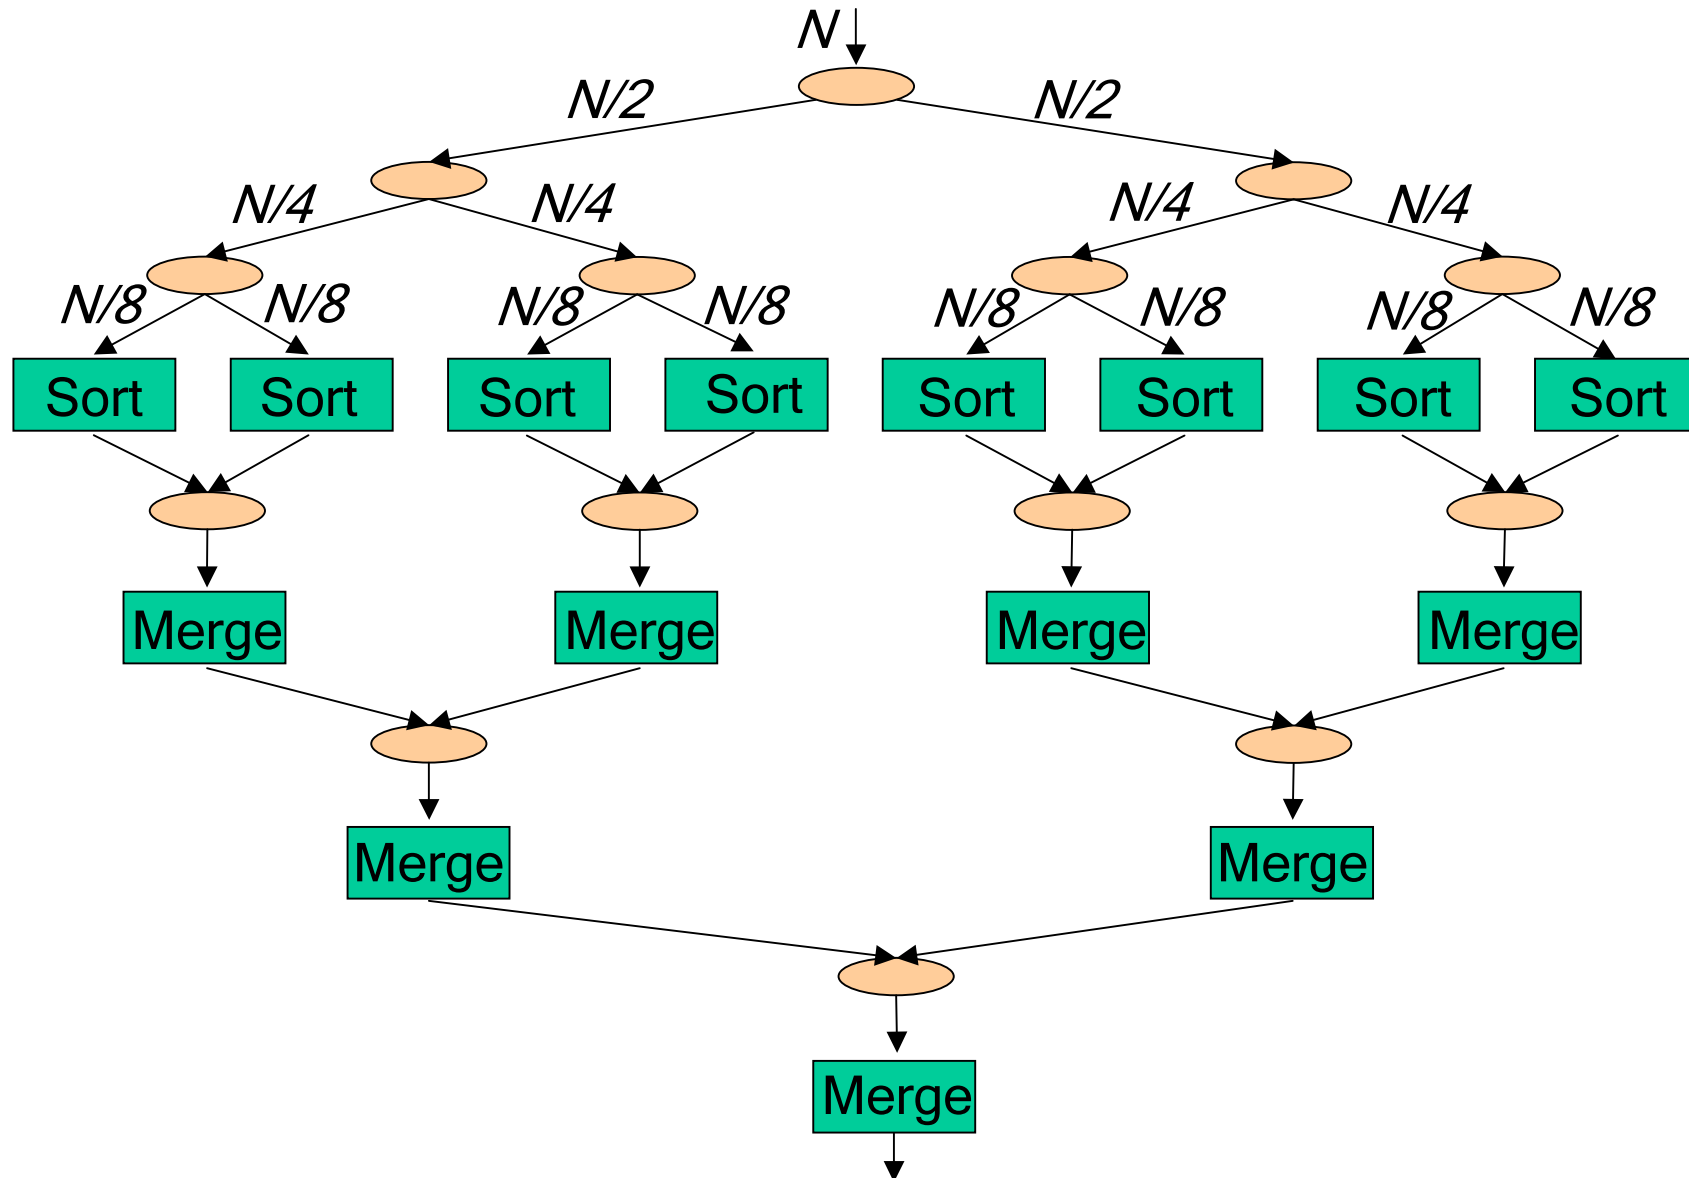
Structure Helps Programmers

- Modules are hierarchical and composable
 - Each structure is single-input, single-output



- Encapsulates common idioms
- Good textual representation
 - Enables parameterizable graphs

N-Element Merge Sort (3-level)




N-Element Merge Sort (K-level)

```
class MergeSort extends Pipeline {
    void init(int N, int K) {
        if (K==1) {
            add(new Sort(N));
        } else {
            add(new SplitJoin() {
                void init() {
                    setSplitter(RoundRobin());
                    add(new MergeSort(N/2, K-1));
                    add(new MergeSort(N/2, K-1));
                    setJoiner(RoundRobin());
                }
            });
            add(new Merge(N));
        }
    }
}
```

Structure Helps Compilers


- Enables local, hierarchical analyses
 - Scheduling
 - Optimization
 - Parallelization
 - Load balancing





Structure Helps Compilers

- Enables local, hierarchical analyses

- Scheduling
- Optimization
- Parallelization
- Load balancing




- Examples:




Structure Helps Compilers

- Enables local, hierarchical analyses

- Scheduling
- Optimization
- Parallelization
- Load balancing




- Examples:



Structure Helps Compilers

- Enables local, hierarchical analyses
 - Scheduling
 - Optimization
 - Parallelization
 - Load balancing



- Disallows non-sensical graphs
- Simplifies separate compilation
 - All blocks single-input, single-output

CON: Restricts Coding Style



- Some graphs need to be re-arranged
- Example: FFT

Bit-reverse order

Butterfly (2 way)

Butterfly (4 way)

Butterfly (8 way)




Outline

- Design of StreamIt
 - Structured Streams
 - Messaging
 - Morphing
- Results
- Conclusions

Control Messages

- Structures for regular, high-bandwidth data
- But also need a control mechanism for irregular, low-bandwidth events




- Change volume on a cell phone
- Initiate handoff of stream
- Adjust network protocol

Supporting Control Messages

- Option 1: Embed message in stream

PRO: - message arrives with data


CON: - complicates filter code
- complicates structure
- runtime overhead



- Option 2: Synchronous method call

PRO: - delivery transparent to user


CON: - timing is unclear
- limits parallelism



StreamIt Messaging System

- Looks like method call, but semantics differ

```
void raiseVolume(int v)  
myVolume += v;  
}
```




- No return value
- Asynchronous delivery
- Can broadcast to multiple targets

StreamIt Messaging System

- Looks like method call, but semantics differ


```
void work () {  
    TargetFilter x;  
  
    ...  
    if (lowVolume())  
        x.raiseVolume(10);  
}
```



- No return value
 - Asynchronous delivery
 - Can broadcast to multiple targets
- Timed relative to data
 - User gains precision; compiler gains flexibility


Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




Message Timing

- A sends message to B with zero latency




General Message Timing

- Latency of N means:
 - Message attached to wavefront that *sender* sees in N executions




General Message Timing

- Latency of N means:
 - Message attached to wavefront that *sender* sees in N executions
- Examples:
 - $A \rightarrow B$, latency 1




General Message Timing

- Latency of N means:
 - Message attached to wavefront that *sender* sees in N executions
- Examples:
 - $A \rightarrow B$, latency 1




General Message Timing

- Latency of N means:
 - Message attached to wavefront that *sender* sees in N executions
- Examples:
 - $A \rightarrow B$, latency 1
 - $B \rightarrow A$, latency 25



General Message Timing

- Latency of N means:
 - Message attached to wavefront that *sender* sees in N executions
- Examples:
 - $A \rightarrow B$, latency 1
 - $B \rightarrow A$, latency 25



Rationale


- Better for the programmer
 - Simplicity of method call
 - Precision of embedding in stream
- Better for the compiler
 - Program is easier to analyze
 - No code for timing / embedding
 - No control channels in stream graph
 - Can reorder filter firings, respecting constraints
 - Implement in most efficient way

Outline

- Design of StreamIt
 - Structured Streams
 - Messaging
 - Morphing
- Results
- Conclusions

Dynamic Changes to Stream

- Stream structure needs to change




- Examples

- Switch radio from AM to FM
- Change from Bluetooth to 802.11

} Program
“Morphing”

Dynamic Changes to Stream

- Stream structure needs to change




- Challenges for programmer:
 - Synchronizing the beginning, end of morphing
 - Preserving live data in the system
 - Efficiency

Morphing in StreamIt

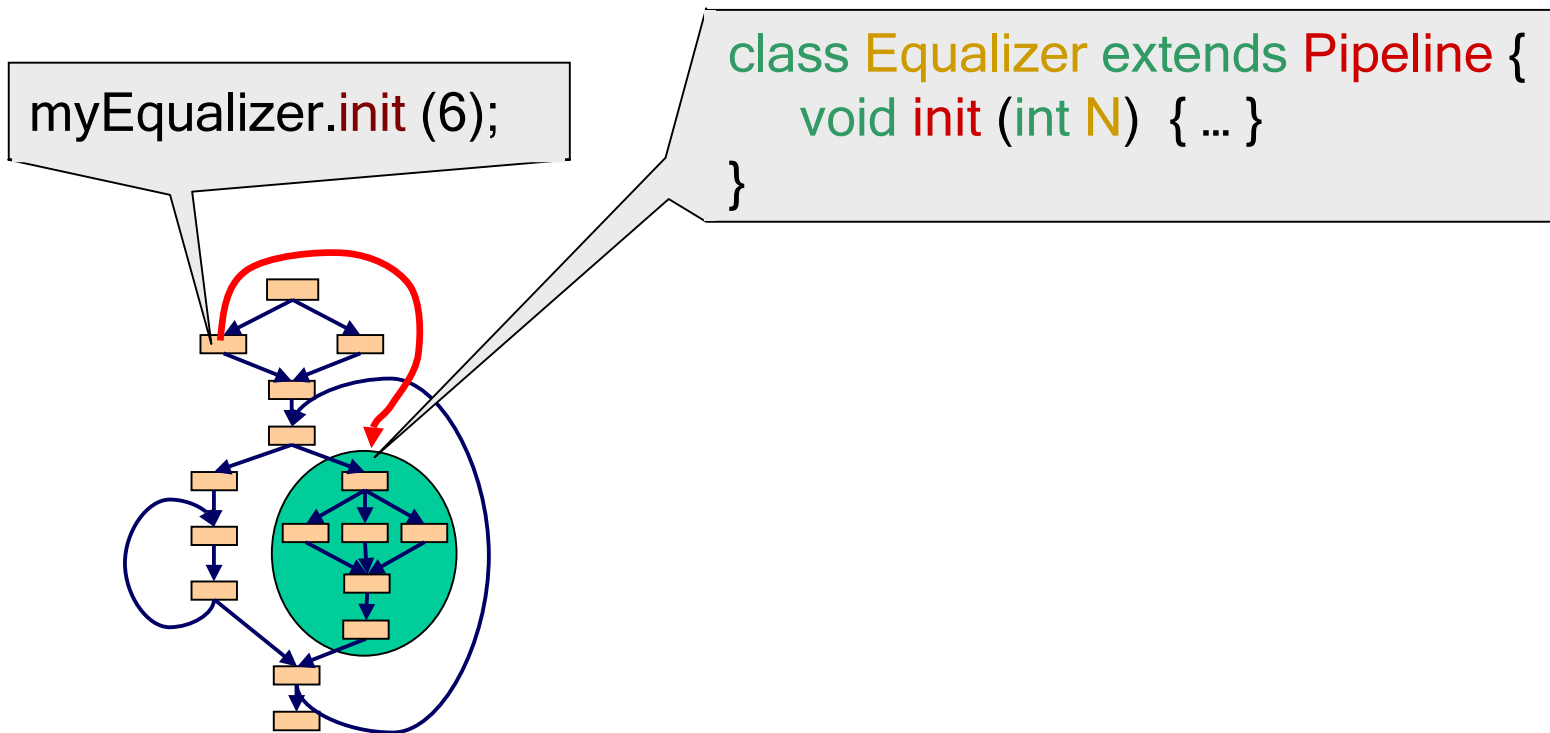
- Send message to “init” to morph a structure

```
class Equalizer extends Pipeline {  
    void init (int N) { ... }  
}
```



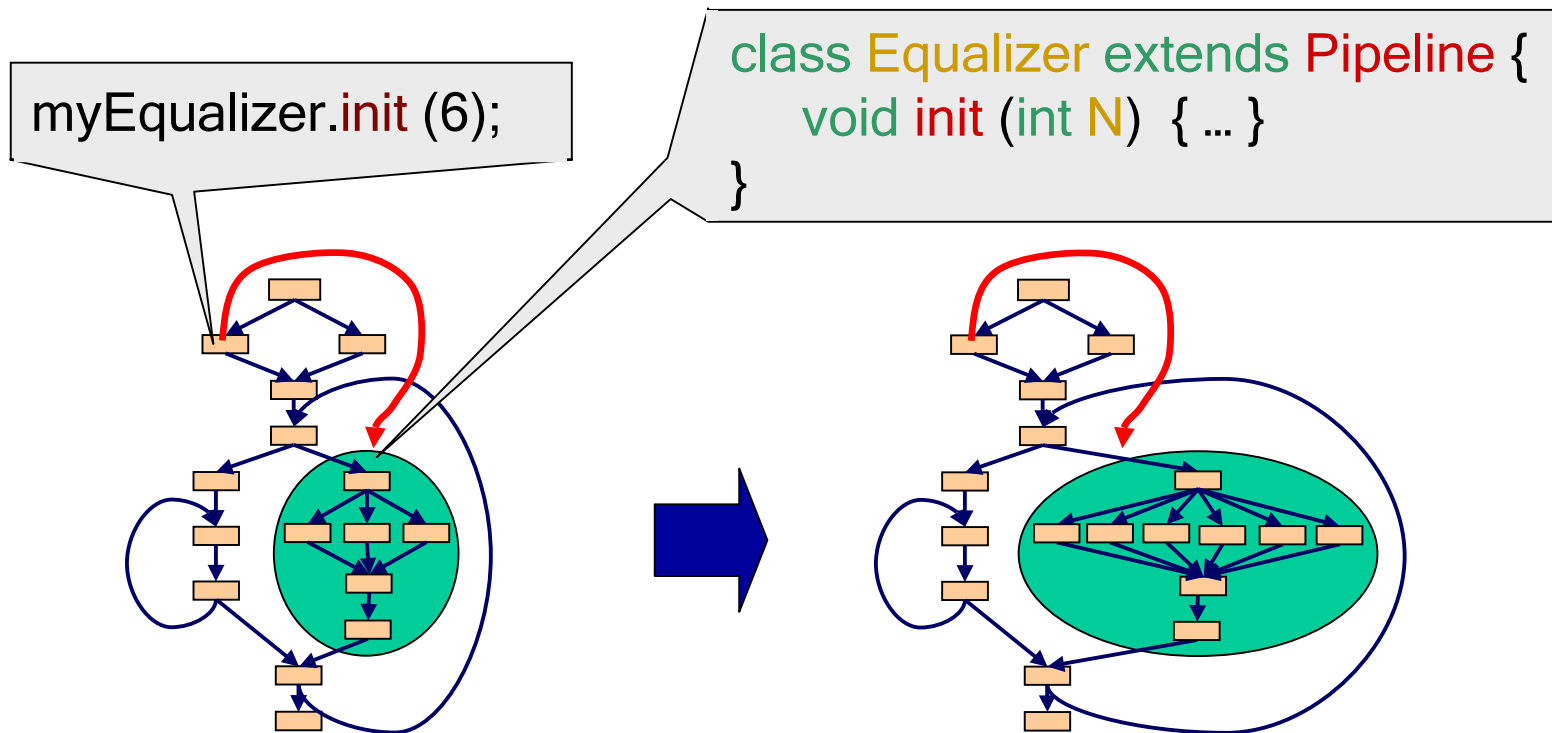
Morphing in StreamIt

- Send message to “init” to morph a structure



Morphing in StreamIt

- Send message to “init” to morph a structure



- When message arrives, structure is replaced
- Live data is automatically drained

Rationale


- Programmer writes “init” only once
 - No need for complicated transitions
- Compiler optimizes each phase separately
 - Benefits from anticipation of phase changes

Outline

- Design of StreamIt
 - Structured Streams
 - Messaging
 - Morphing
- Results
- Conclusions


Implementation

- Prototype StreamIt compiler complete
- Backends:
 - Uniprocessor
 - RAW: A tiled architecture with fine-grained, programmable communication
- Extended KOPI, open-source Java compiler

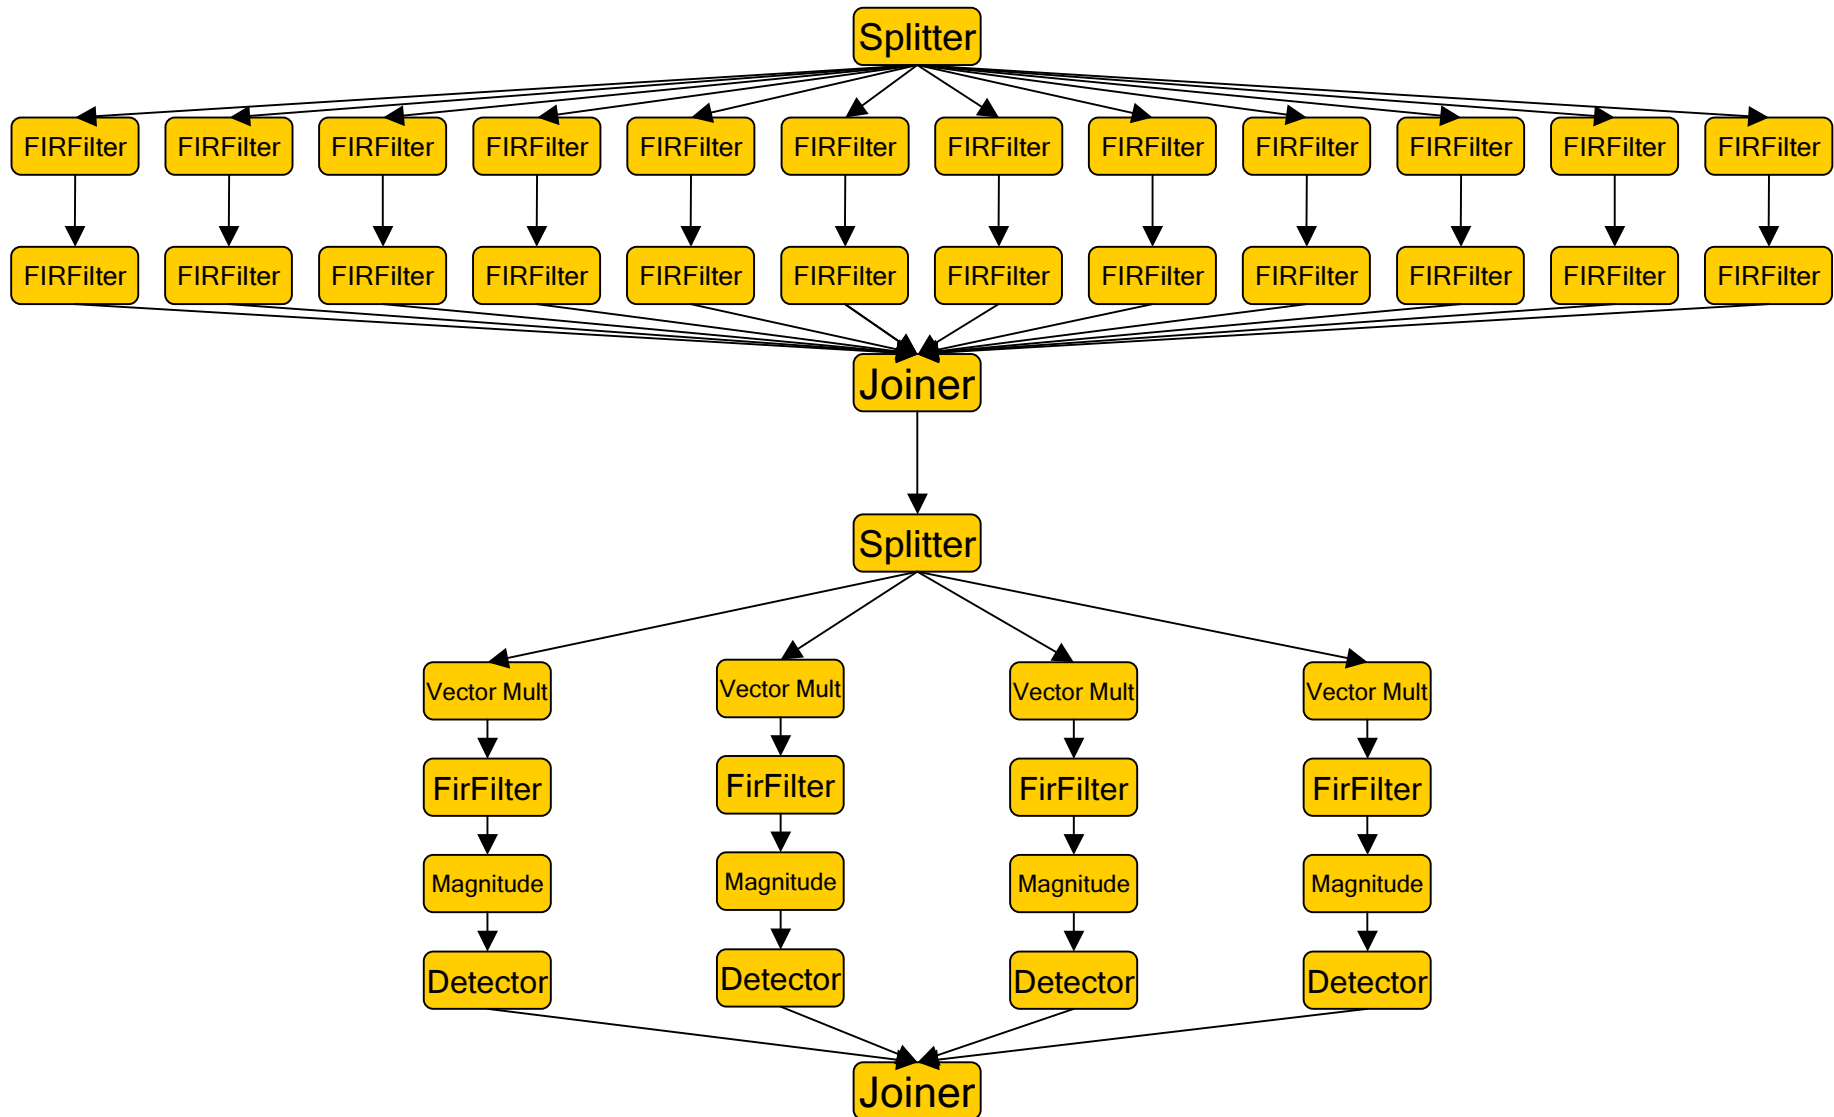


Results


- Developed applications in StreamIt
 - GSM Decoder
 - FM Radio
 - BeamFormer
 - FFT
 - Matrix multiply
 - CRC Encoder/Decoder
- Load-balancing transformations improve performance on RAW



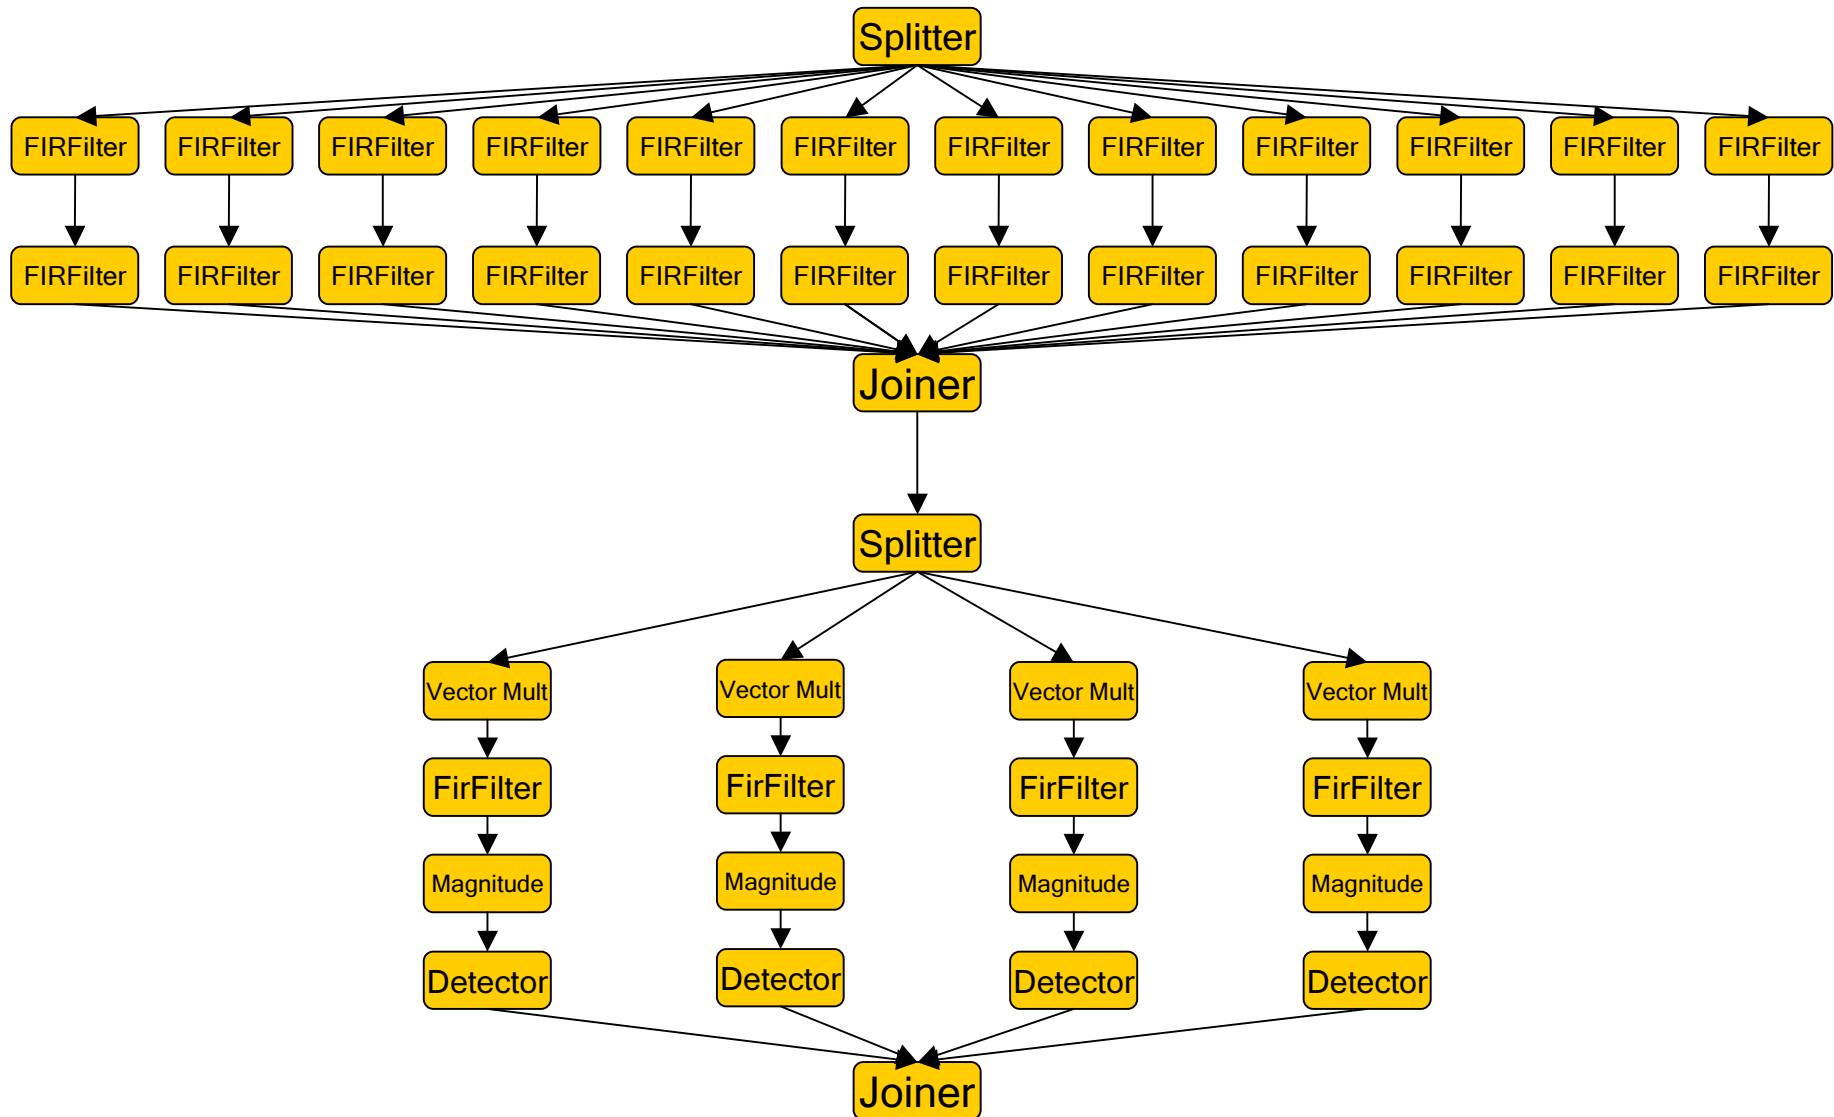
Example: BeamFormer (Original)



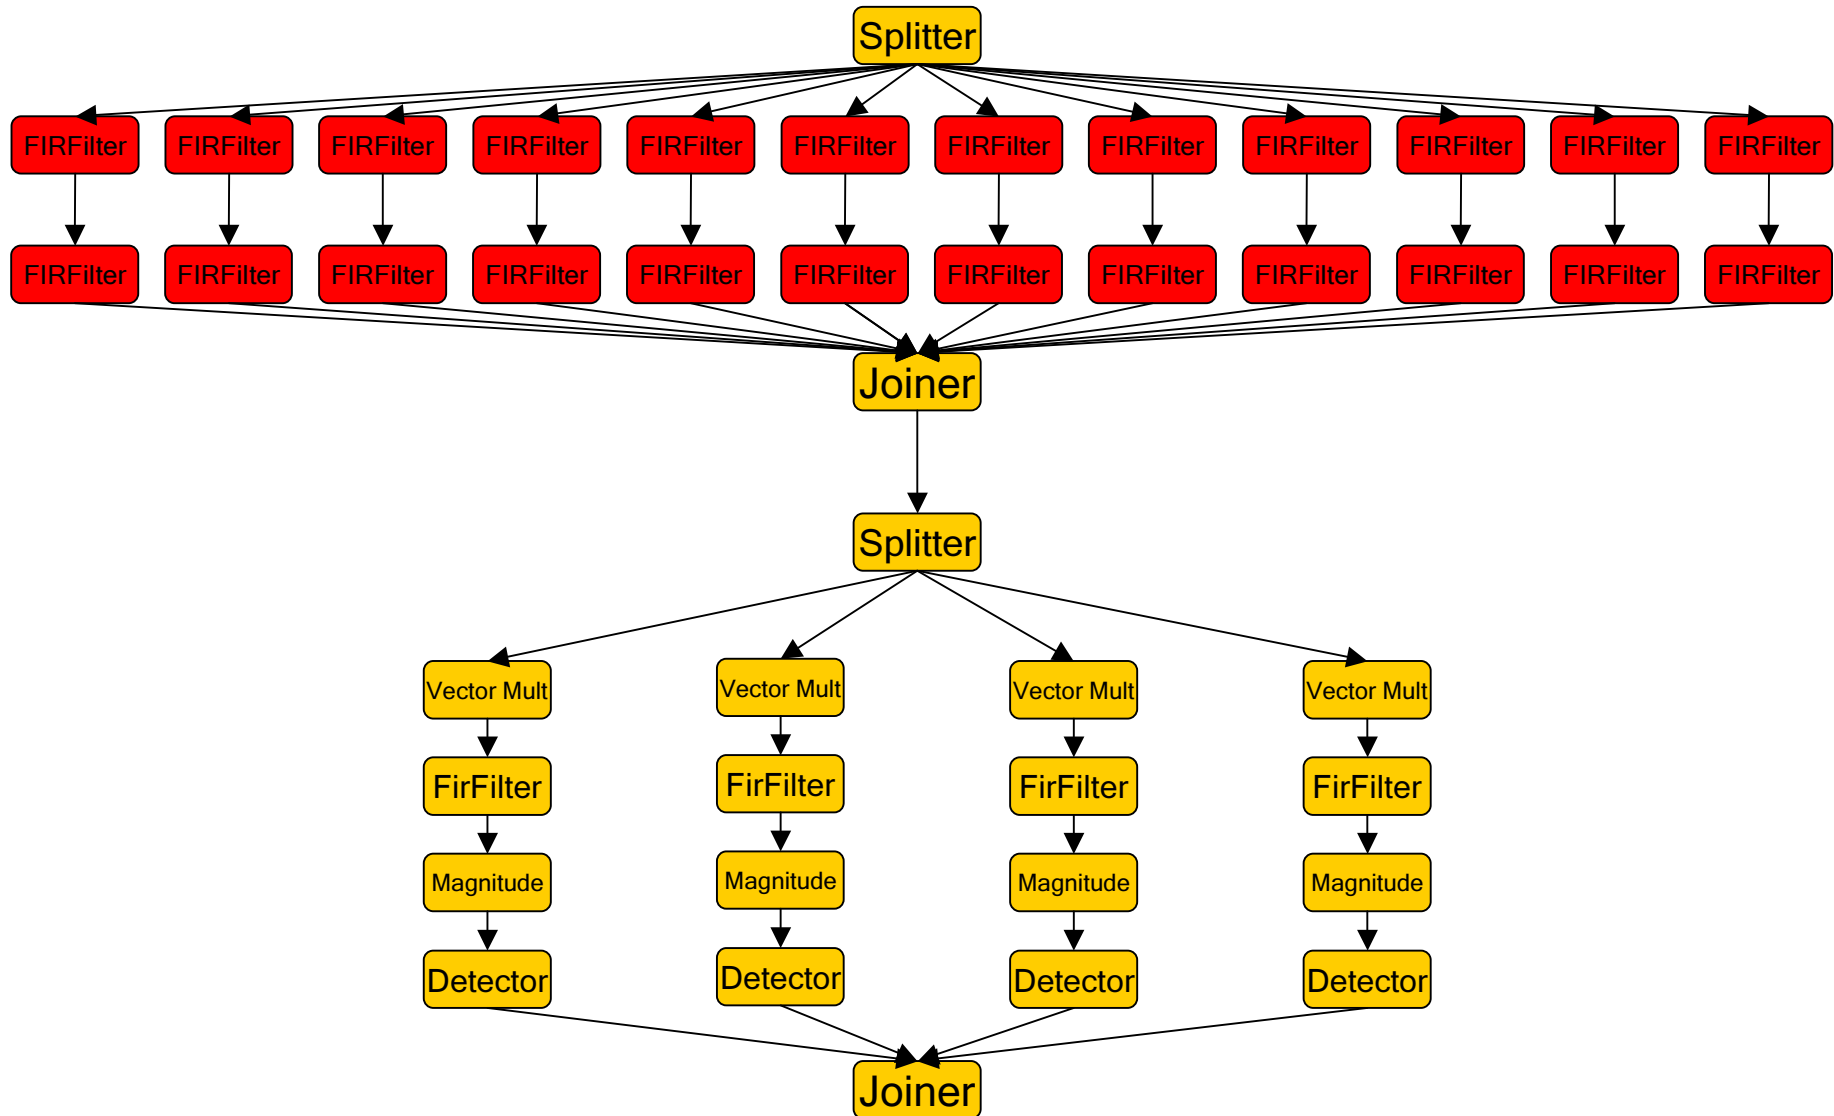
Example: BeamFormer (Original)



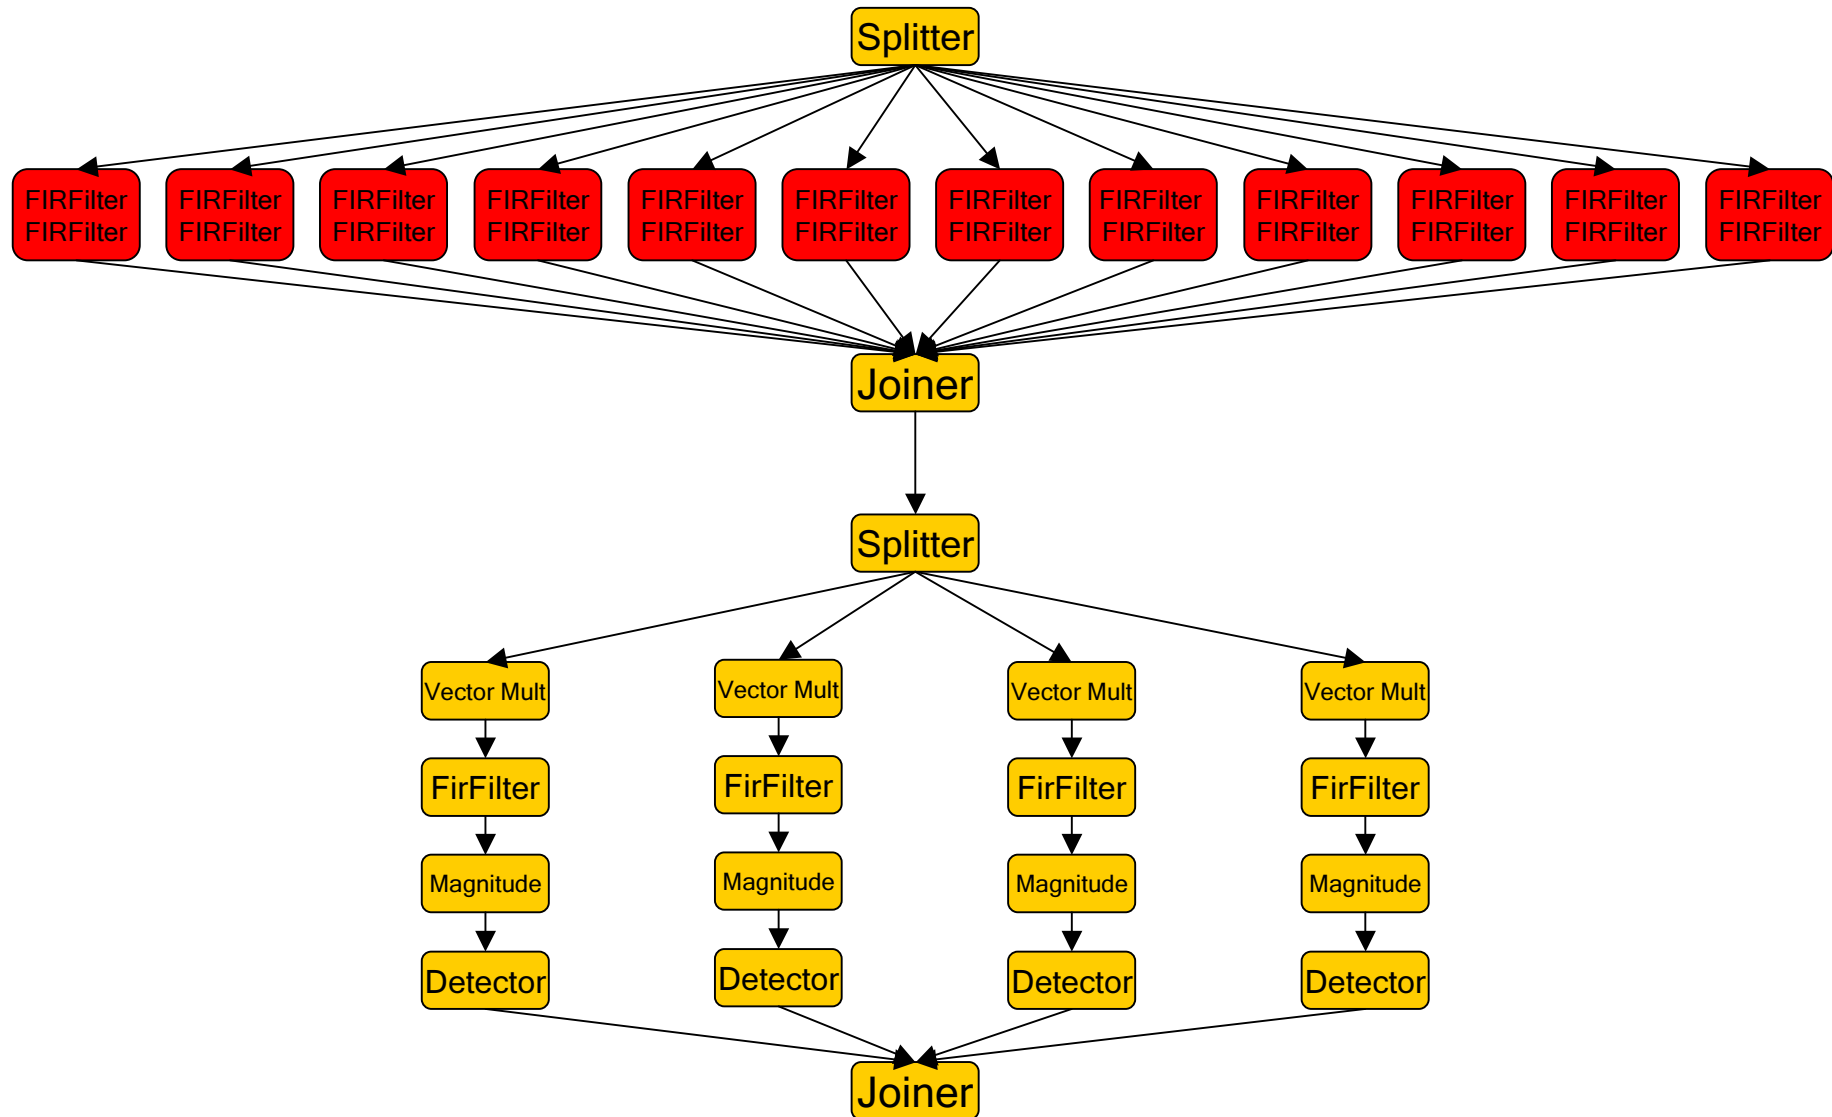
Example: BeamFormer (Original)



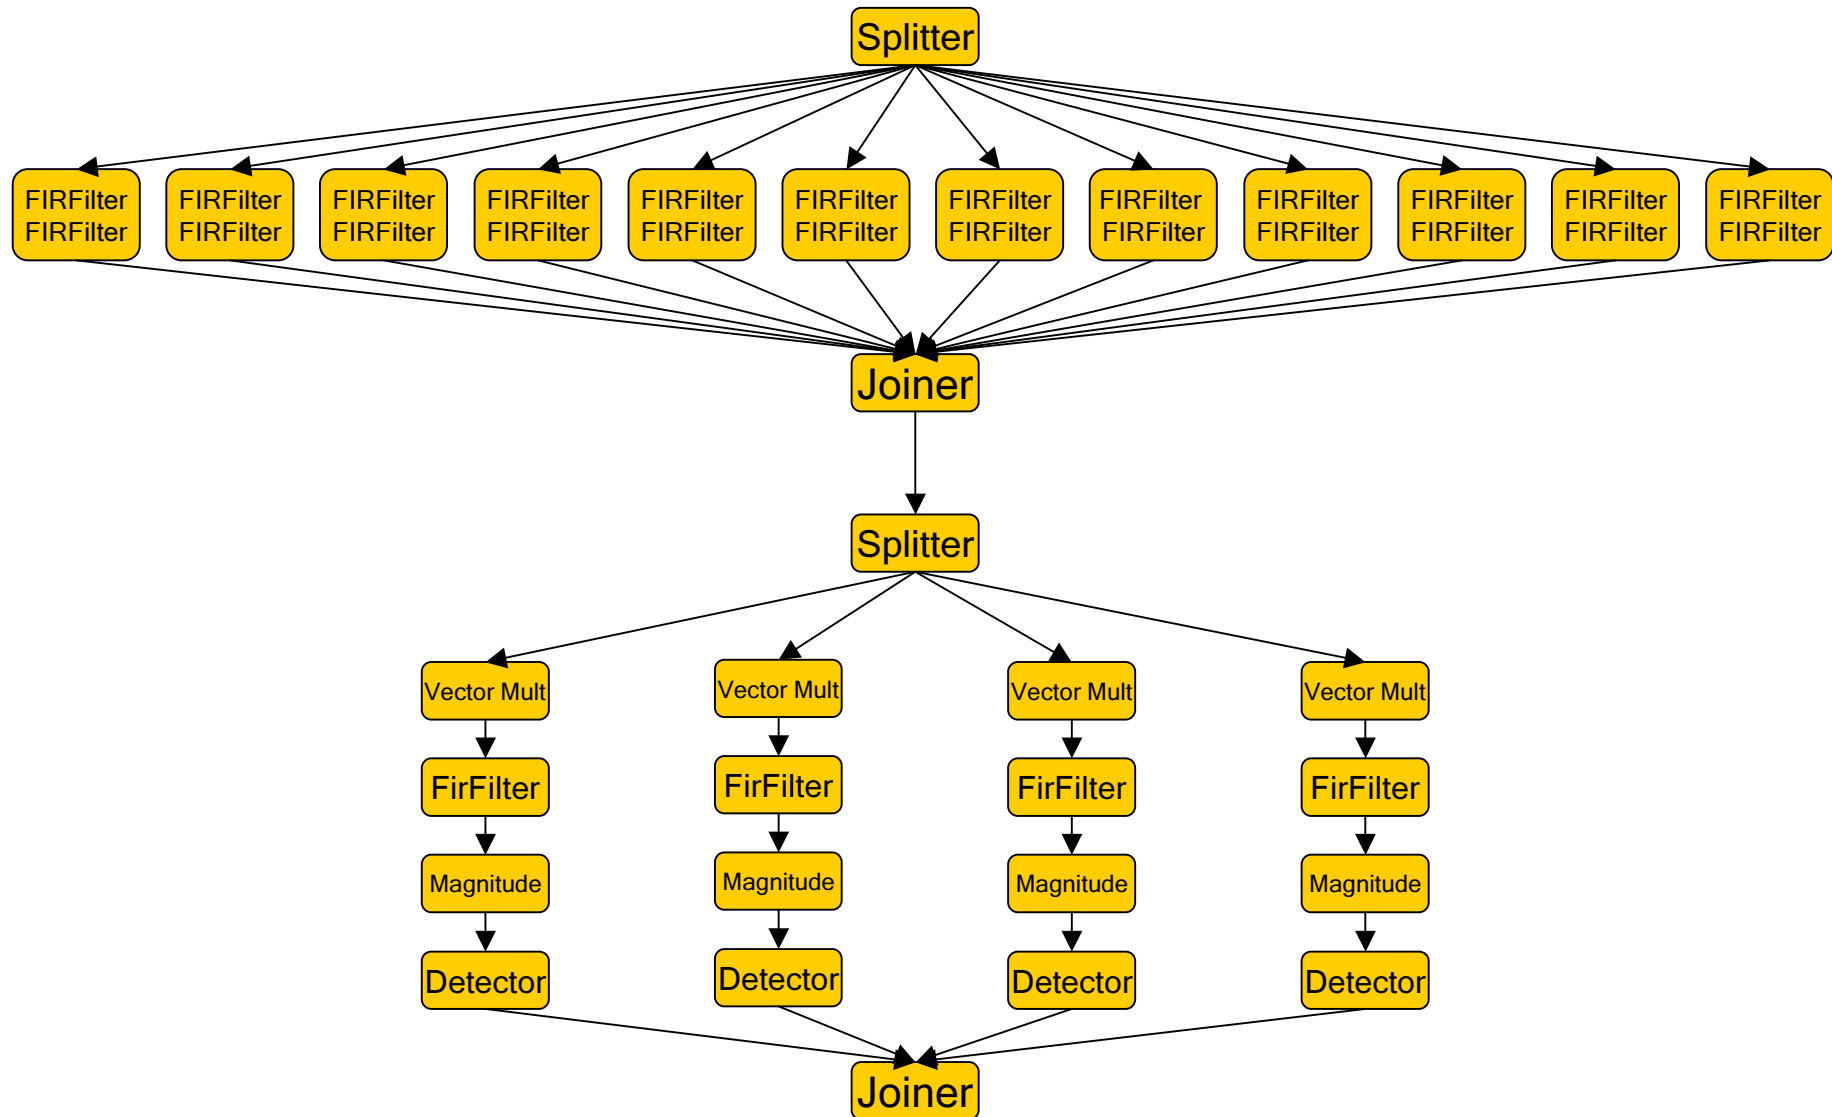
Example: BeamFormer (Original)



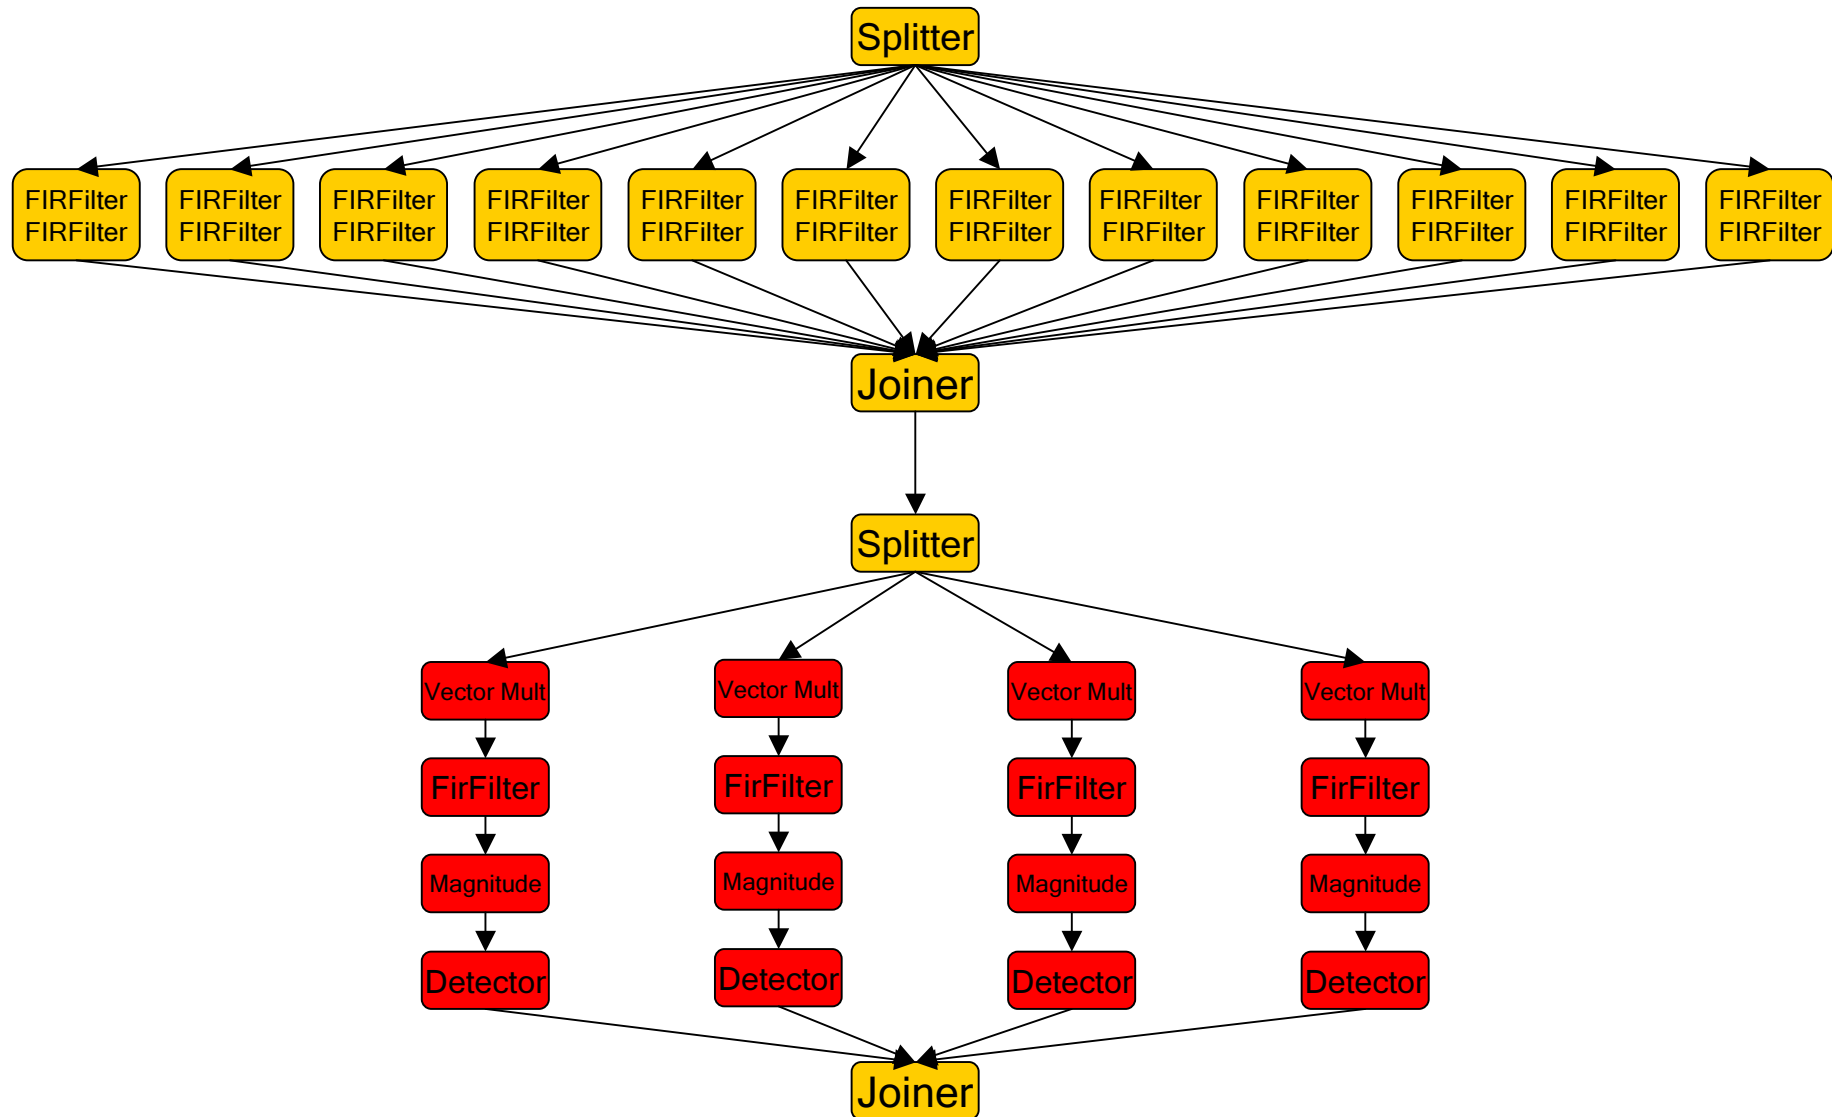
Example: BeamFormer



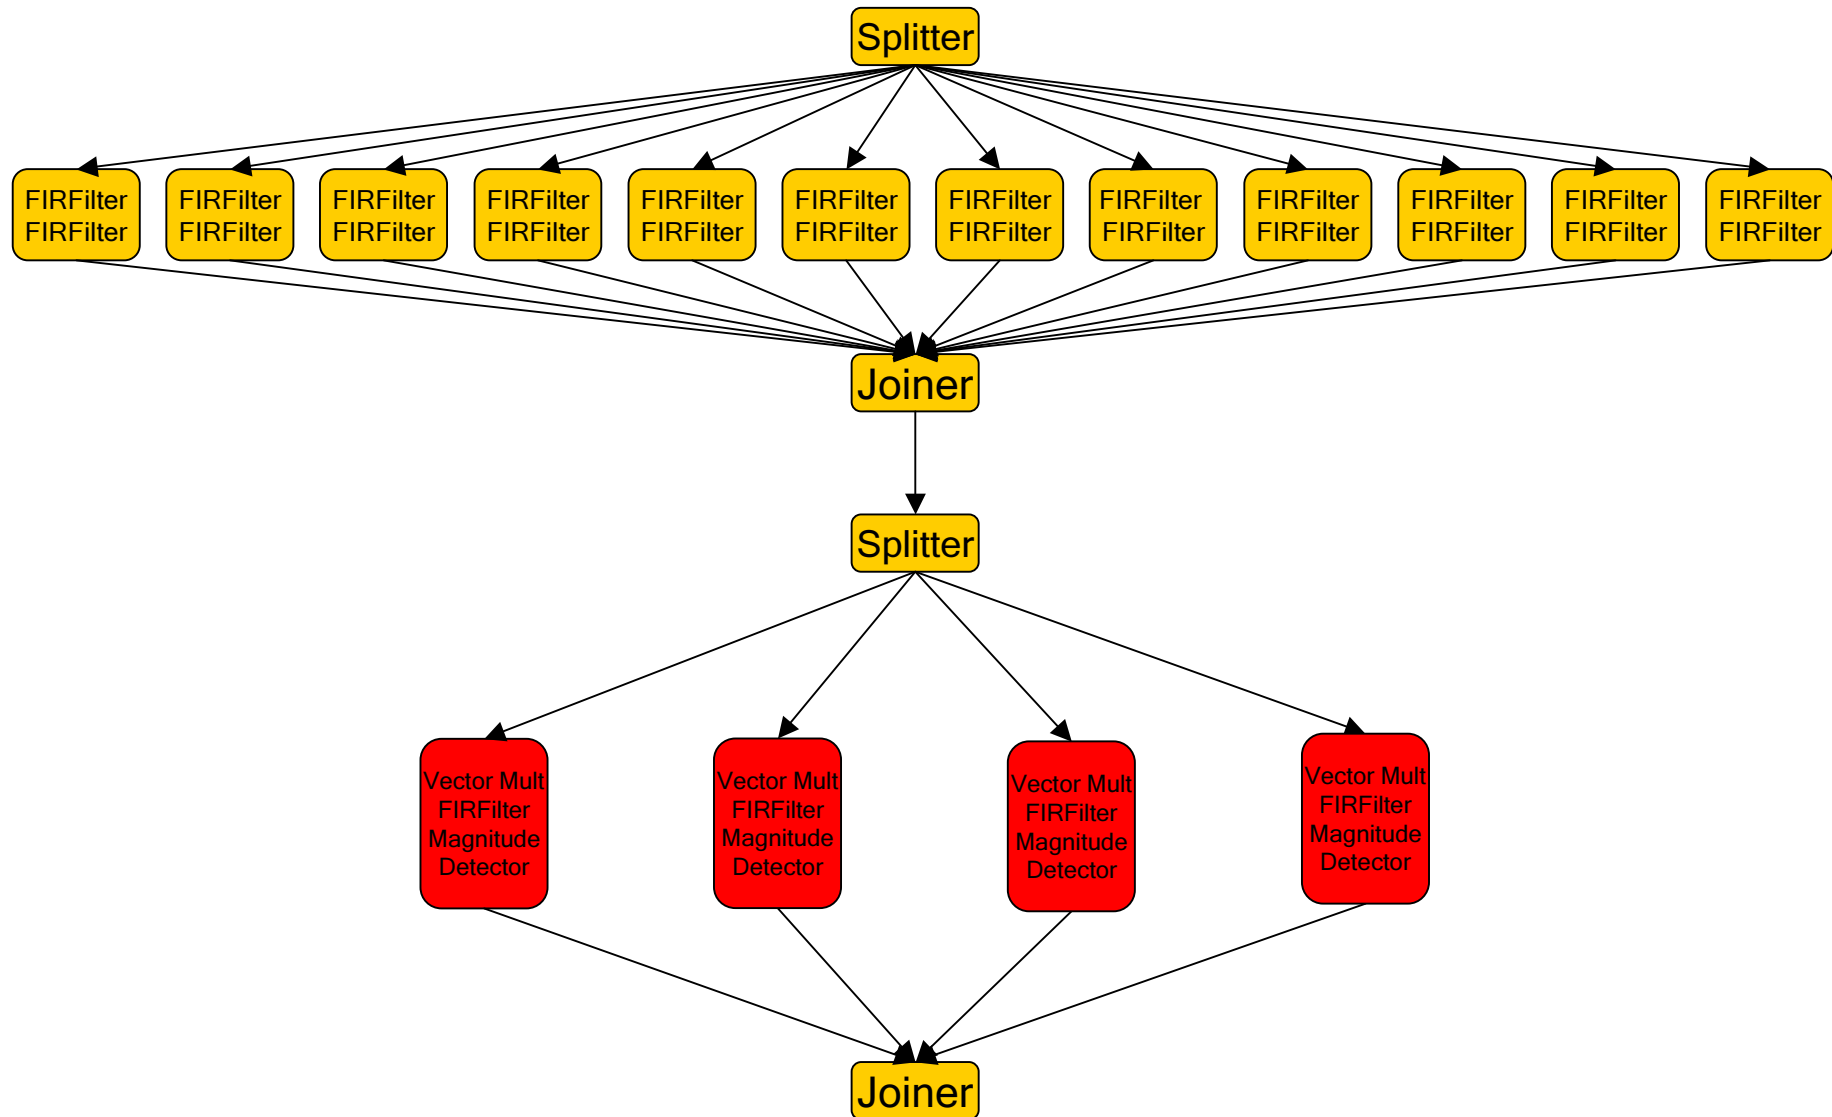
Example: BeamFormer



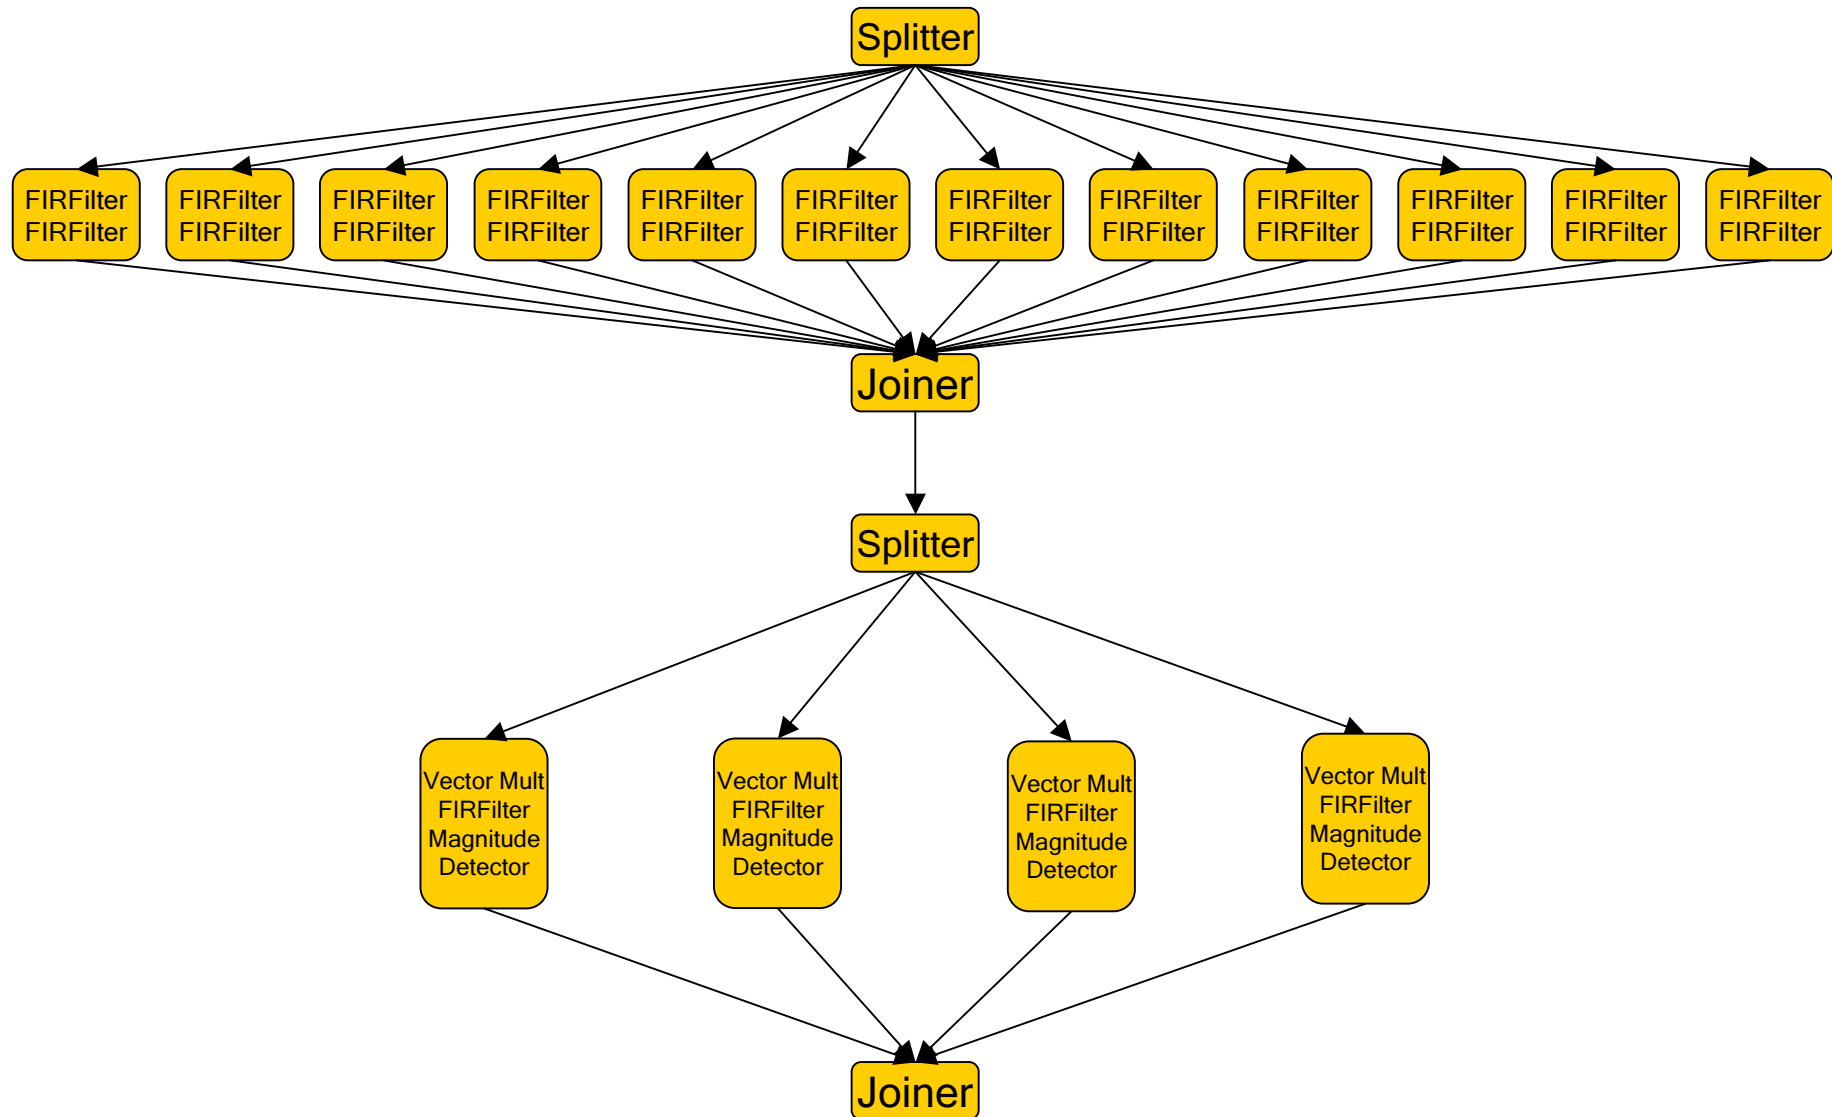
Example: BeamFormer



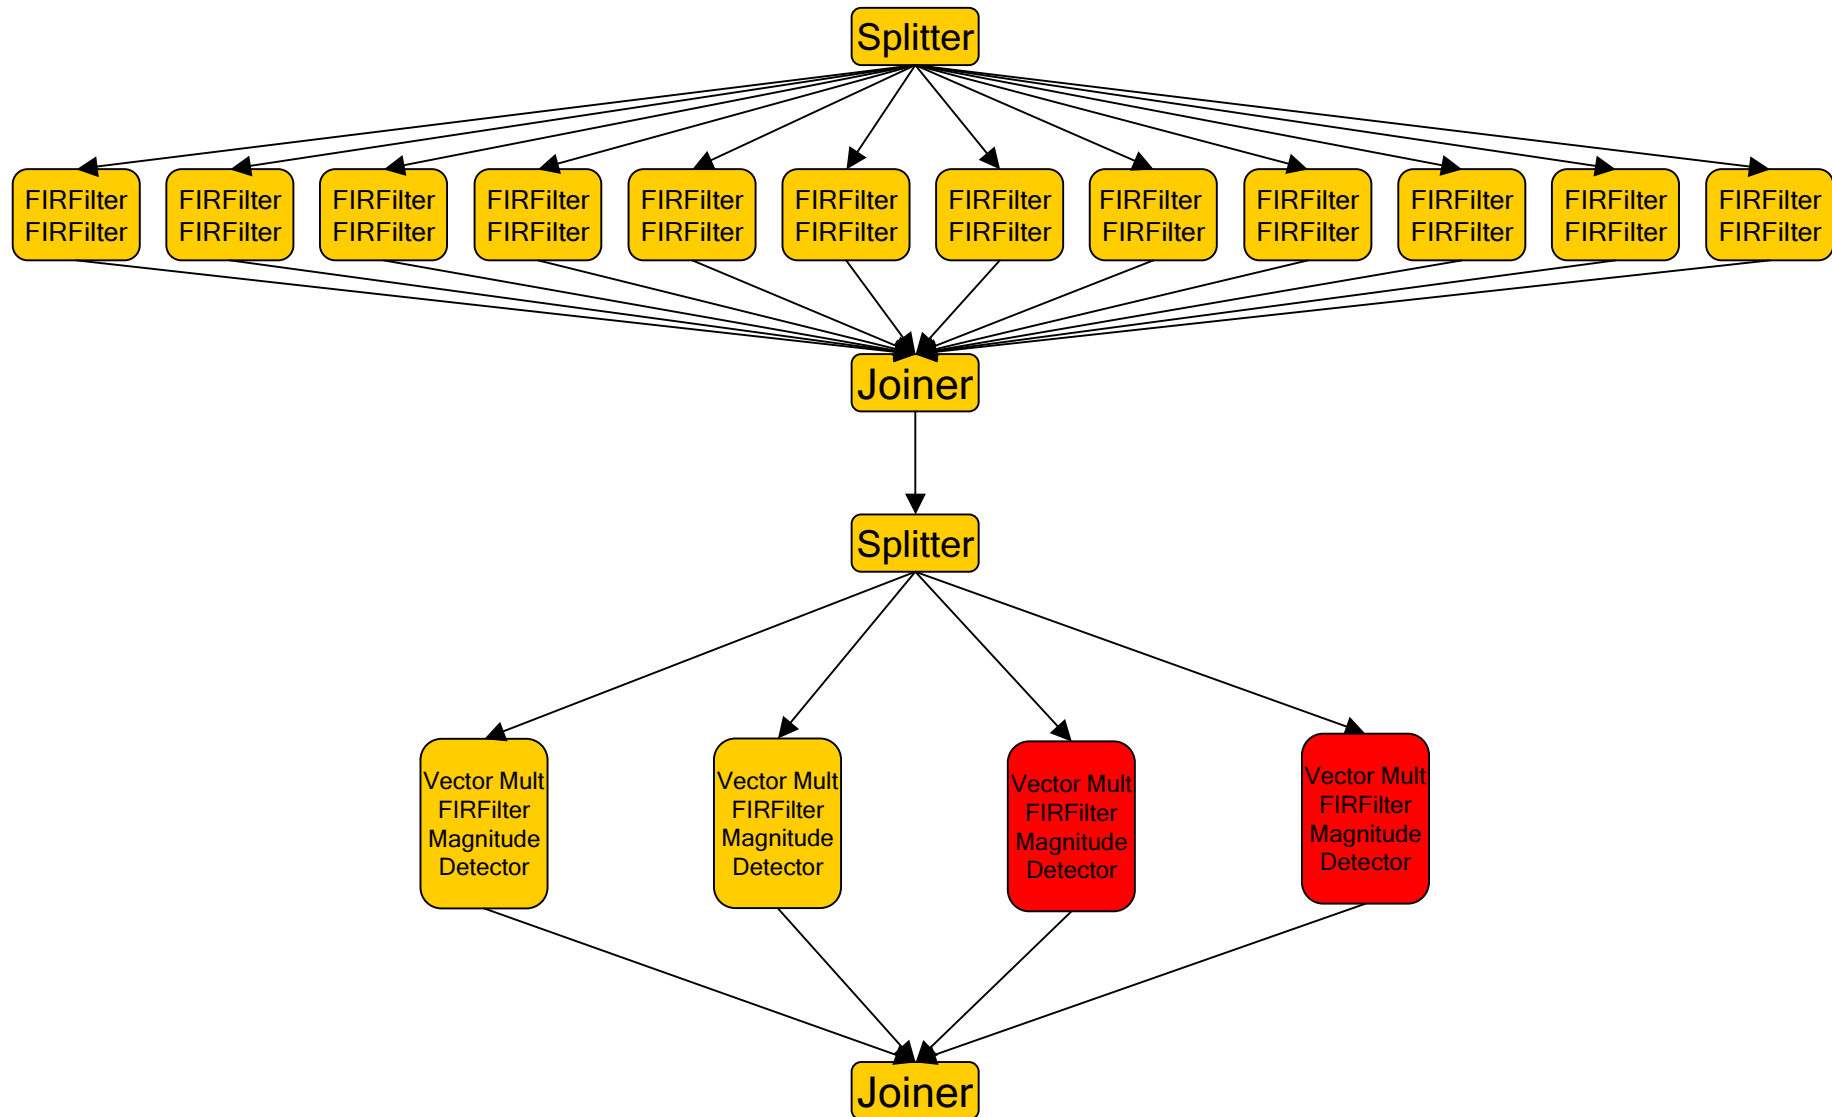
Example: BeamFormer



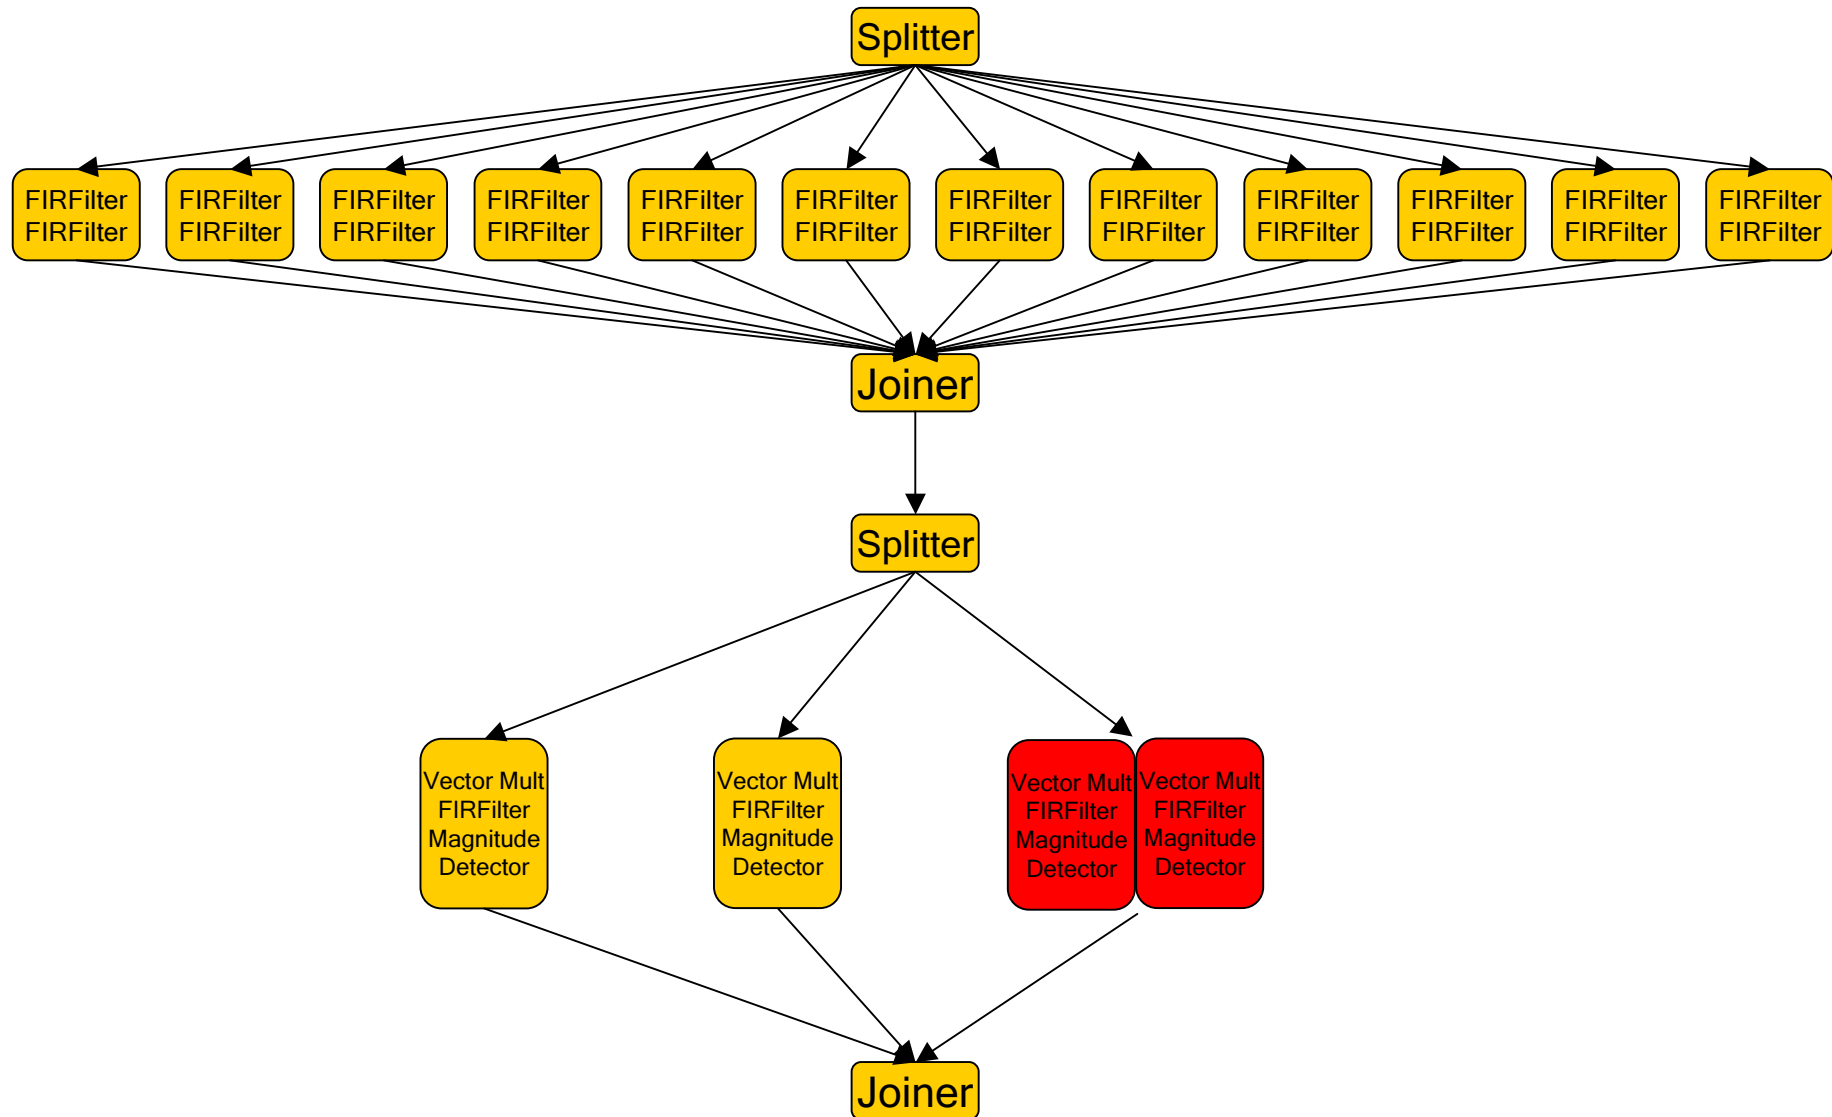
Example: BeamFormer



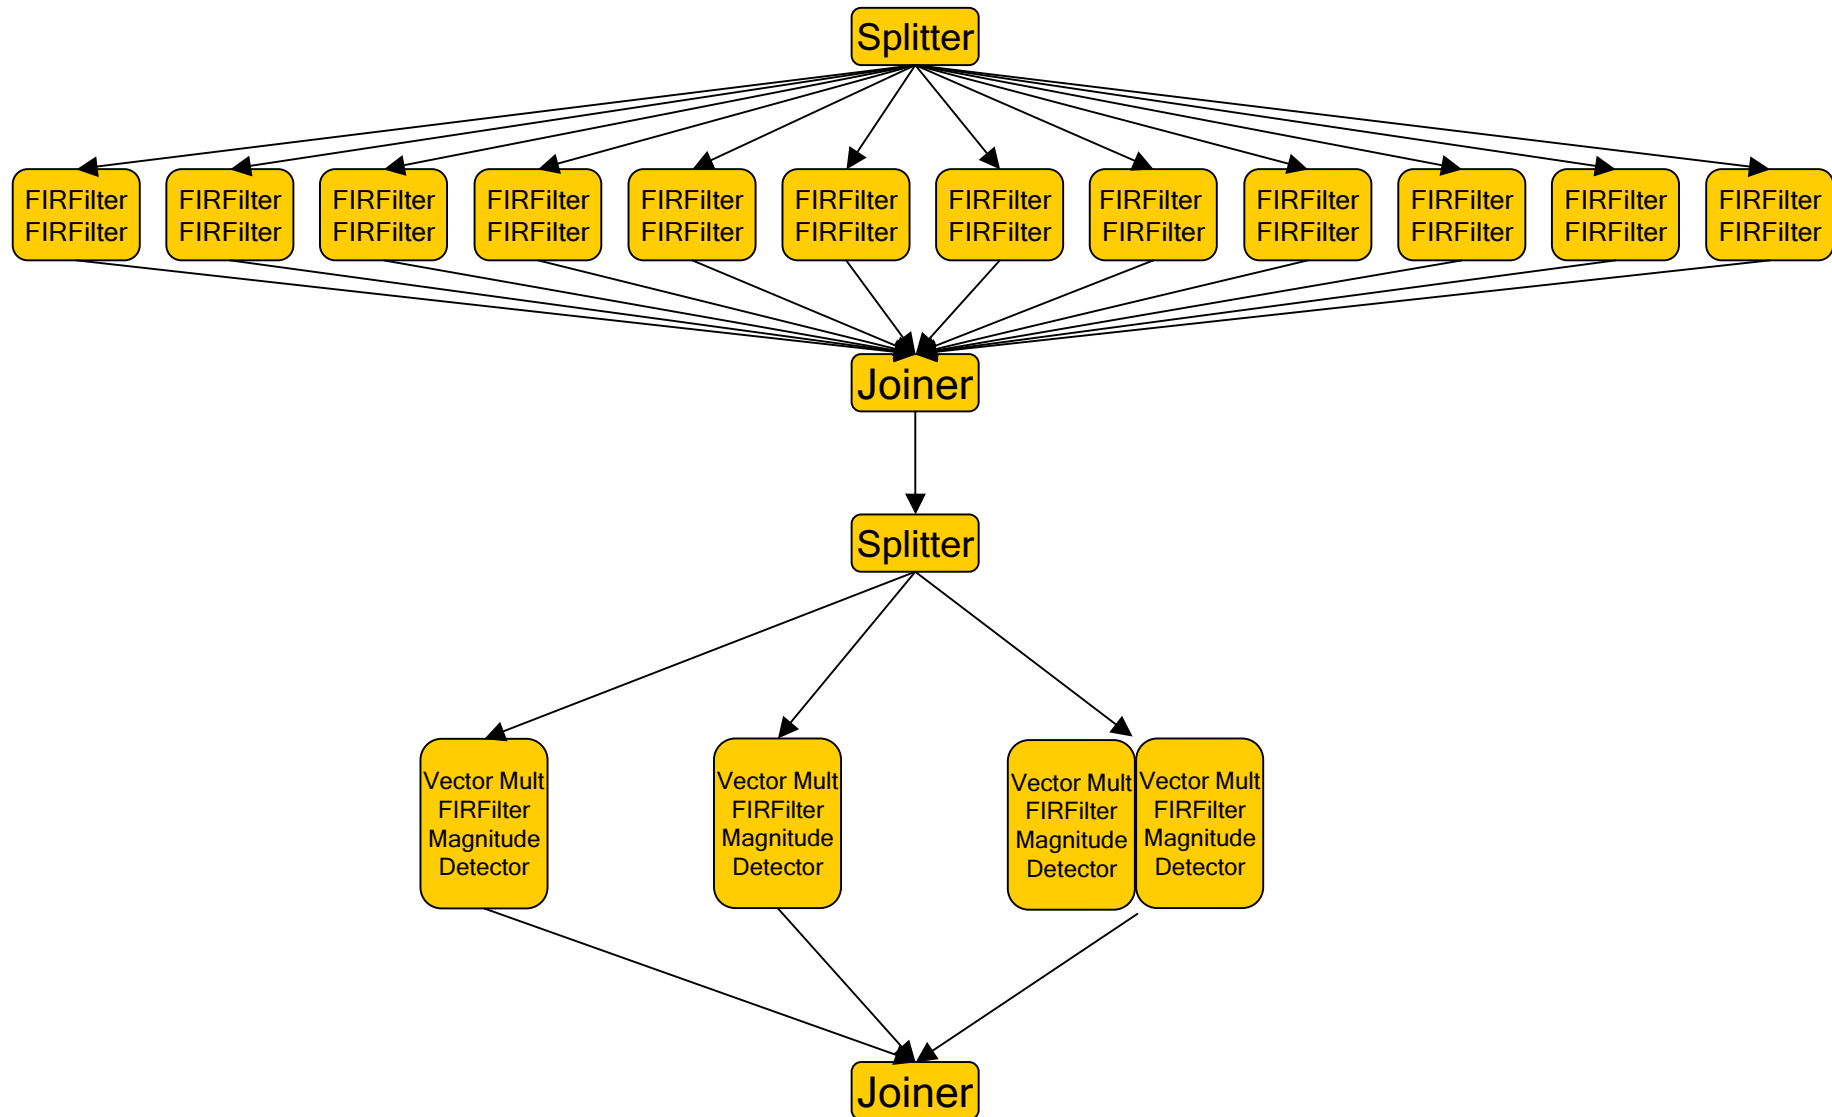
Example: BeamFormer



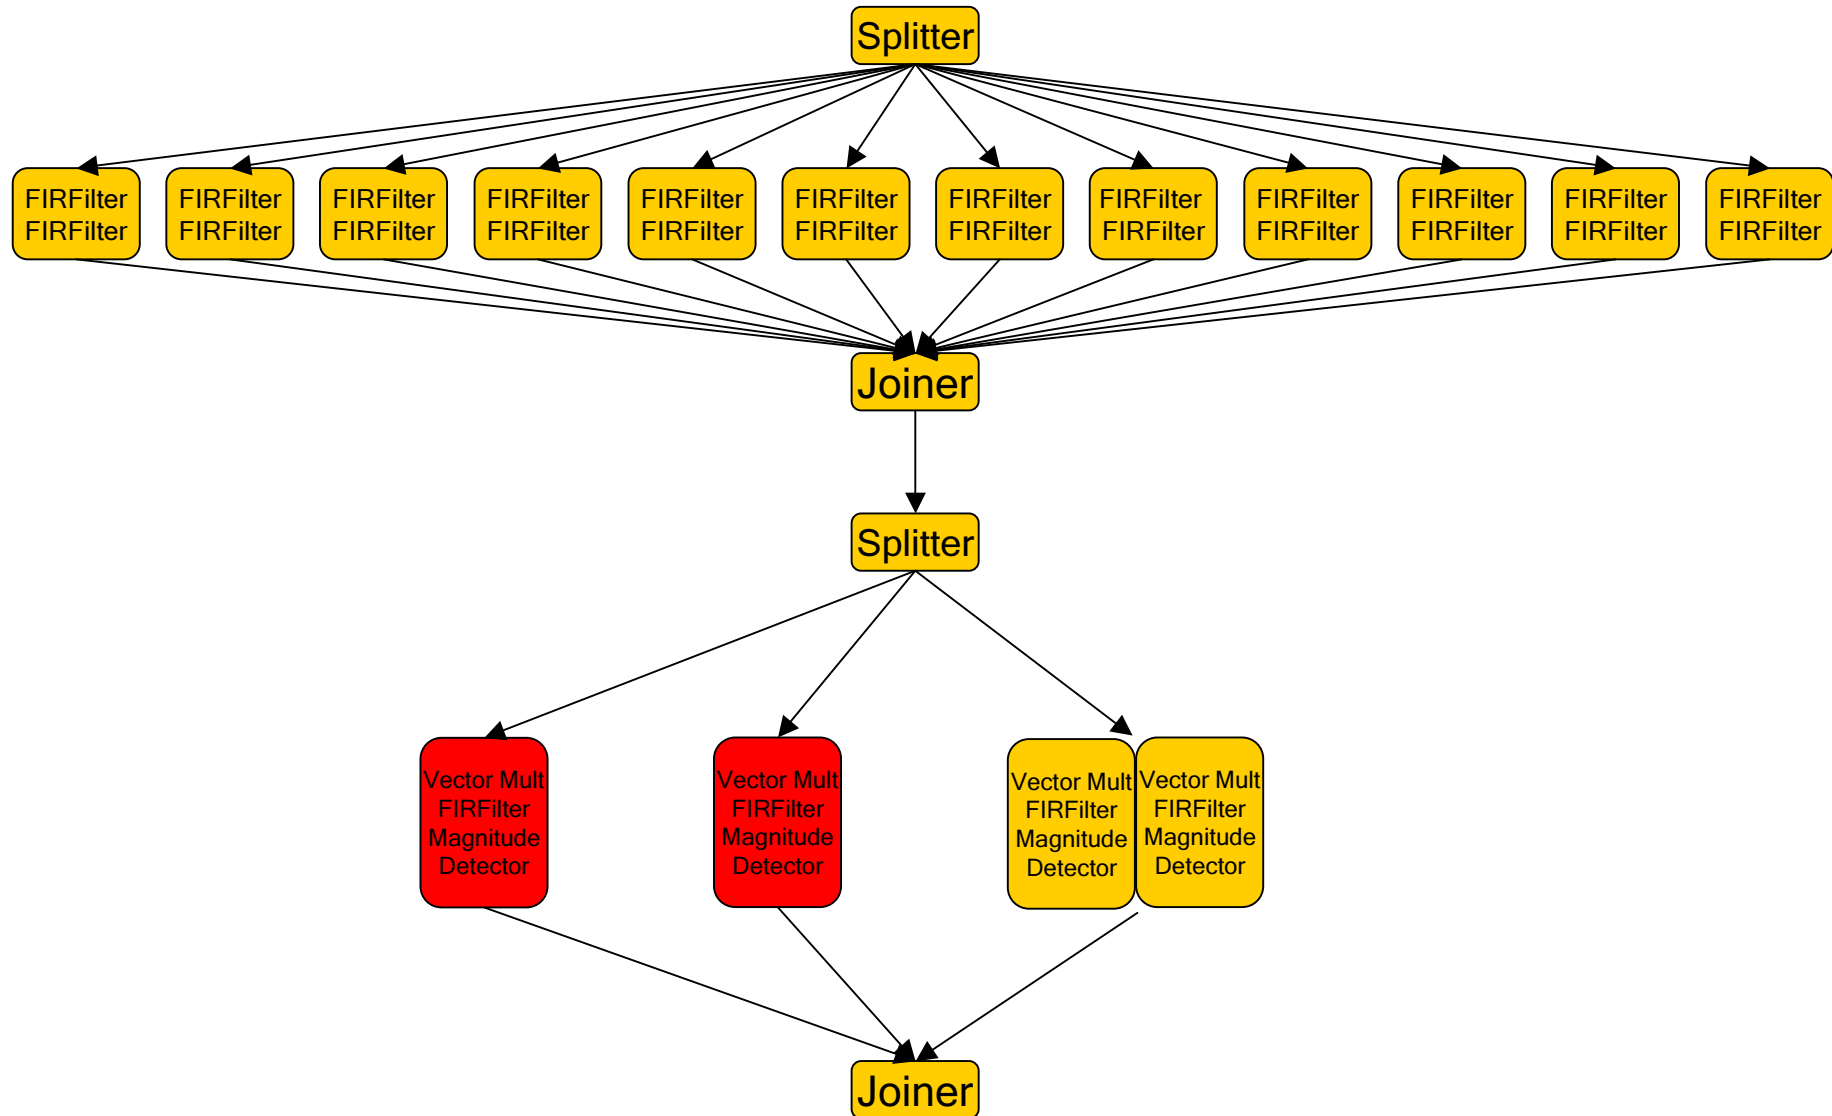
Example: BeamFormer



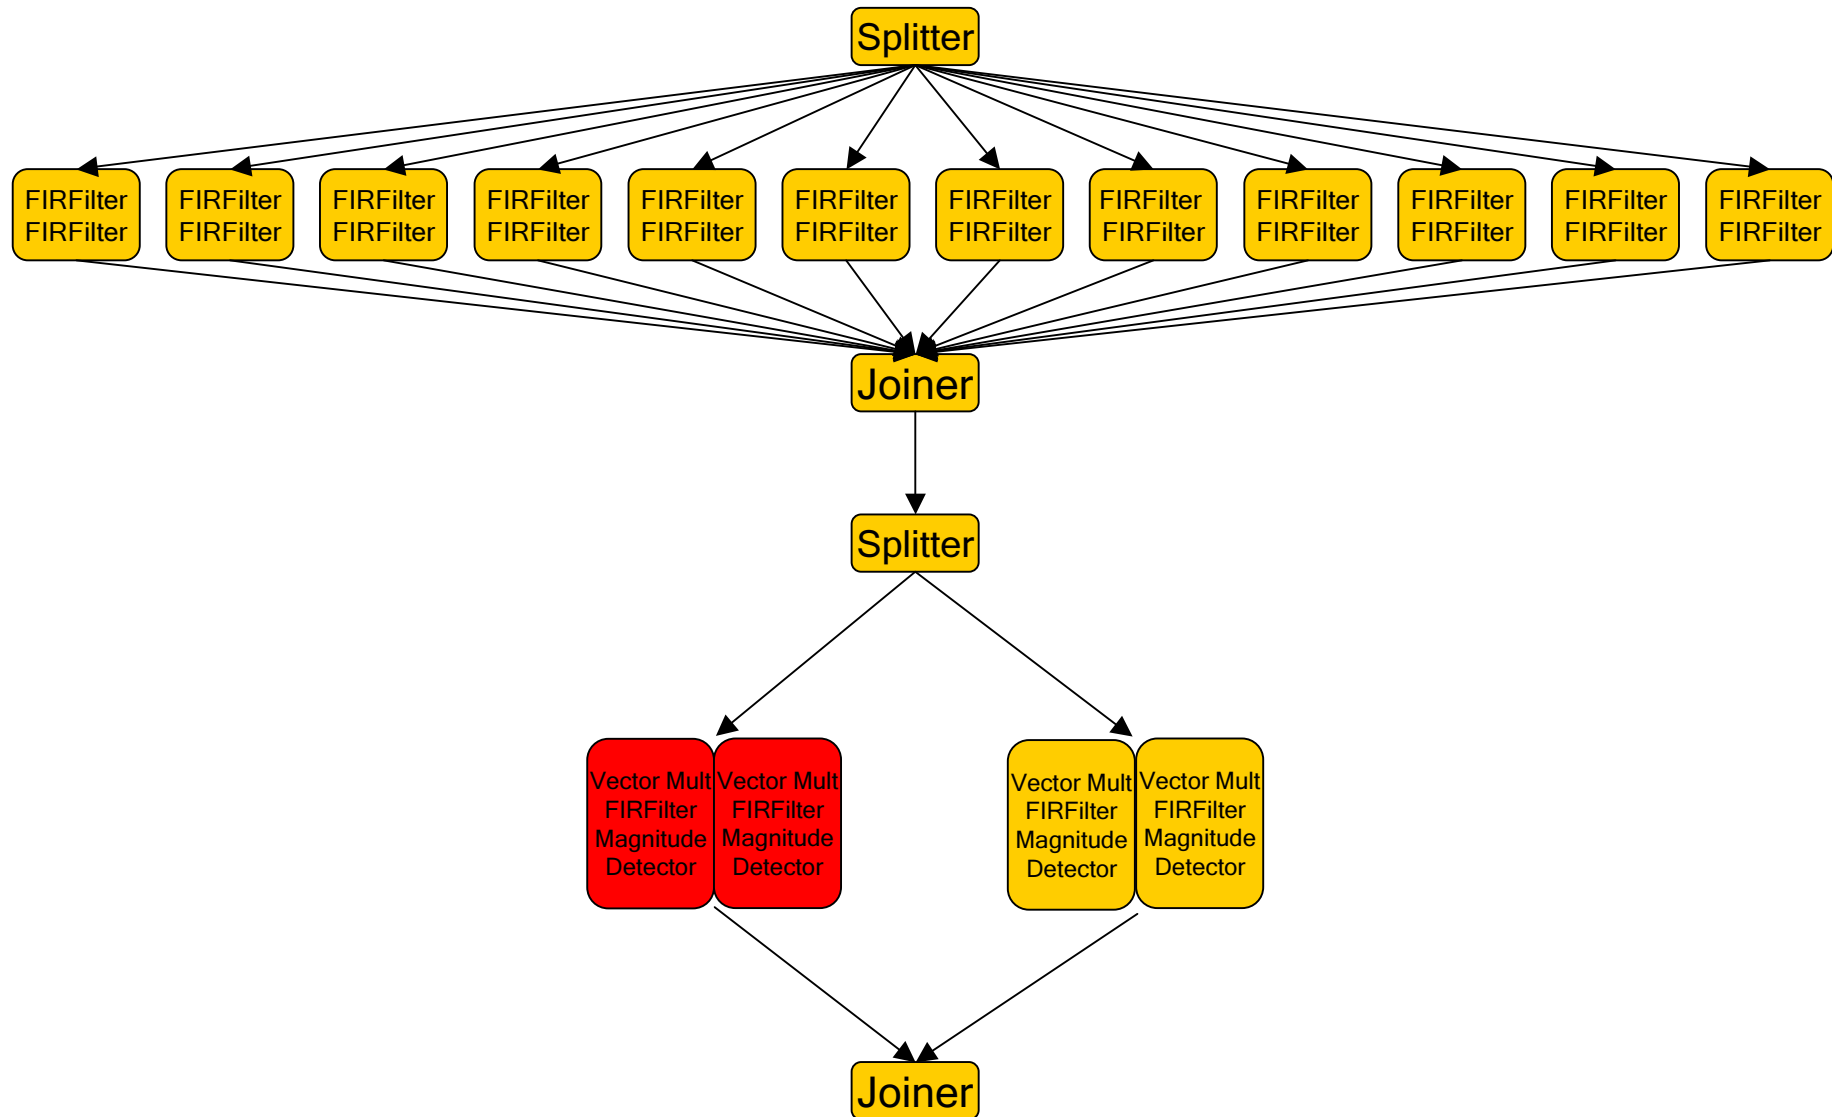
Example: BeamFormer



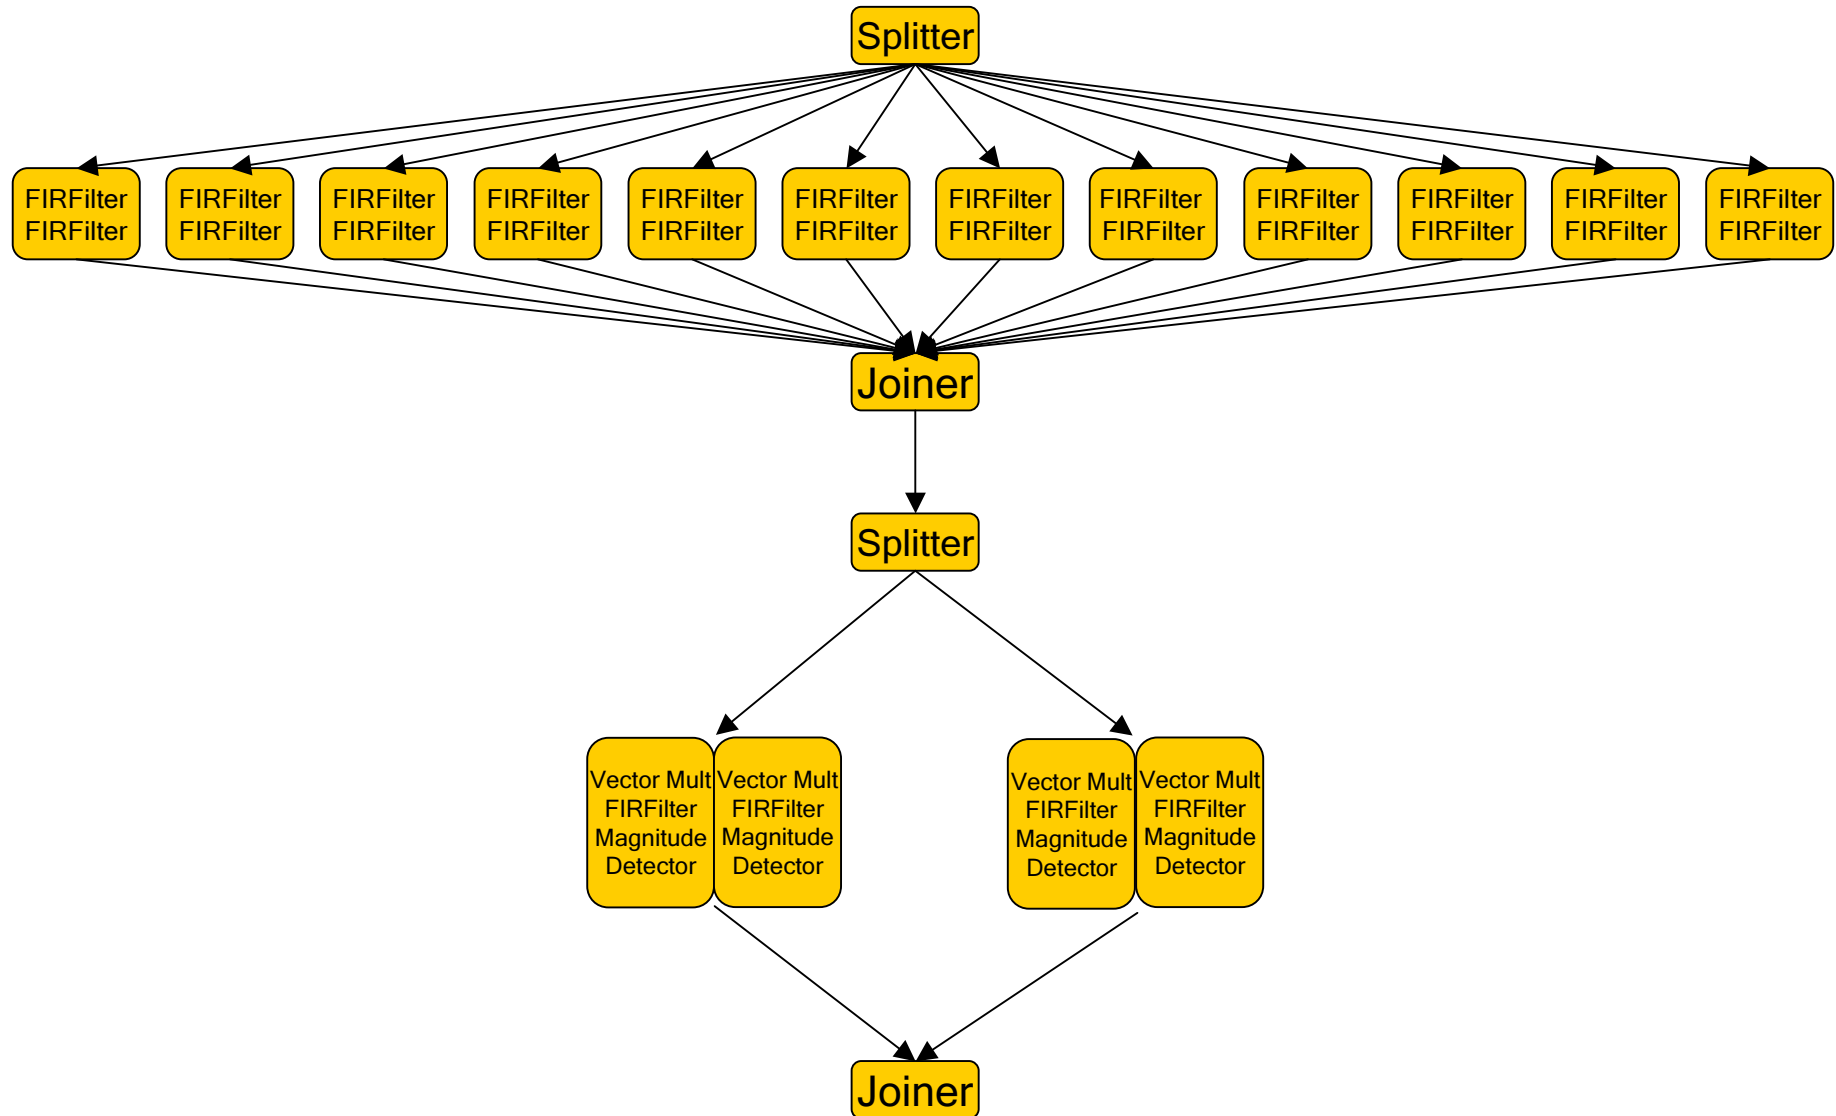
Example: BeamFormer



Example: BeamFormer



Example: BeamFormer (Balanced)



Example: BeamFormer (Balanced)



Outline

- Design of StreamIt
 - Structured Streams
 - Messaging
 - Morphing
- Results
- Conclusions

Conclusions

- Compiler-conscious language design can improve both programmability and performance
 - **Structure** enables local, hierarchical analyses
 - **Messaging** simplifies code, exposes parallelism
 - **Morphing** allows optimization across phases
- Goal: Stream programming at high level of abstraction without sacrificing performance

For More Information

StreamIt Homepage

<http://compiler.lcs.mit.edu/streamit>