

# Optimizing Stream Programs Using Linear State Space Analysis

Sitij Agrawal<sup>1,2</sup>, William Thies<sup>1</sup>, and Saman Amarasinghe<sup>1</sup>

---

<sup>1</sup>Massachusetts Institute of Technology

<sup>2</sup>Sandbridge Technologies


CASES 2005

The logo for StreamIt, featuring the word "StreamIt" in a blue, stylized font. A red arrow points from the "e" to the "t".

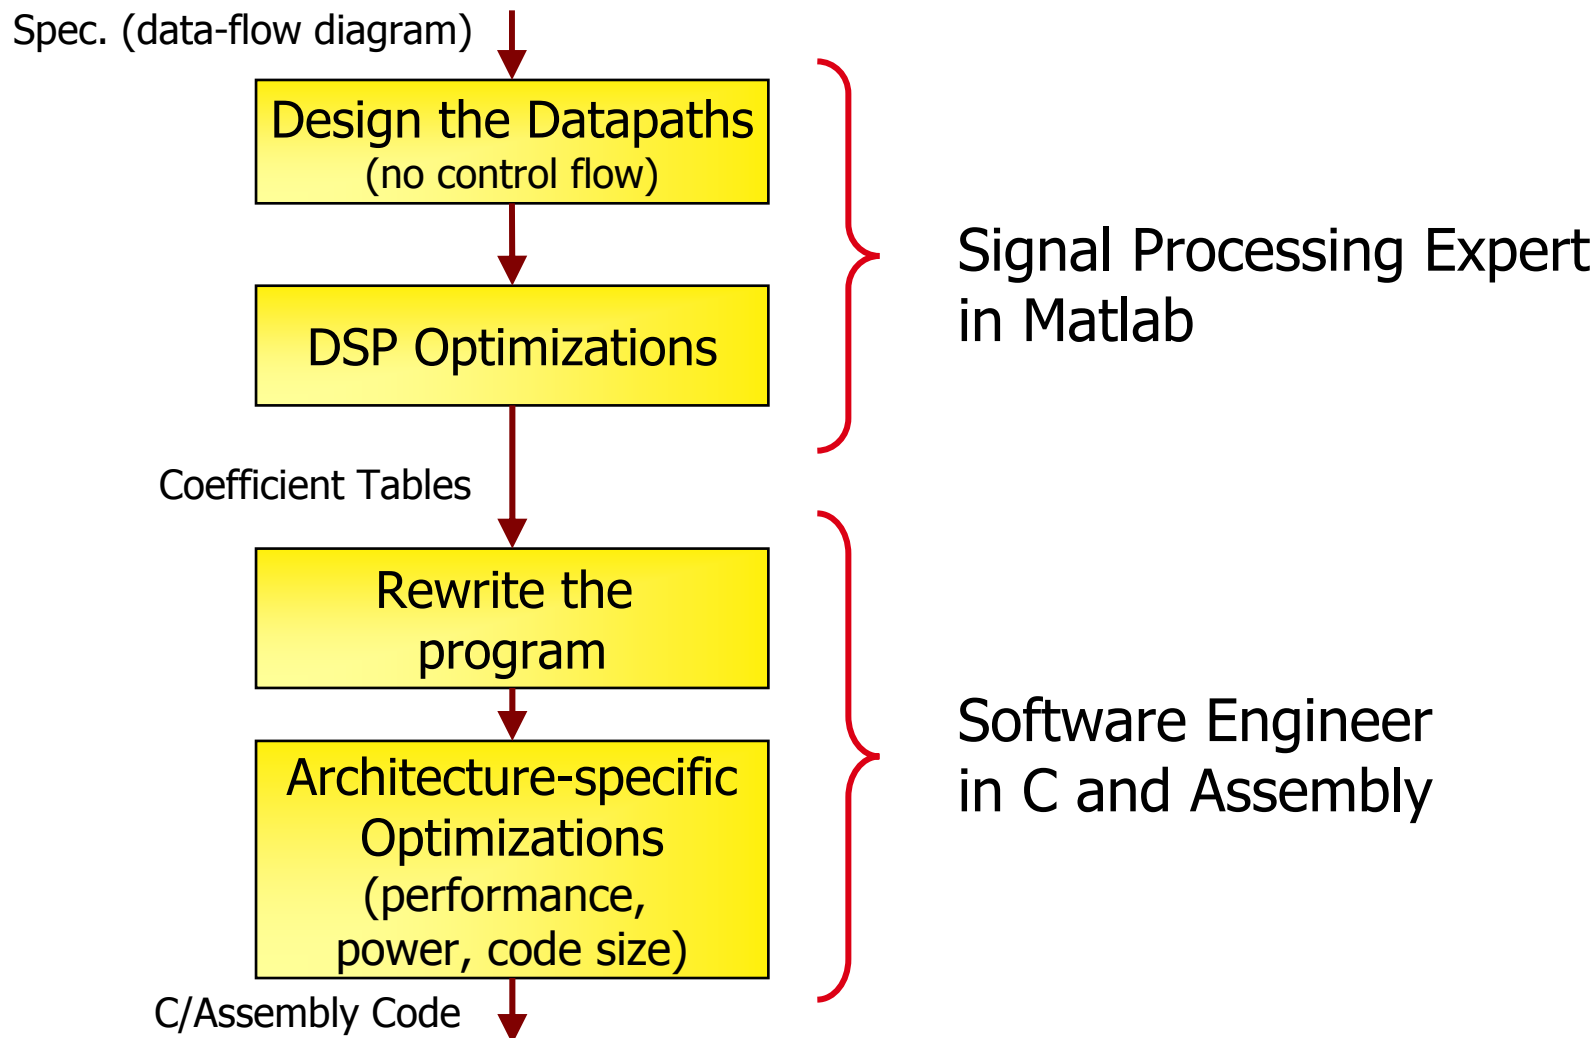
<http://cag.lcs.mit.edu/streamit>

# Streaming Application Domain

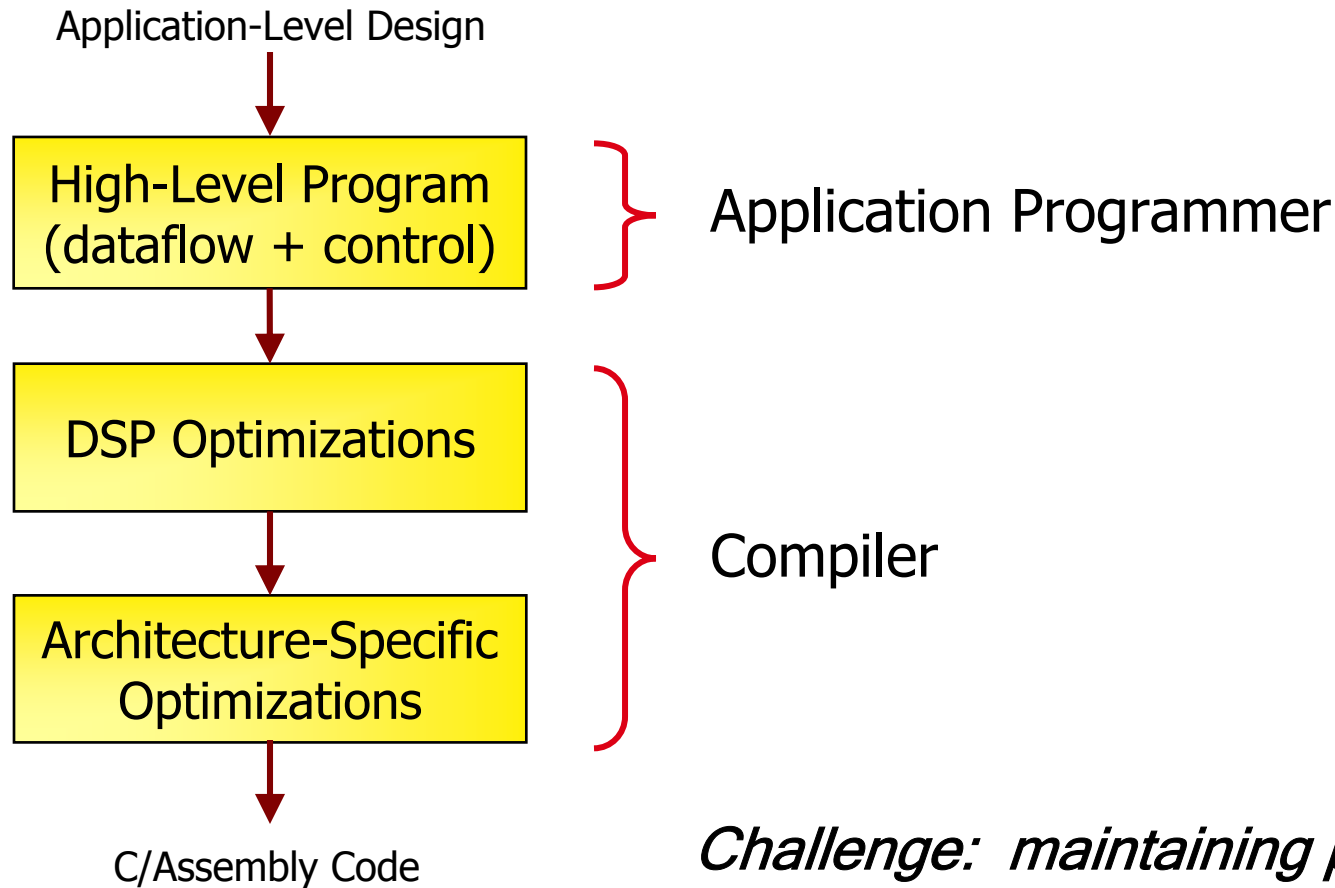
- Based on a stream of data
  - Graphics, multimedia, software radio
  - Radar tracking, microphone arrays, HDTV editing, cell phone base stations
- Properties of stream programs
  - Regular and repeating computation
  - Parallel, independent actors with explicit communication
  - Data items have short lifetimes



# Conventional DSP Design Flow



# Ideal DSP Design Flow



# The StreamIt Language

- Goals:
  - Provide a high-level stream programming model
  - Invent new compiler technology for streams
- Contributions:
  - Language design [CC '02, PPOPP '05]
  - Compiling to tiled architectures [ASPLOS '02, ISCA '04, Graphics Hardware '05]
  - Cache-aware scheduling [LCTES '03, LCTES '05]
  - Domain-specific optimizations [PLDI '03, CASES '05]

# Programming in StreamIt

```
void->void pipeline FMRadio(int N, float lo, float hi) {
```

```
  add AtoD();
```

```
  add FMDemod();
```

```
  add splitjoin {
    split duplicate;
    for (int i=0; i<N; i++) {
      add pipeline {
```

```
        add LowPassFilter(lo + i*(hi - lo)/N);
```


```
        add HighPassFilter(lo + i*(hi - lo)/N);
```

```
      }
    }
    join roundrobin();
  }
```

```
  add Adder();
```


```
  add Speaker();
```

```
}
```



# Example StreamIt Filter

```
float->float filter LowPassButterWorth (float sampleRate, float cutoff) {  
    float coeff;  
    float x;  
  
    init {  
        coeff = calcCoeff(sampleRate, cutoff);  
    }  
  
    work peek 2 push 1 pop 1 {  
        x = peek(0) + peek(1) + coeff * x;  
        push(x);  
        pop();  
    }  
}
```




# Focus: Linear State Space Filters

- Properties:
  1. Outputs are linear function of inputs and states
  2. New states are linear function of inputs and states
- Most common target of DSP optimizations
  - FIR / IIR filters
  - Linear difference equations
  - Upsamplers / downsamplers
  - DCTs



# Representing State Space Filters


- A state space filter is a tuple  $\langle A, B, C, D \rangle$



# Representing State Space Filters

- A state space filter is a tuple  $\langle A, B, C, D \rangle$

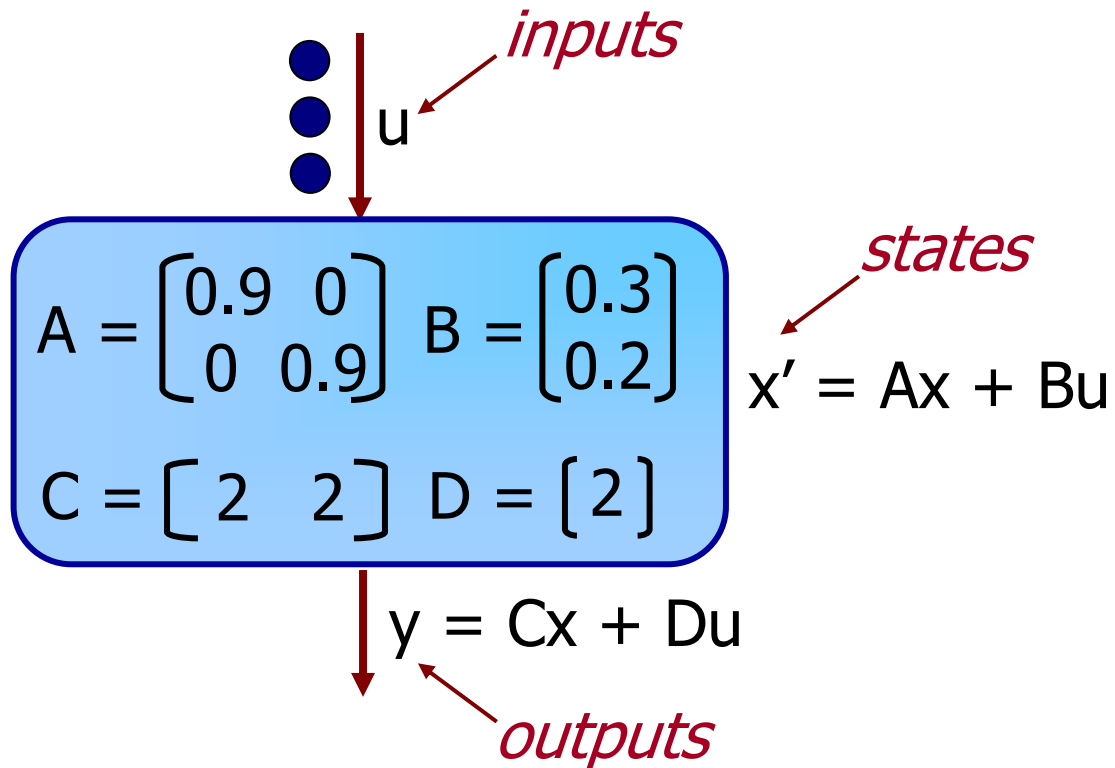
```
float->float filter IIR {  
  float x1, x2;  
  work push 1 pop 1 {  
    float u = pop();  
    push(2*(x1+x2+u));  
    x1 = 0.9*x1 + 0.3*u;  
    x2 = 0.9*x2 + 0.2*u;  
  }  
}
```



# Representing State Space Filters

- A state space filter is a tuple  $\langle A, B, C, D \rangle$


```
float->float filter IIR {  
  float x1, x2;  
  work push 1 pop 1 {  
    float u = pop();  
    push(2*(x1+x2+u));  
    x1 = 0.9*x1 + 0.3*u;  
    x2 = 0.9*x2 + 0.2*u;  
  }  
}
```



# Representing State Space Filters

- A state space filter is a tuple  $\langle A, B, C, D \rangle$

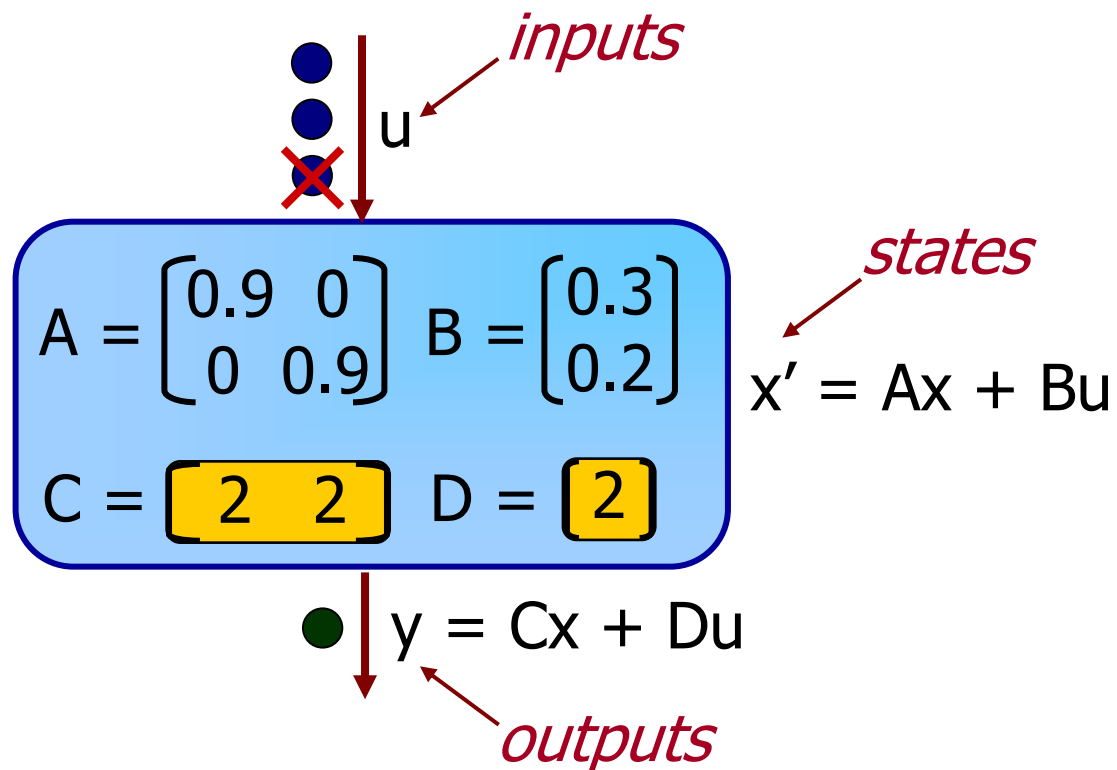
```
float->float filter IIR {  
  float x1, x2;  
  work push 1 pop 1 {  
    float u = pop();  
    push(2*(x1+x2+u));  
    x1 = 0.9*x1 + 0.3*u;  
    x2 = 0.9*x2 + 0.2*u;  
  }  
}
```



# Representing State Space Filters

- A state space filter is a tuple  $\langle A, B, C, D \rangle$


```
float->float filter IIR {  
  float x1, x2;  
  work push 1 pop 1 {  
    float u = pop();  
    push(2*(x1+x2+u));  
    x1 = 0.9*x1 + 0.3*u;  
    x2 = 0.9*x2 + 0.2*u;  
  }  
}
```



# Representing State Space Filters

- A state space filter is a tuple  $\langle A, B, C, D \rangle$


```
float->float filter IIR {  
  float x1, x2;  
  work push 1 pop 1 {  
    float u = pop();  
    push(2*(x1+x2+u));  
    x1 = 0.9*x1 + 0.3*u;  
    x2 = 0.9*x2 + 0.2*u;  
  }  
}
```



# Representing State Space Filters

- A state space filter is a tuple  $\langle A, B, C, D \rangle$

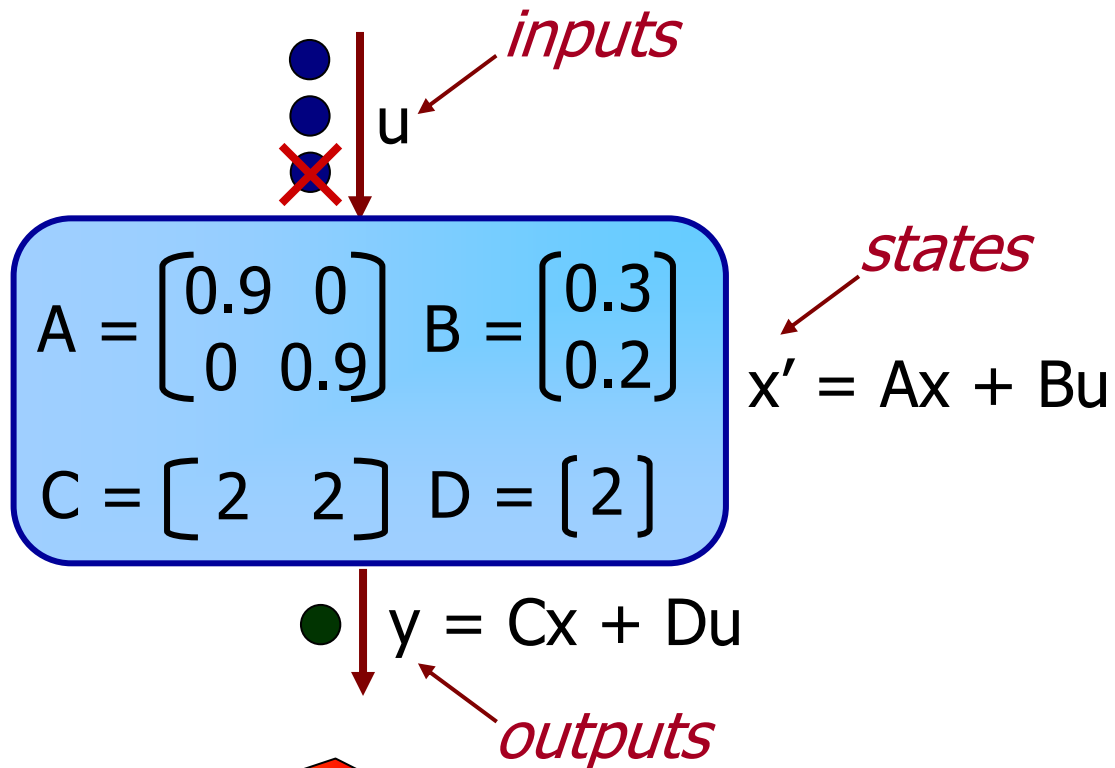
```
float->float filter IIR {  
  float x1, x2;  
  work push 1 pop 1 {  
    float u = pop();  
    push(2*(x1+x2+u));  
    x1 = 0.9*x1 + 0.3*u;  
    x2 = 0.9*x2 + 0.2*u;  
  }  
}
```



# Representing State Space Filters

- A state space filter is a tuple  $\langle A, B, C, D \rangle$

```
float->float filter IIR {  
  float x1, x2;  
  work push 1 pop 1 {  
    float u = pop();  
    push(2*(x1+x2+u));  
    x1 = 0.9*x1 + 0.3*u;  
    x2 = 0.9*x2 + 0.2*u;  
  }  
}
```



Linear dataflow analysis



# State Space Optimizations

1. State removal
2. Reducing the number of parameters
3. Combining adjacent filters

# Change-of-Basis Transformation

$$x' = Ax + Bu$$

$$y = Cx + Du$$

# Change-of-Basis Transformation

$$x' = Ax + Bu$$

$$y = Cx + Du$$



$T$  = invertible matrix

$$Tx' = TAx + TBu$$

$$y = Cx + Du$$

# Change-of-Basis Transformation

$$x' = Ax + Bu$$

$$y = Cx + Du$$



$T$  = invertible matrix

$$Tx' = TA(T^{-1}T)x + TBu$$

$$y = C(T^{-1}T)x + Du$$

# Change-of-Basis Transformation

$$x' = Ax + Bu$$

$$y = Cx + Du$$



$T$  = invertible matrix

$$Tx' = TAT^{-1}(Tx) + TBu$$

$$y = CT^{-1}(Tx) + Du$$

# Change-of-Basis Transformation

$$x' = Ax + Bu$$

$$y = Cx + Du$$



$T =$  invertible matrix,  $z = Tx$

$$Tx' = TAT^{-1}(Tx) + TBu$$

$$y = CT^{-1}(Tx) + Du$$

# Change-of-Basis Transformation

$$x' = Ax + Bu$$

$$y = Cx + Du$$



$T =$  invertible matrix,  $z = Tx$

$$z' = TAT^{-1}z + TBu$$

$$y = CT^{-1}z + Du$$

# Change-of-Basis Transformation

$$x' = Ax + Bu$$

$$y = Cx + Du$$



$T =$  invertible matrix,  $z = Tx$

$$z' = A'z + B'u$$

$$y = C'z + D'u$$

$$A' = TAT^{-1} \quad B' = TB$$

$$C' = CT^{-1} \quad D' = D$$



# Change-of-Basis Transformation

$$x' = Ax + Bu$$

$$y = Cx + Du$$



$T =$  invertible matrix,  $z = Tx$

$$z' = A'z + B'u$$

$$y = C'z + D'u$$

$$A' = TAT^{-1} \quad B' = TB$$

$$C' = CT^{-1} \quad D' = D$$

Can map original states  $x$  to transformed states  $z = Tx$  without changing I/O behavior

# 1) State Removal

- Can remove states which are:
  - a. Unreachable – do not depend on input
  - b. Unobservable – do not affect output
- To expose unreachable states, reduce  $[A \mid B]$  to a kind of row-echelon form
  - For unobservable states, reduce  $[A^T \mid C^T]$
- Automatically finds minimal number of states

# State Removal Example

$$\begin{aligned} x' &= \begin{bmatrix} 0.9 & 0 \\ 0 & 0.9 \end{bmatrix} x + \begin{bmatrix} 0.3 \\ 0.2 \end{bmatrix} u \\ y &= \begin{bmatrix} 2 & 2 \end{bmatrix} x + 2u \end{aligned}$$

$$T = \begin{bmatrix} 1 & 0 \\ 1 & 1 \end{bmatrix}$$



$$\begin{aligned} x' &= \begin{bmatrix} 0.9 & 0 \\ 0 & 0.9 \end{bmatrix} x + \begin{bmatrix} 0.3 \\ \mathbf{0.5} \end{bmatrix} u \\ y &= \begin{bmatrix} \mathbf{0} & 2 \end{bmatrix} x + 2u \end{aligned}$$



```
float->float filter IIR {
  float x1, x2;
  work push 1 pop 1 {
    float u = pop();
    push(2*(x1+x2+u));
    x1 = 0.9*x1 + 0.3*u;
    x2 = 0.9*x2 + 0.2*u;
  } }
```

# State Removal Example

$$\begin{aligned} x' &= \begin{bmatrix} 0.9 & 0 \\ 0 & 0.9 \end{bmatrix} x + \begin{bmatrix} 0.3 \\ 0.2 \end{bmatrix} u \\ y &= \begin{bmatrix} 2 & 2 \end{bmatrix} x + 2u \end{aligned}$$

$$T = \begin{bmatrix} 1 & 0 \\ 1 & 1 \end{bmatrix}$$



$$\begin{aligned} x' &= \begin{bmatrix} 0.9 & 0 \\ 0 & 0.9 \end{bmatrix} x + \begin{bmatrix} 0.3 \\ \mathbf{0.5} \end{bmatrix} u \\ y &= \begin{bmatrix} \mathbf{0} & 2 \end{bmatrix} x + 2u \end{aligned}$$

*x1 is unobservable*

```
float->float filter IIR {
  float x1, x2;
  work push 1 pop 1 {
    float u = pop();
    push(2*(x1+x2+u));
    x1 = 0.9*x1 + 0.3*u;
    x2 = 0.9*x2 + 0.2*u;
  } }
```

# State Removal Example

$$\begin{aligned} x' &= \begin{bmatrix} 0.9 & 0 \\ 0 & 0.9 \end{bmatrix} x + \begin{bmatrix} 0.3 \\ 0.2 \end{bmatrix} u \\ y &= \begin{bmatrix} 2 & 2 \end{bmatrix} x + 2u \end{aligned}$$

$$T = \begin{bmatrix} 1 & 0 \\ 1 & 1 \end{bmatrix}$$



$$\begin{aligned} x' &= 0.9x + 0.5u \\ y &= 2x + 2u \end{aligned}$$



```
float->float filter IIR {
  float x1, x2;
  work push 1 pop 1 {
    float u = pop();
    push(2*(x1+x2+u));
    x1 = 0.9*x1 + 0.3*u;
    x2 = 0.9*x2 + 0.2*u;
  } }
```




```
float->float filter IIR {
  float x;
  work push 1 pop 1 {
    float u = pop();
    push(2*(x+u));
    x = 0.9*x + 0.5*u;
  } }
```

# State Removal Example

9 FLOPs  
12 load/store

---


output



5 FLOPs  
8 load/store

---

output



```
float->float filter IIR {  
  float x1, x2;  
  work push 1 pop 1 {  
    float u = pop();  
    push(2*(x1+x2+u));  
    x1 = 0.9*x1 + 0.3*u;  
    x2 = 0.9*x2 + 0.2*u;  
  }  
}
```

```
float->float filter IIR {  
  float x;  
  work push 1 pop 1 {  
    float u = pop();  
    push(2*(x+u));  
    x = 0.9*x + 0.5*u;  
  }  
}
```

## 2) Parameter Reduction

- Goal:  
Convert matrix entries (parameters) to 0 or 1
- Allows static evaluation:
  - $1 * x \rightarrow x$       Eliminate 1 multiply
  - $0 * x + y \rightarrow y$       Eliminate 1 multiply, 1 add
- Algorithm (Ackerman & Bucy, 1971)
  - Also reduces matrices  $[A \mid B]$  and  $[A^T \mid C^T]$
  - Attains a canonical form with few parameters

# Parameter Reduction Example

$$\begin{aligned}x' &= 0.9x + 0.5u \\ y &= 2x + 2u\end{aligned}$$

$$T = [2]$$



$$\begin{aligned}x' &= 0.9x + \mathbf{1}u \\ y &= \mathbf{1}x + 2u\end{aligned}$$

6 FLOPs  
output

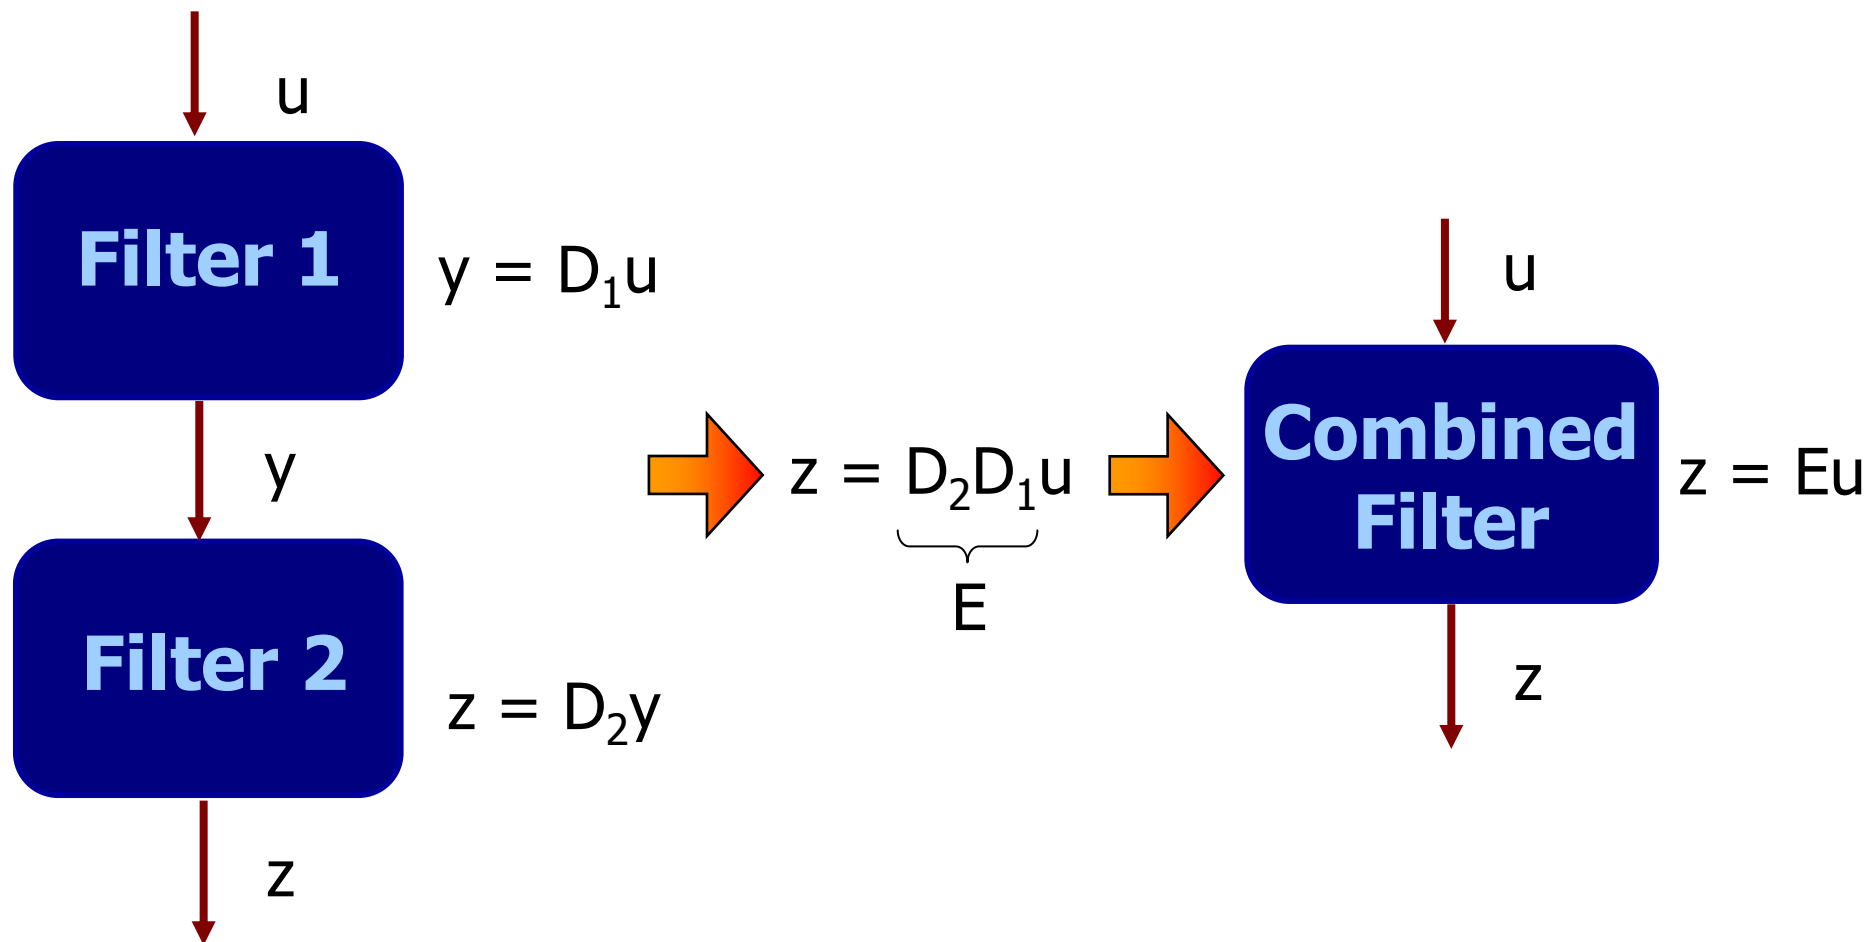


4 FLOPs  
output

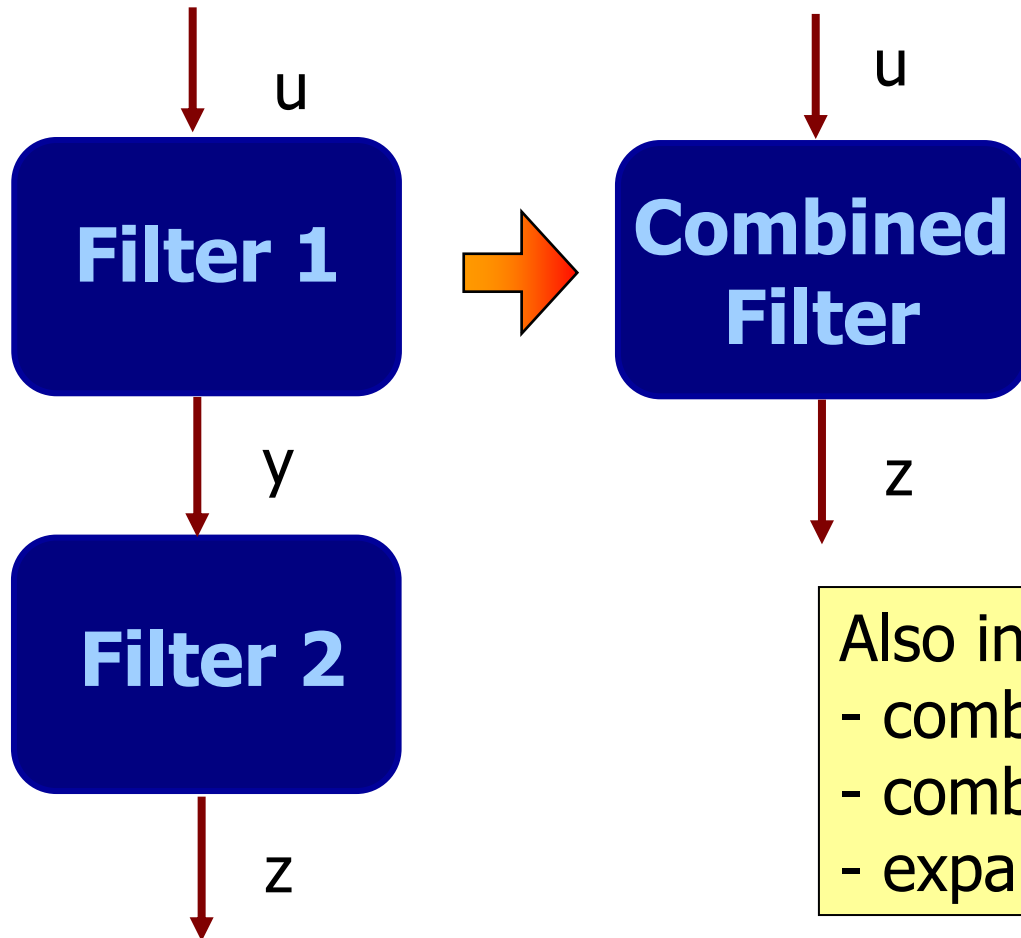




### 3) Combining Adjacent Filters



### 3) Combining Adjacent Filters

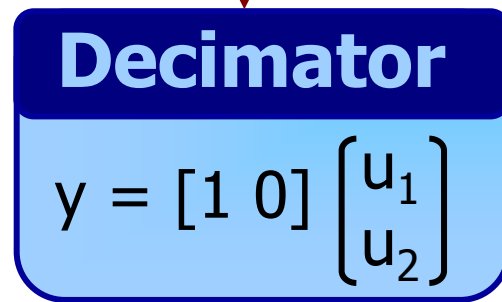



$$x' = \begin{bmatrix} A_1 & 0 \\ B_2 C_1 & A_2 \end{bmatrix} x + \begin{bmatrix} B_1 \\ B_2 D_1 \end{bmatrix} u$$
$$z = \begin{bmatrix} D_2 C_1 & C_2 \end{bmatrix} x + \begin{bmatrix} D_2 D_1 \end{bmatrix} u$$

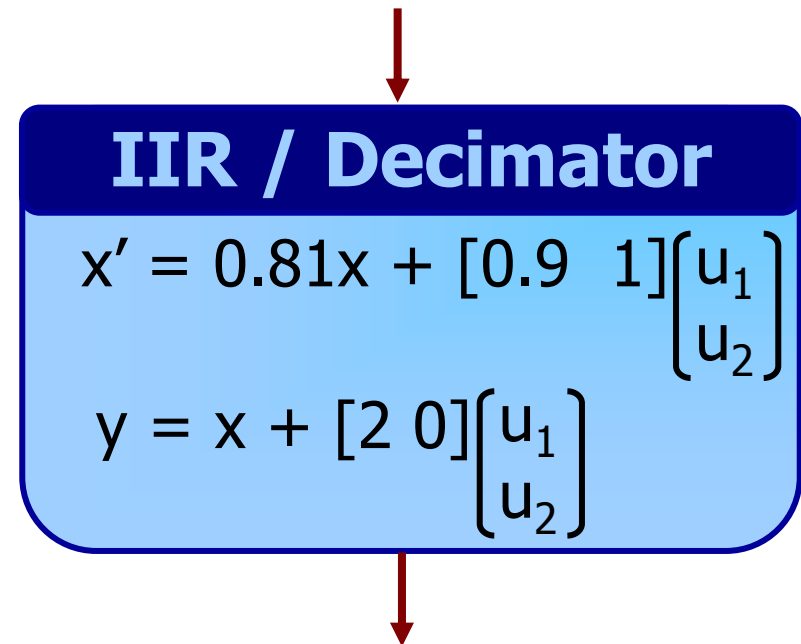

Also in paper:

- combination of parallel streams
- combination of feedback loops
- expansion of mis-matching filters

# Combination Example

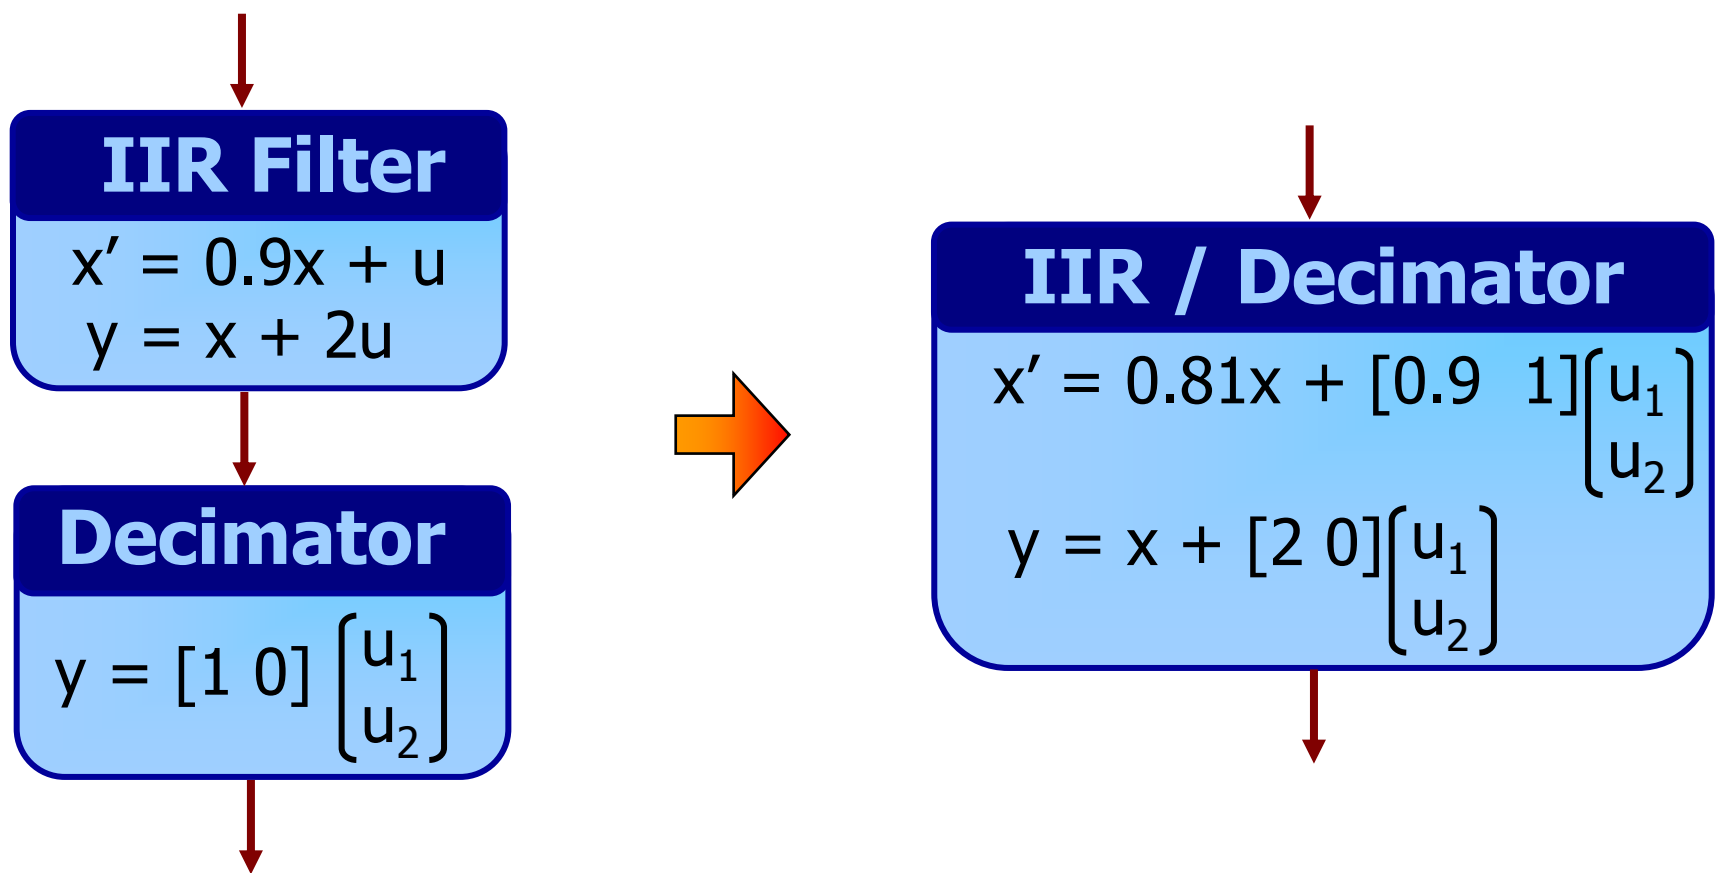


8 FLOPs  
output ☹️



6 FLOPs  
output 😊

# Combination Example



As decimation factor goes to  $\infty$ ,  
eliminate up to 75% of FLOPs.

# Combination Hazards

- Combination sometimes increases FLOPs
- Example: FFT
  - Combination results in DFT
  - Converts  $O(n \log n)$  algorithm to  $O(n^2)$
- Solution: only apply where beneficial
  - Operations known at compile time
  - Using selection algorithm, FLOPs never increase
    - See PLDI '03 paper for details

# Results

- Subsumes combination of linear components
  - Evaluated previously [PLDI '03]
    - **Applications:** FIR, RateConvert, TargetDetect, Radar, FMRadio, FilterBank, Vocoder, Oversampler, DtoA
  - Removed 44% of FLOPs
  - Speedup of 120% on Pentium 4
- Results using state space analysis

	Speedup (Pentium 3)
IIR + 1:2 Decimator	49%
IIR + 1:16 Decimator	87%

# Ongoing Work

- Experimental evaluation
  - Evaluate real applications on embedded machines
  - In progress: MPEG2, JPEG, radar tracker
- Numerical precision constraints
  - Precision often influences choice of coefficients
  - Transformations should respect constraints

# Related Work

- Linear stream optimizations [Lamb et al. '03]
  - Deals with stateless filters
- Automatic optimization of linear libraries
  - SPIRAL, FFTW, ATLAS, Sparsity
- Stream languages
  - Lustre, Esterel, Signal, Lucid, Lucid Synchrone, Brook, Spidle, Cg, Occam , Sisal, Parallel Haskell
- Common sub-expression elimination



# Conclusions

- Linear state space analysis:  
An elegant compiler IR for DSP programs
- Optimizations using state space representation:
  1. State removal
  2. Parameter reduction
  3. Combining adjacent filters
- Step towards adding efficient abstraction layers that remove the DSP expert from the design flow

**StreamIt**

<http://cag.lcs.mit.edu/streamit>